
GloMarGridding

Release 1.2.1

NOC Surface Processes

Aug 14, 2026

CONTENTS:

1	Introduction	1
2	Citation	3
3	Getting Started	5
3.1	Installation	5
3.1.1	Via Pip	5
3.1.2	From Source	5
4	Credits	7
4.1	Development Lead	7
4.2	Contributoring Developers	7
4.3	Acknowledgements	7
4.4	License	7
5	Example Workflow	9
5.1	Load Observations	9
5.2	Output Grid	9
5.3	Align Observations	9
5.4	Create or Load Spatial Covariance	10
5.5	Optionally Load Error Covariance	10
5.6	Ordinary Kriging	10
6	Stationary Interpolation Covariance	13
6.1	Grid	14
6.2	Variograms	16
6.3	Covariance	19
7	Ellipses: Non-Stationary Interpolation Covariance	21
7.1	Ellipse Models	21
7.2	Ellipse Parameter Estimation	23
7.3	Ellipse-based Covariance Estimation	29
8	Error Covariance	33
9	Kriging	39
9.1	Preparation	39
9.2	Simple Kriging	41
9.3	Ordinary Kriging	44
9.4	Perturbed Gridded Fields	47
9.5	Outputs	51

10 Spline Interpolation	53
10.1 Example Workflow	53
10.2 Spline	54
10.3 Thin-Plate-Splines	57
10.4 Spherical Thin-Plate-Splines	60
10.5 Spline Result	62
11 Grid Class	65
12 Auto Regressive Forecasting and Kalman Filter	73
12.1 Introduction and purpose	73
12.2 Producing the prior using a simple autoregressive statistical model	73
12.3 Uncertainty weighted average (Kalman Filter)	74
12.4 Code and usage	75
12.5 <i>autoregressive_kalman.forecast_linear_ar1</i>	75
12.6 <i>autoregressive_kalman.inv_sigma_wgts</i>	79
12.7 <i>autoregressive_kalman.cov_diagonal</i>	81
12.8 <i>autocorrelation</i>	82
13 Miscallaneous Modules	89
13.1 Climatologies	89
13.2 Masking	90
13.3 Distances and Distance Matrices	93
13.4 Covariance Tools and Eigenvalue Clipping	97
13.5 IO	104
13.6 Utilities	105
13.7 Variable Transformation	113
Bibliography	115
Python Module Index	117
Index	119

INTRODUCTION

Global surface temperature datasets - such as those used in the Intergovernmental Panel on Climate Change (IPCC) Assessment Report [IPCC] rely on a range of techniques to generate smoothed, infilled gridded fields from available observations [Lenssen], [Kadow], [RohdeBerkeley], [Morice_2021], [Huang] . The construction of these datasets involves numerous processing decisions, and variations in how each dataset represents the temperature field contribute to what is known as structural uncertainty [Thorne]. This structural uncertainty has two primary sources: (1) differences in the processing of the input temperature measurements, and (2) differences in the spatial interpolation methods applied. Because these steps are often tightly integrated, it is difficult to determine their individual contributions.

glomar_gridding is a software package developed to support the evaluation of structural uncertainty by offering flexible tools for spatial interpolation. The package enables users to spatially interpolate grid-box average observations and their associated uncertainty estimates using Gaussian Process Regression Modelling (GPRM, often called *Kriging*) [Rasmussen], [Cressie] . This technique builds on established methods for generating surface temperature fields [Karspeck], [Morice_2012]. By decoupling interpolation from earlier processing stages - such as homogenization, quality control, and aggregation to grid-cell averages - *glomar_gridding* allows users to create spatially complete fields while independently assessing the effects of upstream data processing choices.

CITATION

Cornes, R. C., S. C.Chan, A.Cable, et al. 2026. "GloMarGridding: A Python Toolkit for Flexible Spatial Interpolation in Climate Applications." *Geoscience Data Journal*, 13 (2): e70064. <https://doi.org/10.1002/gdj3.70064>.

GETTING STARTED

3.1 Installation

3.1.1 Via Pip

GloMarGridding is available on [PyPI](#), and can be installed with *pip* or *uv*:

```
pip install glomar_gridding
```

```
uv add glomar_gridding
```

3.1.2 From Source

Alternatively, you can clone the repository and install using *pip* (or *uv* if preferred).

```
git clone https://github.com/NOCSurfaceProcesses/GloMarGridding.git
cd GloMarGridding
python -m venv venv
source venv/bin/activate
pip install -e .
```

```
git clone https://github.com/NOCSurfaceProcesses/GloMarGridding.git
cd GloMarGridding
uv sync --all-extras --python 3.11
```


4.1 Development Lead

- Agnieszka Faulkner <agfaul@noc.ac.uk> @agfaul
- Joseph T. Siddons <josidd@noc.ac.uk> @josidd

4.2 Contributing Developers

- Archie Cable <acable@noc.ac.uk> @acable
- Steven Chan <stchan@noc.ac.uk> @stchan
- Richard C. Cornes <rcornes@noc.ac.uk> @ricorne
- Elizabeth C. Kent <eck@noc.ac.uk> @eck
- Duo Chan <Duo.Chan@soton.ac.uk>

4.3 Acknowledgements

- Thanks to Simon Williams for providing the original Matlab code that GloMarGridding is based upon.
- Supported by the Natural Environmental Research Council through National Capability funding (AtlantiS: NE/Y005589/1)

4.4 License

Copyright 2025 National Oceanography Centre

Licensed under the Apache License, Version 2.0 (the “License”); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an “AS IS” BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

EXAMPLE WORKFLOW

Here we present a simple example, where we use Ordinary Kriging to interpolate observational data.

A full example can be found in the notebooks directory of the repository.

5.1 Load Observations

In this hypothetical example, we are working with data from a *csv* file. This is assumed to be point-observation data.

```
obs = pl.read_csv("/path/to/obs.csv", ...)
```

5.2 Output Grid

The first step is to create the output grid, using *glomar_gridding.grid.grid_from_resolution()* to specify a global grid with a fixed resolution. Here, the grid has a 5-degree resolution.

The grid is essentially an empty *xarray.DataArray* with a coordinate system defined by the input parameters. This can be used to map to the observations, and to create a covariance matrix with consistent coordinates.

```
grid = grid_from_resolution(  
    resolution=5,  
    bounds=[(-87.5, 90), (-177.5, 180)],  
    coord_names=["latitude", "longitude"],  
)
```

5.3 Align Observations

The input observations may not be located at grid-box locations, i.e. they may be located somewhere between grid-box centres. *glomar_gridding.grid.map_to_grid()* can be used to map each point observation to a grid-box by identifying the nearest grid-box position from an input grid.

This adds an index column to the observational data-frame, which is the 1-dimensional index value of the coordinate. This allows for easy mapping to the indices of covariance matrices, etc.

```
obs = map_to_grid(  
    obs=obs,  
    grid=grid,  
)
```

In this example, it is assumed that the observations are such that at most one observation is associated with a grid-box. If this is not the case, the observations can be combined into a grid-box *super*-observation with a weighting using `glomar_gridding.kriging.prep_obs_for_kriging()`.

Extract the observation values, and the grid-box index for each observation. In this example, the observation value is stored in the “val” column.

```
grid_obs = obs["val"]
grid_idx = obs["grid_idx"]
```

5.4 Create or Load Spatial Covariance

The grid can be converted in to a distance matrix, and finally to a covariance matrix using a `glomar_gridding.variogram.Variogram` object, for example `glomar_gridding.variogram.GaussianVariogram`.

```
dist = grid_to_distance_matrix(
    grid=grid,
    lat_coord="latitude",
    lon_coord="longitude",
)

variogram = GaussianVariogram(
    range=1200,
    psill=1.2,
    nugget=0.0,
).fit(dist)

covariance = variogram_to_covariance(variogram, variance=1.2)
```

Alternatively, the covariance matrix can be loaded from disk. A non-stationary (varying parameter) covariance matrix can be estimated using ellipse-based models. See `glomar_gridding.ellipse.EllipseModel`.

5.5 Optionally Load Error Covariance

In this example an error covariance matrix is loaded from a netCDF file on disk, using `glomar_gridding.io.load_array()`

```
error_cov = load_array("/path/to/error_cov.nc", var="error_covariance")
```

Alternatively, an error covariance matrix can be computed component wise.

5.6 Ordinary Kriging

In this example, we will infill the observations using Ordinary Kriging. For this, we use `glomar_gridding.kriging.OrdinaryKriging`, which requires a spatial covariance matrix, observation grid indices, observation values, and (optionally) error covariance as inputs.

```
ok = OrdinaryKriging(
    covariance.values,
    idx=grid_idx,
    obs=grid_obs,
    error_cov=error_cov,
)
```

We can now use this class-instance to solve the system, using the *solve* method.

```
result = ok.solve()
```

Finally, the output can be mapped back on to the grid using *glomar_gridding.grid.assign_to_grid()*

```
gridded_result = assign_to_grid(  
    values=result,  
    grid_idx=np.arange(grid.size),  
    grid=grid,  
)
```


STATIONARY INTERPOLATION COVARIANCE

Outside of the observation values and positions, the spatial covariance structure is the most important component for Kriging. The Kriging classes in this library all require this matrix as the input to the class. This covariance matrix can be provided as a pre-computed matrix, loaded into the environment with `glommar_gridding.interpolation_covariance.load_covariance()`. Alternatively the covariance structure can be estimated using functions and classes contained within `glommar_gridding`.

A commonly used approach is to compute a stationary covariance structure, using a *Variogram* with fixed scales. The use of *stationary* here suggests that the range of covariance is constant across all positions - each position has the same influence over a fixed distance.

```
from glommar_gridding.grid import grid_from_resolution, grid_to_distance_matrix
from glommar_gridding.variogram import GaussianVariogram, variogram_to_covariance

# Initialise a grid
grid = grid_from_resolution(
    resolution=5,
    bounds=[(-87.5, 90), (-177.5, 180)],
    coord_names=["latitude", "longitude"],
)

# Compute a distance matrix
dist = grid_to_distance_matrix(
    grid=grid,
    lat_coord="latitude",
    lon_coord="longitude",
)

# Define and compute a Variogram
variogram = GaussianVariogram(
    nugget=0.0,
    psill=1.2,
    range=1300,
).fit(dist)

# Convert to covariance
covariance = variogram_to_covariance(variogram, sill=1.2)
```

6.1 Grid

`glomar_gridding.grid.grid_from_resolution(resolution, bounds, coord_names, definition='center')`

Generate a grid from a resolution value, or a list of resolutions for given boundaries and coordinate names.

Note that all list inputs must have the same length, the ordering of values in the lists is assumed align.

The constructed grid will be regular, in the sense that the grid spacing is constant. However, the resolution in each direction can be different, allowing for finer resolution in some direction.

Parameters

- **resolution** (float | list[float]) – Resolution of the grid. Can be a single resolution value that will be applied to all coordinates, or a list of values mapping a resolution value to each of the coordinates.
- **bounds** (list[tuple[float, float]]) – A list of bounds of the form (*lower_bound*, *upper_bound*) indicating the bounding box of the returned grid. For a grid between -180 and 180 degrees in Longitude, one would set the bounds in this dimension to (-180, 180) (if using “left” for the ‘definition’ argument).
- **coord_names** (list[str]) – List of coordinate names in the same order as the bounds and resolution(s).
- **definition** (Literal['center', 'left']) – Each grid box is defined by the centre of the box in each coordinate. This argument defines whether the lower bound in each coordinate direction defines the “left”-edge or the “center” of the 1st grid-box.
 - “left”: The lower-bound can be used to define the left-side of each grid-box. The grid-box value will be the left-edge + 1/2 of resolution in each direction.
 - “center”: The lower-bound can be used to define the center of each grid-box. The grid-box value will be the center in each direction

Returns

grid – The grid defined by the resolution and bounding box.

Return type

DataArray

Examples

```
>>> grid_from_resolution(
    resolution=5,
    bounds=[(-87.5, 90), (-177.5, 180)], # Lower bound is centre
    coord_names=["lat", "lon"]
)
<xarray.DataArray (lat: 36, lon: 72)> Size: 21kB
array([[nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       ...,
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan]], shape=(36, 72))
Coordinates:
  * lat      (lat) float64 288B -87.5 -82.5 -77.5 ... 77.5 82.5 87.5
  * lon      (lon) float64 576B -177.5 -172.5 ... 172.5 177.5
```

```
glomar_gridding.grid.grid_to_distance_matrix(grid, dist_func=<function
                                             haversine_distance_from_frame>, **dist_kwargs)
```

Calculate a distance matrix between all positions in a grid. Orientation of latitude and longitude will be maintained in the returned distance matrix.

Parameters

- **grid** (DataArray) – A 2-d grid containing latitude and longitude indexes specified in decimal degrees.
- **dist_func** (Callable) – Distance function to use to compute pairwise distances. See `glomar_gridding.distances.calculate_distance_matrix` for more information.
- ****dist_kwargs** – Keyword arguments to pass to the distance function. This may include requirements for the name of specific coordinates, for example “lat_coord” and “lon_coord”.

Returns

dist – A DataArray containing the distance matrix with coordinate system defined with grid cell index (“index_1” and “index_2”). The coordinates of the original grid are also kept as coordinates related to each index (the coordinate names are suffixed with “_1” or “_2” respectively).

Return type

DataArray

Examples

```
>>> grid = grid_from_resolution(
    resolution=5,
    bounds=[(-90, 90), (-180, 180)], # Lower bound is left-edge
    coord_names=["lat", "lon"],
    definition="left",
)
>>> grid_to_distance_matrix(grid, lat_coord="lat", lon_coord="lon")
<xarray.DataArray 'dist' (index_1: 2592, index_2: 2592)> Size: 54MB
array([[ 0.,          24.24359308,  48.44112457, ...,
        19463.87158499, 19461.22915012, 19459.64166305],
       [ 24.24359308,  0.,          24.24359308, ...,
        19467.56390938, 19463.87158499, 19461.22915012],
       [ 48.44112457,  24.24359308,  0.,          ...,
        19472.29905588, 19467.56390938, 19463.87158499],
       ...,
       [19463.87158499, 19467.56390938, 19472.29905588, ...,
        0.,          24.24359308,  48.44112457],
       [19461.22915012, 19463.87158499, 19467.56390938, ...,
        24.24359308,  0.,          24.24359308],
       [19459.64166305, 19461.22915012, 19463.87158499, ...,
        48.44112457,  24.24359308,  0.          ]],
      shape=(2592, 2592))
Coordinates:
* index_1    (index_1) int64 21kB 0 1 2 3 4 ... 2587 2588 2589 2590 2591
* index_2    (index_2) int64 21kB 0 1 2 3 4 ... 2587 2588 2589 2590 2591
lat_1        (index_1) float64 21kB -87.5 -87.5 -87.5 ... 87.5 87.5
lon_1        (index_1) float64 21kB -177.5 -172.5 ... 172.5 177.5
lat_2        (index_2) float64 21kB -87.5 -87.5 -87.5 ... 87.5 87.5 87.5
lon_2        (index_2) float64 21kB -177.5 -172.5 ... 172.5 177.5
```

6.2 Variograms

```
class glomar_gridding.variogram.ExponentialVariogram(psill, nugget, range=None,
                                                    effective_range=None)
```

Exponential Model

Parameters

- **psill** (float | ndarray) – Sill of the variogram where it will flatten out. Values in the variogram will not exceed psill + nugget. This value is the variance.
- **nugget** (float | ndarray) – The value of the independent variable at distance 0
- **effective_range** (float | ndarray | None) – Effective Range, this is the distance where 95% of the sill are exceeded. This is not the range parameter, which is defined as $r/3$.
- **range** (float | ndarray | None) – The range parameter. One of range and effective_range must be set. If range is not set, it will be computed from effective_range.

```
fit(distance_matrix)
```

Fit the ExponentialVariogram model to a distance matrix

Parameters

distance_matrix (ndarray | DataArray) – The distance matrix indicating the distance between each pair of points in the grid.

Returns

A matrix containing the variogram values at each distance.

Return type

ndarray | DataArray

Examples

```
>>> ExponentialVariogram(range=1200, psill=1.2, nugget=0.0).fit(dist)
```

```
class glomar_gridding.variogram.SphericalVariogram(psill, nugget, effective_range=None,
                                                    range=None)
```

Spherical Model

Parameters

- **psill** (float | ndarray) – Sill of the variogram where it will flatten out. Values in the variogram will not exceed psill + nugget. This value is the variance.
- **nugget** (float | ndarray) – The value of the independent variable at distance 0
- **effective_range** (float | ndarray | None) – Effective Range, this is the distance where 95% of the sill are exceeded. This is not the range parameter, which is equal to the effective range in the SphericalVariogram case.
- **range** (float | ndarray | None) – The range parameter. One of range and effective_range must be set. If range is not set, it will be computed from effective_range.

```
fit(distance_matrix)
```

Fit the SphericalVariogram model to a distance matrix

Parameters

distance_matrix (ndarray | DataArray) – The distance matrix indicating the distance between each pair of points in the grid.

Returns

A matrix containing the variogram values at each distance.

Return type

ndarray | DataArray

Examples

```
>>> SphericalVariogram(range=1200, psill=1.2, nugget=0.0).fit(dist)
```

class glomar_gridding.variogram.**GaussianVariogram**(psill, nugget, effective_range=None, range=None)
Gaussian Model

Parameters

- **psill** (float | ndarray) – Sill of the variogram where it will flatten out. Values in the variogram will not exceed psill + nugget. This value is the variance.
- **nugget** (float | ndarray) – The value of the independent variable at distance 0
- **effective_range** (float | ndarray | None) – Effective Range, this is the distance where 95% of the sill are exceeded. This is not the range parameter, which is defined as $r/2$.
- **range** (float | ndarray | None) – The range parameter. One of range and effective_range must be set. If range is not set, it will be computed from effective_range.

fit(distance_matrix)

Fit the GaussianVariogram model to a distance matrix

Parameters

distance_matrix (ndarray | DataArray) – The distance matrix indicating the distance between each pair of points in the grid.

Returns

A matrix containing the variogram values at each distance.

Return type

ndarray | DataArray

Examples

```
>>> GaussianVariogram(range=1200, psill=1.2, nugget=0.0).fit(dist)
```

class glomar_gridding.variogram.**MaternVariogram**(psill, nugget, effective_range=None, range=None, nu=0.5, method='sklearn')

Matern Models

Same args as the Variogram classes with additional nu (ν), method parameters.

Sklearn:

- 1) This is called “sklearn” because if $d/\text{range} = 1.0$ and $\nu = 0.5$, it gives $1/e$ correlation...
- 2) This is NOT the same formulation as in GSTAT nor in papers about non-stationary anisotropic covariance models (aka Karspeck paper).
- 3) It is perhaps the most intuitive (because of (1)) and is used in sklearn GP and HadCRUT5 and other UKMO dataset.
- 4) nu defaults to 0.5 (exponential; used in HADSST4 and our kriging). HadCRUT5 uses 1.5.
- 5) The “2” is inside the square root for middle and right.

GeoStatic:

Similar to Sklearn MaternVariogram model but uses the range scaling in gstat. Note: there are no square root 2 or nu in middle and right

Yields the same answer to sklearn MaternVariogram if $\nu = 0.5$ but are otherwise different.

Karspeck:

Similar to Sklearn MaternVariogram model but uses the form in Karspeck paper Note: Note the 2 is outside the square root for middle and right e-folding distance is now at $d/\text{SQRT}(2)$ for $\nu = 0.5$

Parameters

- **psill** (float | ndarray) – Sill of the variogram where it will flatten out. Values in the variogram will not exceed psill + nugget. This value is the variance.
- **nugget** (float | ndarray) – The value of the independent variable at distance 0
- **effective_range** (float | ndarray | None) – Effective Range, this is the lag where 95% of the sill are exceeded. This is not the range parameter, which is defined as $r/3$ if $\nu < 0.5$ or $\nu > 10$, otherwise $r/2$ (where r is the effective range). One of effective_range and range must be set.
- **range** (float | ndarray | None) – The range parameter. One of range and effective_range must be set. If range is not set, it will be computed from effective_range.
- **nu** (float | ndarray) – ν , smoothing parameter, shapes to a smooth or rough variogram function
- **method** (Literal['sklearn', 'gstat', 'karspeck']) – One of “sklearn”, “gstat”, or “karspeck”:
 - sklearn: https://scikit-learn.org/stable/modules/generated/sklearn.gaussian_process.kernels.Matern.html#sklearn.gaussian_process.kernels.Matern
 - gstat: <https://scikit-gstat.readthedocs.io/en/latest/reference/models.html#matern-model>
 - karspeck: <https://rmets.onlinelibrary.wiley.com/doi/10.1002/qj.900>

References

see chapter 4.2 of Rasmussen, C. E., & Williams, C. K. I. (2005). Gaussian Processes for Machine Learning. The MIT Press. <https://doi.org/10.7551/mitpress/3206.001.0001>

fit(distance_matrix)

Fit the MaternVariogram model to a distance matrix.

Parameters

distance_matrix (ndarray | DataArray) – The distance matrix indicating the distance between each pair of points in the grid.

Returns

A matrix containing the variogram values at each distance.

Return type

ndarray | DataArray

Examples

```
>>> MaternVariogram(
    range=1200, psill=1.2, nugget=0.0, nu=1.5, method="karspeck"
).fit(dist)
```

`glomar_gridding.variogram.variogram_to_covariance(variogram, variance)`

Convert a variogram matrix to a covariance matrix.

This is given by:

$\text{covariance} = \text{variance} - \text{variogram}$

Parameters

- **variogram** (ndarray | DataArray) – The variogram matrix, output of Variogram.fit.
- **variance** (ndarray | float) – The variance

Returns

cov – The covariance matrix

Return type

ndarray | DataArray

6.3 Covariance

I/O functionality for loading a covariance matrix from disk.

`glomar_gridding.interpolation_covariance.load_covariance(path, cov_var_name='covariance', **kwargs)`

Load a covariance matrix from a netCDF file. Can input a filename or a string to format with keyword arguments.

Parameters

- **path** (str) – Full filename (including path), or filename with replacements using str.format with named replacements. For example: `/path/to/global_covariance_{month:02d}.nc`
- **cov_var_name** (str) – Name of the variable for the covariance matrix
- ****kwargs** – Keywords arguments matching the replacements in the input path.

Returns

covariance – A numpy matrix containing the covariance matrix loaded from the netCDF file determined by the input arguments.

Return type

ndarray

ELLIPSES: NON-STATIONARY INTERPOLATION COVARIANCE

7.1 Ellipse Models

```
class glomar_gridding.ellipse.EllipseModel(anisotropic, rotated, physical_distance, v,  
                                             unit_sigma=False)
```

The class that contains variogram/ellipse fitting methods and parameters

This class assumes your input to be a standardised correlation matrix. They are easier to handle because stdevs in the covariance function become 1

Parameters

- **anisotropic** (bool) – Should the output be an ellipse? Set to False for circle.
- **rotated** (bool) – Can the ellipse be rotated. If anisotropic is False this value cannot be True.
- **physical_distance** (bool) – Use physical distances rather than lat/lon distance.
- **v** (float) – Matern Shape Parameter. Must be > 0.0.
- **unit_sigma=True** (bool) – When MLE fitting the Matern parameters, assuming the Matern parameters themselves are normally distributed, there is standard deviation within the log likelihood function.

See Wikipedia entry for Maximum Likelihood under: - Continuous distribution, continuous parameter space

Its actual value is not important to the best (MLE) estimate of the Matern parameters. If one assumes the parameters are normally distributed, the mean (best estimate) is independent of its variance. In fact in Karspeck et al 2012 [Karspeck], it is simply set to 1 Eq B1). This value can however be computed. It serves a similar purpose as the original standard deviation: in this case, how the actual observed semivariance disperses around the fitted variogram.

The choice to default to 1 follows Karspeck et al. 2012 [Karspeck]

```
fit(X, y, guesses=None, bounds=None, opt_method='Nelder-Mead', tol=None,  
    estimate_SE='bootstrap_parallel', n_sim=500, n_jobs=4, backend='loky', random_seed=1234,  
    options=None)
```

Default solver in Nelder-Mead as used in Karspeck et al. 2012 [Karspeck] i.e. <https://docs.scipy.org/doc/scipy/reference/optimize.minimize-neldermead.html> default max-iter is 200 x (number_of_variables) for 3 variables (Lx, Ly, theta) → 200x3 = 600 note: unlike variogram fitting, no nugget, no sill, and no residue variance (normalised data but Fisher transform needed?) can be adjusted using “maxiter” within “options” kwargs

Much of the variable names are defined the same way as earlier

Parameters

- **X** (ndarray) – Array of displacements. Expected to be 1-dimensional if the ellipse model is not anisotropic, 2-dimensional otherwise. In units of km if the ellipse uses physical distances, otherwise in degrees. The displacements are from each position within the test region to the centre of the ellipse.
- **y** (ndarray) – Vector of observed correlations between the centre of the ellipse and each test point.
- **guesses=None** (*list[float] | None*) – List of initial values to `scipy.optimize.minimize`, default guesses for the ellipse model are used if not set.
- **bounds=None** (*list[tuple[float, float]] | None*) – Tuples/lists of bounds for fitted parameters. Default bounds for the ellipse model are used if not set.
- **opt_method** (str) – `scipy.optimize.minimize` optimisation method. Defaults to “Nelder-Mead”. See <https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.minimize.html> for valid values.
- **tol=None** (*float | None*) – `scipy.optimize.minimize` convergence tolerance
- **estimate_SE='bootstrap_parallel'** (*str | None*) – How to estimate standard error if needed. If not set no standard error is computed.
- **n_sim=500** (*int*) – Number of bootstrap to estimate standard error
- **n_jobs=DEFAULT_N_JOBS** (*int*) – Number of threads for bootstrapping if *estimate_SE* is set to “bootstrap_parallel”.
- **backend=DEFAULT_BACKEND** (*str*) – joblib backend for bootstrapping.
- **random_seed=1234** (*int*) – Random seed for bootstrap
- **options** (*dict[str, Any] | None*) – Options to use in fitting.

Return type

`tuple[OptimizeResult, float | None, list[tuple[float, float]]]`

Returns

- **results** (*OptimizeResult*) – Output of `scipy.optimize.minimize`
- **SE** (*float | None*) – Standard error of the fitted parameters
- **bounds** (*list[tuple[float, ...]]*) – Bounds of fitted parameters

negative_log_likelihood(*params, X, y, arctanh_transform=True*)

Compute the negative log-likelihood given observed *X* independent observations (displacements) and *y* dependent variable (the observed correlation), and Matern parameters *params*. Namely does the Matern covariance function using *params*, how close it explains the observed displacements and correlations.

$\log(LL) = \text{SUM} (f(y, x | \text{params}))$ *params* = Maximise ($\log(LL)$) *params* = Minimise ($-\log(LL)$) which is how usually the computer solves it assuming errors of *params* are normally distributed

There is a hidden scale/standard deviation in `stats.norm.logpdf(scale, ...)` which defaults to 1) but since we have scaled our values to covariance to correlation (and even used Fisher transform) as part of the function, it can be dropped

Otherwise, you need to have *stdev* as the last value of *params*, and should be set to the scale parameter

Parameters

- **X** (ndarray) – Observed displacements to the centre of the ellipse.
- **y** (ndarray) – Observed correlation against the centre of the ellipse.

- **params** (list[float]) – Ellipse parameters (in the current optimize iteration) or if you want to compute the actual negative log-likelihood.
- **arctanh_transform** (bool) – Should the Fisher (arctanh) transform be used This is usually option, but it does make the computation more stable if they are close to 1 (or -1; doesn't apply here)

Returns

nLL – The negative log likelihood

Return type

float

negative_log_likelihood_function(X, y)

Creates a function that can be fed into `scipy.optimize.minimize`

Return type

Callable[[list[float]], float]

7.2 Ellipse Parameter Estimation

class `glomar_gridding.ellipse.EllipseBuilder`(data_array, coords)

Class to build spatial covariance and correlation matrices used to estimate ellipse parameters using an instance of `EllipseModel` which sets up the defaults for a given configuration.

To fit ellipse parameters to the correlation of the input `data_array`, call `self.fit_ellipse_model`.

Parameters

- **data_array** (ndarray) – Training data stored within a numpy array. In general, this input should be extracted from an `xarray.DataArray`, and masked appropriately.
- **coords** (Coordinates) – The coordinates associated with the `data_array` value. It is expected that these are ["time", "latitude", "longitude"]

calc_cov(rounding=None)

Calculate covariance and correlation matrices.

Parameters

rounding (int | None) – Round the values of the output.

Return type

None

compute_params(default_value, matern_ellipse, max_distance=6000, min_distance=0.3, delta_x_method='Modified_Met_Office', guesses=None, bounds=None, opt_method='Nelder-Mead', tol=0.001, estimate_SE=None, n_jobs=4, n_sim=500, physical_distance_selection=True, options=None)

Fit ellipses/covariance models using local covariances to all unmasked grid points.

The form of the covariance model depends on the 'anisotropic' and 'rotated' attributes of the ellipse model:

- `anisotropic=False` (radial distance only)
- `anisotropic=True` and `rotated=False` (x and y are different, but not rotated)
- `anisotropic=True` and `rotated=True` (x and y are different, ellipse can be rotated)

The ellipse model attribute `v` = matern covariance function shape parameter. Karspeck et al. [Karspeck] and Paciorek & Schervish [PaciorekSchervish] use 3 and 4 but 0.5 and 1.5 are popular. 0.5 gives an exponential decay: $\lim v \rightarrow \infty$, Gaussian shape

delta_x_method: only meaningful for physical_distance ellipses:

- “Met_Office”: Cylindrical Earth $\Delta x = 6400\text{km} \times \Delta \text{lon}$ (in radians)
- “Modified_Met_Office”: uses the average zonal dist at different lat

Parameters

- **default_value** (Any) – Default value(s) to fill arrays where parameter estimation is not possible (typically due to masking). Typically, one should set a value that is appropriate to the type of the field. If a single value is provided, this is used for all fields. If not, the length of the list of default values must equal the number of parameters of the *EllipseModel*
- **matern_ellipse** (*EllipseModel*) – EllipseModel to use for parameter estimation
- **max_distance** (float) – Maximum separation in distance unit that data will be fed into parameter fitting Units depend on physical_distance attribute (km if True, otherwise degrees).
- **min_distance** (float) – Minimum separation in distance unit that data will be fed into parameter fitting Units depend on physical_distance attribute (km if True, otherwise degrees). Note: Due to the way we compute the Matern function, it is undefined at $\text{dist} == 0$ even if the limit $\rightarrow 0$ is obvious.
- **delta_x_method=“Modified_Met_Office”** (str) – How to compute distances between grid points For isotropic variogram/covariances, this is a trivial problem; you can just take the haversine or Euclidean (“tunnel”) distance as they are non-directional.

But it is non trivial for anisotropic cases, you have to define a set of orthogonal space. In HadSST4, Earth is assumed to be cylindrical “tin can” Earth, so you can just define the orthogonal space by lines of constant lat and lon (delta_x_method=“Met_Office”).

The modified “Modified_Met_Office” is a variation to that, but allow the tin can get squished at the poles. (Sinusoidal projection). This does results in a problem: the zonal displacement now depends in which latitude you compute on (at the beginning latitude or at the end latitude). Here we take the average of the two.

- **guesses=None** (tuple of floats; None uses default guess values) – Initial guess values that get feeds in the optimizer for MLE. In scipy, you are required to do so (but R often doesn’t). You should anyway; sometimes they do funny things if you don’t (per recommendation of David Stephenson)
- **bounds=None** (tuple of floats; None uses default bounds values) – This is essentially a Bayesian “uninformative prior” that forces convergence if the optimizer hits the bound. For lower resolution fitting, this is rarely a problem. For higher resolution fits, this often interacts with the limit of the data you can put into the fit the optimizer may fail to converge if the input data is very smooth (aka ENSO region, where anomalies are smooth over very large (~10000km) scales).
- **opt_method=‘Nelder-Mead’** (str) – scipy.optimize method. Nelder-Mead is the one used by [Karspeck]. See <https://docs.scipy.org/doc/scipy/tutorial/optimize.html> for valid options
- **tol=0.001** (float) – Set convergence tolerance for scipy optimize. See <https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.minimize.html#scipy.optimize.minimize>

Note on new tol kwarg: For N-M, this sets the value to both xatol and fatol Default is 1E-3 Since it affects accuracy of all values including rotation rotation angle 0.001 rad ~ 0.05 deg

- **estimate_SE=None** (*str* / *None*) – The code can estimate the standard error if the Matern parameters. This is not usually used or discussed for the purpose of kriging. Certain *opt_method* (gradient descent) can do this automatically using Fisher Info for certain covariance function, but is not possible for some nasty functions (aka Bessel func) gets involved nor it is possible for some optimisers (such as Nelder-Mead). The code does it using bootstrapping.
- **n_jobs=DEFAULT_N_JOBS** (*int*) – If parallel processing, number of threads to use.
- **n_sim** (*int*) – Number of simulations to bootstrap for SE estimation.
- **physical_distance_selection** (*bool*) – Select training data using physical distance (haversine distance) or Euclidean distance of degree difference between each position and all possible training data. Data that falls within the min and max distance values using the selected distance method are selected as training data.
- **options** (*dict[str, Any]* | *None*) – Options to use in fitting.

Returns

params – Containing arrays for each parameter in the ellipse model class. Note that one array is likely to be “qc_code”, which takes values:

- -1: Not computed (masked)
- 0: success
- 2: success but with one parameter reaching upper boundaries
- 3: success with multiple parameters reaching the boundaries (aka both Lx and Ly), can be both at lower or upper bound.
- 9: fail, probably due to running out of maxiter (see *scipy.optimize.minimize* kwargs “options”)

Return type

Dataset

find_nearest_xy_index_in_cov_matrix(*lonlat*, *use_full=False*)

Find the nearest column/row index of the covariance that corresponds to a specific lat lon

Return type

tuple[int, ndarray]

fit_ellipse_model(*xy_point*, *matern_ellipse*, *max_distance=6000*, *min_distance=0.3*, *delta_x_method='Modified_Met_Office'*, *guesses=None*, *bounds=None*, *opt_method='Nelder-Mead'*, *tol=0.001*, *estimate_SE=None*, *n_jobs=4*, *n_sim=500*, *physical_distance_selection=True*, *options=None*)

Fit ellipses/covariance models using local covariances.

The form of the covariance model depends on the ‘anisotropic’ and ‘rotated’ attributes of the ellipse model:

- *anisotropic=False* (radial distance only)
- *anisotropic=True* and *rotated=False* (x and y are different, but not rotated)
- *anisotropic=True* and *rotated=True* (x and y are different, ellipse can be rotated)

The ellipse model attribute *v* = matern covariance function shape parameter. Karspeck et al. [Karspeck] and Paciorek & Schervish [PaciorekSchervish] use 3 and 4 but 0.5 and 1.5 are popular. 0.5 gives an exponential decay: $\lim v \rightarrow \infty$, Gaussian shape

delta_x_method: only meaningful for *physical_distance* ellipses:

- “Met_Office”: Cylindrical Earth $\Delta x = 6400\text{km} \times \Delta \text{lon}$ (in radians)
- “Modified_Met_Office”: uses the average zonal dist at different lat

Parameters

- **xy_point** (int) – The index point where ellipses will be fitted to
- **max_distance** (float) – Maximum separation in distance unit that data will be fed into parameter fitting Units depend on physical_distance attribute (km if True, otherwise degrees).
- **min_distance** (float) – Minimum separation in distance unit that data will be fed into parameter fitting Units depend on physical_distance attribute (km if True, otherwise degrees). Note: Due to the way we compute the Matern function, it is undefined at $\text{dist} == 0$ even if the limit $\rightarrow$ zero is obvious.
- **delta_x_method=“Modified_Met_Office”** (str) – How to compute distances between grid points For isotropic variogram/covariances, this is a trivial problem; you can just take the haversine or Euclidean (“tunnel”) distance as they are non-directional.

But it is non trivial for anisotropic cases, you have to define a set of orthogonal space. In HadSST4, Earth is assumed to be cylindrical “tin can” Earth, so you can just define the orthogonal space by lines of constant lat and lon ($\Delta x_{\text{method}} = \text{“Met_Office”}$).

The modified “Modified_Met_Office” is a variation to that, but allow the tin can get squished at the poles. (Sinusoidal projection). This does results in a problem: the zonal displacement now depends in which latitude you compute on (at the beginning latitude or at the end latitude). Here we take the average of the two.

- **guesses=None** (tuple of floats; None uses default guess values) – Initial guess values that get feeds in the optimizer for MLE. In scipy, you are required to do so (but R often doesn’t). You should anyway; sometimes they do funny things if you don’t (per recommendation of David Stephenson)
- **bounds=None** (tuple of floats; None uses default bounds values) – This is essentially a Bayesian “uninformative prior” that forces convergence if the optimizer hits the bound. For lower resolution fitting, this is rarely a problem. For higher resolution fits, this often interacts with the limit of the data you can put into the fit the optimizer may fail to converge if the input data is very smooth (aka ENSO region, where anomalies are smooth over very large (~10000km) scales).
- **opt_method=‘Nelder-Mead’** (str) – scipy.optimize method. Nelder-Mead is the one used by [Karspeck]. See <https://docs.scipy.org/doc/scipy/tutorial/optimize.html> for valid options
- **tol=0.001** (float) – Set convergence tolerance for scipy optimize. See <https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.minimize.html#scipy.optimize.minimize>

Note on new tol kwarg: For N-M, this sets the value to both xatol and fatol Default is 1E-3 (?) Since it affects accuracy of all values including rotation rotation angle 0.001 rad ~ 0.05 deg

- **estimate_SE=None** (str / None) – The code can estimate the standard error if the Matern parameters. This is not usually used or discussed for the purpose of kriging. Certain opt_method (gradient descent) can do this automatically using Fisher Info for certain covariance function, but is not possible for some nasty functions (aka Bessel func) gets involved nor it is possible for some optimisers (such as Nelder-Mead). The code does it using bootstrapping.

- **n_jobs=DEFAULT_N_JOBS** (*int*) – If parallel processing, number of threads to use.
- **n_sim** (*int*) – Number of simulations to bootstrap for SE estimation.
- **physical_distance_selection** (*bool*) – Select training data using physical distance (haversine distance) or Euclidean distance of degree difference between each position and all possible training data. Data that falls within the min and max distance values using the selected distance method are selected as training data.
- **options** (*dict[str, Any] | None*) – Options to use in fitting.

Returns

Dictionary with results of the fit and the observed correlation matrix.

Return type

dict[str, Any] | None

```
glomar_gridding.ellipse.estimate.get_ellipse_params(ellipse_model, ellipse_builder, n_par_jobs=1,
                                                    par_backend='loky', batch_size=None,
                                                    max_distance=6000, min_distance=0.3,
                                                    delta_x_method='Modified_Met_Office',
                                                    physical_distance_selection=True,
                                                    bounds=None, guesses=None,
                                                    method='Nelder-Mead', tol=0.001,
                                                    default_value=nan, options=None)
```

Get ellipse parameters using a batched parallel process.

Compute all ellipse parameters at all unmasked positions for an input EllipseBuilder object.

Makes use of a batched approach to precompute training data for a batch of positions, allowing parallel processing with duplication of the EllipseBuilder class.

Parameters

- **ellipse_model** (*EllipseModel*) – The ellipse model used to define the parameters of the system.
- **ellipse_builder** (*EllipseBuilder*) – The ellipse build object used to define the training data used for estimating the ellipse parameters.
- **n_par_jobs** (*int*) – Number of parallel jobs to spawn, defaults to 1.
- **par_backend** (*str*) – Parallel backend to use with ‘joblib.Parallel’. Defaults to “loky”.
- **batch_size** (*int | None*) – Size of the batches used to precompute training data. Defaults to 8 x n_par_jobs.
- **max_distance** (*float*) – Maximum separation in distance unit that data will be fed into parameter fitting Units depend on physical_distance attribute (km if True, otherwise degrees). Defaults to 6000 km.
- **min_distance** (*float*) – Minimum separation in distance unit that data will be fed into parameter fitting Units depend on physical_distance attribute (km if True, otherwise degrees). Note: Due to the way we compute the Matern function, it is undefined at dist == 0 even if the limit -> zero is obvious. Defaults to 0.3 km.
- **delta_x_method** (*Literal['Met_Office', 'Modified_Met_Office']*) – How to compute distances between grid points For isotropic variogram/covariances, this is a trivial problem; you can just take the haversine or Euclidean (“tunnel”) distance as they are non-directional.

But it is non trivial for anistropic cases, you have to define a set of orthogonal space. In HadSST4, Earth is assumed to be cylindrical “tin can” Earth, so you can just define the orthogonal space by lines of constant lat and lon (delta_x_method=“Met_Office”).

The modified “Modified_Met_Office” is a variation to that, but allow the tin can get squished at the poles. (Sinusoidal projection). This does results in a problem: the zonal displacement now depends in which latitude you compute on (at the beginning latitude or at the end latitude). Here we take the average of the two.

Defaults to “Modified_Met_Office”.

- **physical_distance_selection** (bool) – Select training data using physical distance (haversine distance) or Euclidean distance of degree difference between each position and all possible training data. Data that falls within the min and max distance values using the selected distance method are selected as training data.
 - **guesses** (list[float] | None) – Initial guess values that get feeds in the optimizer for MLE. In scipy, you are required to do so (but R often doesn’t). You should anyway; sometimes they do funny things if you don’t (per recommendation of David Stephenson). Defaults to the defaults for the input ellipse_model.
 - **bounds** (list[tuple[float, float]] | None) – This is essentially a Bayesian “uninformative prior” that forces convergence if the optimizer hits the bound. For lower resolution fitting, this is rarely a problem. For higher resolution fits, this often interacts with the limit of the data you can put into the fit the optimizer may fail to converge if the input data is very smooth (aka ENSO region, where anomalies are smooth over very large (~10000km) scales). Defaults to defaults for the input ellipse_model.
 - **method** (str) – scipy.optimize method. Nelder-Mead is the one used by [Karspeck]. See <https://docs.scipy.org/doc/scipy/tutorial/optimize.html> for valid options. Defaults to “Nelder-Mead”.
 - **tol** (float) – Set convergence tolerance for scipy optimize. See <https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.minimize.html#scipy.optimize.minimize>
- Note on new tol kwarg: For N-M, this sets the value to both xatol and fatol Default is 1E-3 Since it affects accuracy of all values including rotation angle 0.001 rad ~ 0.05 deg,
- **default_value** (Any) – Default value(s) to be used for masked values in the returned dataset. Can be specified as a single value used for each parameter, or a list of values.
 - **options** (dict[str, Any] | None) – Options to use in fitting.

Returns

Containing the ellipse parameters, the standard deviation, fit success defined by

- -1: Not computed (masked)
- 0: success
- 2: success but with one parameter reaching upper boundaries
- 3: success with multiple parameters reaching the boundaries (aka both Lx and Ly), can be both at lower or upper bound.
- 9: fail, probably due to running out of maxiter (see scipy.optimize.minimize kwargs “options”)

and the number of iterations.

Return type

Dataset

7.3 Ellipse-based Covariance Estimation

```
class glomar_gridding.ellipse.EllipseCovarianceBuilder(Lx, Ly, theta, stdev, lats, lons, v,
                                                    delta_x_method='Modified_Met_Office',
                                                    max_dist=None, precision=<class
                                                    'numpy.float32'>,
                                                    covariance_method='array',
                                                    batch_size=None)
```

Compute covariance from Ellipse parameters and positions.

v = Matern covariance shape parameter

Lx - an numpy array of horizontal length scales Ly - an numpy array of meridonal length scales or None for circles. theta - an numpy array of rotation angles (RADIANS ONLY) or None for unrotated ellipses.

sdev - standard deviation – right now it just takes a numeric array if you have multiple contribution to sdev (uncertainties derived from different sources), you need to put them into one array

Rules: Valid point: 1) cov_ns and cor_ns are computed out to max_dist; out of range = 0.0 2) Masked points are ignored

Invalid (masked) points: 1) Skipped over

Parameters

- **Lx** (ndarray) – Long length scales of the ellipses (length of semi-major axis) or radius of circles.
- **Ly** (ndarray | None) – Short length scales of the ellipses (length of semi-minor axis). If set to None then the ellipses are isotropic and they are circles.
- **theta** (ndarray | None) – Rotation angles of the ellipses in radians. If set to None then the ellipses are unrotated.
- **stdev** (ndarray) – Standard deviations of the ellipses.
- **lats** (ndarray) – Array containing the latitude values
- **lons** (ndarray) – Array containing the longitude values
- **v** (float) – Matern shape parameter
- **delta_x_method** (Optional[Literal['Met_Office', 'Modified_Met_Office']]) – How are displacements computed between points
- **max_dist** (float | None) – If the Haversine distance between 2 points exceed max_dist, covariance is set to 0. If set to None then an infinite max dist is assumed and covariance between all pairs of positions will be computed.
- **precision** (type) – Floating point precision of the output covariance numpy defaults to np.float32.
- **covariance_method** (Literal['batched', 'low_memory', 'array']) – Set the covariance method used:
 - array (default): faster but uses significantly more memory as more pre-computation is performed. Values are computed in a vectorised method. If the number of points is >= 10_000 then it will use the 'low_memory' method.
 - low_memory: slower iterative process, computes each value individually.
 - batched: combines the above approaches.

If the number of grid-points exceeds 10_000 and “array” method is used, the method will be overwritten to “low_memory”.

- **batch_size** (int | None) – Size of the batch to use for the “batched” method. Must be set if the covariance_method is set to “batched”.

c_ij_anisotropic_array(*i_s*, *j_s*)

Compute the covariances between pairs of ellipses, at displacements.

Each ellipse is defined by values from Lxs, Lys, and thetas, with standard deviation in stdevs.

The displacements between each pair of ellipses are x_is and x_js.

For N ellipses, the number of displacements should be $1/2 * N * (N - 1)$, i.e. the displacement between each pair combination of ellipses. This function will return the upper triangular values of the covariance matrix (excluding the diagonal).

itertools.combinations is used to handle ordering, so the displacements must be ordered in the same way.

Parameters

- **i_s** (*numpy.ndarray*) – The row indices for the covariance matrix.
- **j_s** (*numpy.ndarray*) – The column indices for the covariance matrix.

Returns

c_ij – A vector containing the covariance values between each pair of ellipses. This will return the components of the upper triangle of the covariance matrix as a vector (excluding the diagonal).

Return type
ndarray

References

1. Paciorek and Schevish 2006 [[PaciorekSchervish](#)] Equation 8
2. Karspeck et al. 2012 [[Karspeck](#)] Equation 17

c_ij_isotropic_array(*i_s*, *j_s*, *dist*)

Compute the covariances between pairs of circle, at distance.

Each circle is defined by values from Lxs, with standard deviation in stdevs.

The distances between each pair of circles is the haversine distance.

itertools.combinations is used to handle ordering, so the displacements must be ordered in the same way.

Parameters

- **i_s** (*numpy.ndarray*) – The row indices for the covariance matrix.
- **j_s** (*numpy.ndarray*) – The column indices for the covariance matrix.
- **dist** (*numpy.ndarray*) – The Haversine distance between each coordinate.

Returns

c_ij – A vector containing the covariance values between each pair of circles. This will return the components of the upper triangle of the covariance matrix as a vector (excluding the diagonal).

Return type
ndarray

References

1. Paciorek and Schevish 2006 [[PaciorekSchervish](#)] Equation 8
2. Karspeck et al. 2012 [[Karspeck](#)] Equation 17

calculate_cor()

Calculate correlation matrix from the covariance matrix

Return type

None

calculate_covariance_array()

Calculate the covariance matrix from the ellipse parameters

Return type

None

calculate_covariance_batched()

Compute the covariance matrix from ellipse parameters, using a batched approach. This approach is more memory safe and appropriate for low-memory operations, but is slower than `self.calculate_covariance` which pre-computes values at all upper triangle points. This approach performs pre-computation at all points within the current batch.

Each ellipse is defined by values from `Lxs`, `Lys`, and `thetas`, with standard deviation in `stdevs`.

Requires a `batch_size` parameter.

Return type

None

References

1. Paciorek and Schevish 2006 [[PaciorekSchervish](#)] Equation 8
2. Karspeck et al. 2012 [[Karspeck](#)] Equation 17

calculate_covariance_loop()

Compute the covariance matrix from ellipse parameters, using a loop. This approach is more memory safe and appropriate for low-memory operations, but is significantly slower than `self.calculate_covariance` which uses a lot of pre-computation and a vectorised approach.

Each ellipse is defined by values from `Lxs`, `Lys`, and `thetas`, with standard deviation in `stdevs`.

Return type

None

References

1. Paciorek and Schevish 2006 [[PaciorekSchervish](#)] Equation 8
2. Karspeck et al. 2012 [[Karspeck](#)] Equation 17

uncompress_cov(diag_fill_value=nan, fill_value=nan)

Convert the covariance matrix to full grid size.

Optionally, fill the array with along the diagonal with a `diag_fill_value` and off the diagonal with a `fill_value`, which both default to `np.nan`.

Overwrites the `cov_ns` attribute.

Parameters

- **diag_fill_value** (Any) – Value to assign to diagonal masked values. Defaults to *np.nan*
- **fill_value** (Any) – Value to assign to off-diagonal masked values. Defaults to *np.nan*

Return type

None

ERROR COVARIANCE

Functions for computing correlated and uncorrelated components of the error covariance. These values are determined from standard deviation (sigma) values assigned to groupings within the observational data.

The correlated components will form a matrix that is permutationally equivalent to a block diagonal matrix (i.e. the matrix will be block diagonal if the observational data is sorted by the group).

The uncorrelated components will form a diagonal matrix.

Further a distance-based component can be constructed, where distances between records within the same grid box are evaluated.

The functions in this module are valid for observational data where there could be more than 1 observation in a gridbox.

`glommar_gridding.error_covariance.correlated_components(df, group_col, bias_sig_col=None, bias_sig_map=None)`

Returns measurements covariance matrix updated by adding bias uncertainty to the measurements based on a grouping within the observational data.

The result is equivalent to a block diagonal matrix via permutation. If the input observational data is sorted by the group column then the resulting matrix is block diagonal, where the blocks are the size of each grouping. The values in each block are the square of the sigma value associated with the grouping.

Note that in most cases the output is not a block-diagonal, as the input is not usually sorted by the group column. In most processing cases, the input dataframe will be sorted by the gridbox index.

The values can either be pre-defined in the observational dataframe, and can be indicated by the “bias_val_col” argument. Alternatively, a mapping can be passed, the values will be then assigned by this mapping of group to sigma.

Parameters

- **df** (DataFrame) – Observational DataFrame including group information and bias uncertainty values for each grouping. It is assumed that a single bias uncertainty value applies to the whole group, and is applied as cross terms in the covariance matrix (plus to the diagonal).
- **group_col** (str) – Name of the column that can be used to partition the observational DataFrame.
- **bias_sig_col** (str | None) – Name of the column containing bias uncertainty values for each of the groups identified by ‘group_col’. It is assumed that a single bias uncertainty value applies to the whole group, and is applied as cross terms in the covariance matrix (plus to the diagonal).
- **bias_sig_map** (dict[str, float] | None) – Mapping between values in the group_col and bias uncertainty values, if bias_val_col is not in the DataFrame.

Return type
ndarray

`glomar_gridding.error_covariance.dist_weight(df, dist_fn, grid_idx='grid_idx', **dist_kwargs)`

Compute the distance and weight matrices over gridboxes for an input Frame.

This function acts as a wrapper for a distance function, allowing for computation of the distances between positions in the same gridbox using any distance metric.

The weightings from this function are for the gridbox mean of the observations within a gridbox.

Parameters

- **df** (DataFrame) – The observation DataFrame, containing the columns required for computation of the distance matrix. Contains the “grid_idx” column which indicates the gridbox for a given observation. The index of the DataFrame should match the index ordering for the output distance matrix/weights.
- **dist_fn** (Callable) – The function used to compute a distance matrix for all points in a given grid-cell. Takes as input a polars.DataFrame as first argument. Any other arguments should be constant over all gridboxes, or can be a look-up table that can use values in the DataFrame to specify values specific to a gridbox. The function should return a numpy matrix, which is the distance matrix for the gridbox only. This wrapper function will correctly apply this matrix to the larger distance matrix using the index from the DataFrame.

If `dist_fn` is None, then no distances are computed and None is returned for the `dist` value.

- **grid_idx** (str) – Name of the column containing the grid index values
- ****dist_kwargs** – Arguments to be passed to `dist_fn`. In general these should be constant across all gridboxes. It is possible to pass a look-up table that contains pre-computed values that are gridbox specific, if the keys can be matched to a column in `df`.

Return type

tuple[ndarray, ndarray]

Returns

- **dist** (*numpy.matrix*) – The distance matrix, which contains the same number of rows and columns as rows in the input DataFrame `df`. The values in the matrix are 0 if the indices of the row/column are for observations from different gridboxes, and non-zero if the row/column indices fall within the same gridbox. Consequently, with appropriate re-arrangement of rows and columns this matrix can be transformed into a block-diagonal matrix. If the DataFrame input is pre-sorted by the gridbox column, then the result is a block-diagonal matrix.

If `dist_fn` is None, then this value will be None.

- **weights** (*numpy.matrix*) – A matrix of weights. This has dimensions $n \times p$ where n is the number of unique gridboxes and p is the number of observations (the number of rows in `df`). The values are 0 if the row and column do not correspond to the same gridbox and equal to the inverse of the number of observations in a gridbox if the row and column indices fall within the same gridbox. The rows of weights are in a sorted order of the gridbox. Should this be incorrect, one should re-arrange the rows after calling this function.

`glomar_gridding.error_covariance.get_weights(df, grid_idx='grid_idx')`

Get just the weight matrices over gridboxes for an input Frame.

The weightings from this function are for the gridbox mean of the observations within a gridbox.

Parameters

- **df** (DataFrame) – The observation DataFrame, containing the columns required for computation of the distance matrix. Contains the “grid_idx” column which indicates the gridbox for a given observation. The index of the DataFrame should match the index ordering for the output weights.
- **grid_idx** (str) – Name of the column containing the gridbox index from the output grid.

Returns

weights – A matrix of weights. This has dimensions $n \times p$ where n is the number of unique gridboxes and p is the number of observations (the number of rows in df). The values are 0 if the row and column do not correspond to the same gridbox and equal to the inverse of the number of observations in a gridbox if the row and column indices fall within the same gridbox. The rows of weights are in a sorted order of the gridbox. Should this be incorrect, one should re-arrange the rows after calling this function.

Return type

ndarray

```
glomar_gridding.error_covariance.grid_box_weighted_sum(grid_box_frame, group_col,
                                                       error_group_correlated,
                                                       error_uncorrelated, lat_col='lat',
                                                       lon_col='lon', date_col='datetime',
                                                       sill='sill', space_range='space_range',
                                                       time_range='time_range',
                                                       bad_groups=None)
```

Get the weighted mean weights for a grid-box. The weights are computed using uncertainty from a simple spatio-temporal exponential variogram model, the correlated error covariance components (group-wise) and the uncorrelated error covariance components (group-wise).

The correlated components of the error covariance are 0 between pairs of records with `_different_` groups, and the value from the ‘error_group_correlated’ column squared when the groups match.

This computes the weights for a grid-box using the local inverse error covariance.

A factor of 1/4 can be applied to the cross-record pairs in the computation of correlated components for a subset of groupings considered “bad” - for example generic ids.

Parameters

- **grid_box_frame** (DataFrame) – DataFrame containing the observations for all records within a grid-box.
- **group_col** (str) – Name of the column by which records are grouped for the purpose of correlated and uncorrelated components of the error covariance structure.
- **error_group_correlated** (str) – Name of the column containing correlated error values by group.
- **error_uncorrelated** (str) – Name of the column containing uncorrelated error values by group.
- **lat_col** (str) – Name of the latitude column.
- **lon_col** (str) – Name of the longitude column.
- **date_col** (str) – Name of the datetime column.
- **sill** (str) – The name of the column containing sill values of the simplified varigram model.
- **space_range** (str) – The name of the column containing space range values for the simplified varigram model. [km]

- **time_range** (str) – The name of the column containing time range values for the simplified variogram model. [days]
- **bad_groups** (str | list[str] | None) – Values in the groups that required lower priority in the weightings. For example, this could be records with invalid, or generic, ids.

Return type

tuple[ndarray, ndarray, ndarray, ndarray]

Returns

- **weights** (*np.ndarray*) – A vector containing the weights for each record in the grid-box. Retaining the order within the grid-cell.
- **vario** (*np.ndarray*) – A matrix containing the output of the spatio-temporal variogram model for the grid-box records.
- **correlated** (*np.ndarray*) – Correlated error covariance matrix for the grid-box. Contains correlated error values when the group between each pair of records matches, zero otherwise.
- **uncorrelated** (*np.ndarray*) – Uncorrelated error matrix for the grid-box. Diagonal matrix.

```
glomar_gridding.error_covariance.uncorrelated_components(df, group_col='data_type',  
                                                         obs_sig_col=None, obs_sig_map=None)
```

Calculates the covariance matrix of the measurements (observations). This is the uncorrelated component of the covariance.

The result is a diagonal matrix. The diagonal is formed by the square of the sigma values associated with the values in the grouping.

The values can either be pre-defined in the observational dataframe, and can be indicated by the “bias_val_col” argument. Alternatively, a mapping can be passed, the values will be then assigned by this mapping of group to sigma.

Parameters

- **df** (DataFrame) – The observational DataFrame containing values to group by.
- **group_col** (str) – Name of the group column to use to set observational sigma values.
- **obs_sig_col** (str | None) – Name of the column containing observational sigma values. If set and present in the DataFrame, then this column is used as the diagonal of the returned covariance matrix.
- **obs_sig_map** (dict[str, float] | None) – Mapping between group and observational sigma values used to define the diagonal of the returned covariance matrix.

Return type

ndarray

Returns

- *A diagonal matrix representing the uncorrelated components of the error*
- *covariance matrix.*

```
glomar_gridding.error_covariance.weighted_sum(df, grid_idx, group_col, error_group_correlated,  
                                              error_uncorrelated, lat_col='lat', lon_col='lon',  
                                              date_col='datetime', sill='sill',  
                                              space_range='space_range', time_range='time_range',  
                                              bad_groups=None)
```

Get the weights matrix for the weighted sum approach, accounting for correlated and uncorrelated error components. The weights are computed using uncertainty from a simple spatio-temporal exponential variogram

model, the correlated error covariance components (group-wise) and the uncorrelated error covariance components (group-wise).

The correlated components of the error covariance are 0 between pairs of records with `_different_` groups, and the value from the `'error_group_correlated'` column squared when the groups match.

This computes the weights for each grid-box using the local inverse error covariance.

A factor of 1/4 can be applied to the cross-record pairs in the computation of correlated components for a subset of groupings considered “bad” - for example generic ids.

Parameters

- **df** (DataFrame) – The observation DataFrame. Contains the “grid_idx” column which indicates the gridbox for a given observation. The index of the DataFrame should match the index ordering for the output weights.
- **grid_idx** (str) – Name of the column containing the gridbox index from the output grid.
- **group_col** (str) – Name of the column by which records are grouped for the purpose of correlated and uncorrelated components of the error covariance structure.
- **error_group_correlated** (str) – Name of the column containing correlated error values by group.
- **error_uncorrelated** (str) – Name of the column containing uncorrelated error values by group.
- **lat_col** (str) – Name of the latitude column.
- **lon_col** (str) – Name of the longitude column.
- **date_col** (str) – Name of the datetime column.
- **bad_groups** (str | list[str] | None) – Values in the groups that required lower priority in the weightings. For example, this could be records with invalid, or generic, ids.
- **sill** (str) – The name of the column containing sill values of the simplified variogram model.
- **space_range** (str) – The name of the column containing space range values for the simplified variogram model. [km]
- **time_range** (str) – The name of the column containing time range values for the simplified variogram model. [days]

Returns

weights – A matrix containing the weights for each record in the data. Retaining the order within the frame.

Return type

ndarray

KRIGING

The `glommar_gridding.kriging` module contains classes and functions for interpolation via Kriging. Two methods of Kriging are supported by `glommar_gridding`:

- Simple Kriging
- Ordinary Kriging

For each Kriging method there is a class and a function. The recommended approach is to use the classes, the functions will be deprecated in a future version of `glommar_gridding`. The classes require the full grid spatial covariance structure, the observation values, and the grid index of each observation, and an optional error covariance matrix as inputs. Each grid index should be a single index value, and represents the flattened index, and connects directly to the corresponding index of the covariance matrices. If an error covariance matrix is provided, the covariance matrix will be automatically subset to the grid index values, if the resulting matrix contains *nan* or 0 values on the diagonal, then the observation values and indices are filtered to exclude these points, and the error covariance matrix is subset again. Just initialising the class does not solve the system, this requires the `solve` method to be called.

9.1 Preparation

`glommar_gridding` provides functionality for preparing your data for the interpolation. The `grid` module has functionality for defining the output grid (`glommar_gridding.grid.grid_from_resolution()`) which allows the user to create a coordinate system for the output, that can easily be mapped to a covariance matrix. The grid object is an `xarray.DataArray` object, with a coordinate system. Once the grid is defined, the observations can be mapped to the grid. This creates a 1-dimensional index value that should match to the covariance matrices used in the interpolation.

```
glommar_gridding.grid.map_to_grid(obs, grid, obs_coords=['lat', 'lon'], grid_coords=['latitude', 'longitude'],
                                  sort=True, bounds=None, add_grid_pts=True, grid_prefix='grid_')
```

Align an observation dataframe to a grid defined by an `xarray DataArray`.

Maps observations to the nearest grid-point, and sorts the data by the 1d index of the `DataArray` in a row-major format.

The grid defined by the latitude and longitude coordinates of the input `DataArray` is then used as the output grid of the Gridding process.

Parameters

- **obs** (`DataFrame`) – The observational `DataFrame` containing positional data with latitude, longitude values within the `obs_latname` and `obs_lonname` columns respectively. Observations are mapped to the nearest grid-point in the grid.
- **grid** (`DataArray`) – Contains the grid coordinates to map observations to.
- **obs_coords** (`list[str]`) – Names of the column containing positional values in the input observational `DataFrame`.

- **grid_coords** (list[str]) – Names of the coordinates in the input grid DataArray used to define the grid.
- **sort** (bool) – Sort the observational DataFrame by the grid index
- **bounds** (list[tuple[float, float]] | None) – Optionally filter the grid and DataFrame to fall within spatial bounds. This list must have the same size and ordering as *obs_coords* and *grid_coords* arguments.
- **add_grid_pts** (bool) – Add the grid positional information to the observational DataFrame.
- **grid_prefix** (str) – Prefix to use for the new grid columns in the observational DataFrame.

Returns

obs – Containing additional *grid_**, and *grid_idx* values indicating the positions and grid index of the observation respectively. The DataFrame is also sorted (ascendingly) by the *grid_idx* columns for consistency with the gridding functions.

Return type

DataFrame

Examples

```
>>> obs = pl.read_csv("/path/to/obs.csv")
>>> grid = grid_from_resolution(
    resolution=5,
    bounds=[(-87.5, 90), (-177.5, 180)], # Lower bound is centre
    coord_names=["lat", "lon"]
)
>>> obs = map_to_grid(obs, grid, grid_coords=["lat", "lon"])
```

For Kriging, the interpolation requires at most a single observation value in each grid box. If the data contains multiple values in a single grid cell then these need to be combined.

`glomar_gridding.kriging.prep_obs_for_kriging(unmask_idx, unique_obs_idx, weights, obs,`
`remove_obs_mean=0, obs_bias=None, error_cov=None)`

Prep masked observations for Kriging. Combines observations in the same grid box to a single averaged observation using a weighted average.

Parameters

- **unmask_idx** (ndarray) – Indices of all un-masked points for chosen date.
- **unique_obs_idx** (ndarray) – Unique indices of all measurement points for a chosen date, representative of the indices of gridboxes, which have => 1 measurement.
- **weights** (ndarray) – Weight matrix (inverse of counts of observations).
- **obs** (ndarray) – All point observations/measurements for the chosen date.
- **remove_obs_mean** (int) – Should the mean or median from obs be removed and added back onto obs?
 - 0 = No (default action)
 - 1 = the mean is removed
 - 2 = the median is removed
 - 3 = the spatial mean is removed

Note that the mean will need to be reapplied to the Kriging result.

- **obs_bias** (ndarray | None) – Bias of all measurement points for a chosen date (corresponds to x_obs).

Return type

tuple[ndarray, ndarray]

Returns

- **obs_idx** (numpy.ndarray[int]) – Subset of grid-box indices containing observations that are unmasked.
- **grid_obs** (numpy.ndarray[float]) – Unmasked and combined observations

9.2 Simple Kriging

class glomar_gridding.kriging.**SimpleKriging**(covariance, idx, obs, error_cov=None)

Class for SimpleKriging.

The equation for simple Kriging is:

$$(C_{obs} + E)^{-1} \times C_{cross} \times y + \mu$$

Where μ is a constant known mean, typically this is 0.

Parameters

- **covariance** (ndarray) – The spatial covariance matrix. This can be a pre-computed matrix loaded into the environment, or computed from a Variogram class or using Ellipse methods.
- **idx** (ndarray) – The 1d indices of observation grid points. These values should be between 0 and (N * M) - 1 where N, M are the number of longitudes and latitudes respectively. Note that these values should also be computed using “C” ordering in numpy reshaping. They can be computed from a grid using glomar_gridding.grid.map_to_grid. Each value should only appear once. Points that contain more than 1 observation should be averaged
- **obs** (ndarray) – The observation values. If there are multiple observations in any grid box then these values need to be averaged into one value per grid box.
- **error_cov** (ndarray | None) – Optionally add error covariance values to the covariance between observation grid points.

constraint_mask()

Compute the observational constraint mask (A14 in [Morice_2021]) to determine if a grid point should be masked/weights modified by how far it is to its near observed point

Note: typo in Section A4 in [Morice_2021] (confirmed by authors).

Equation to use is A14 is incorrect. Easily noticeable because dimensionally incorrect is wrong, but the correct answer is easy to figure out.

Correct Equation (extra matrix inverse for $C_{obs} + E$):

$$1 - \frac{\text{diag}(C - C_{cross}^T \times (C_{obs} + E)^{-1} \times C_{cross})}{\text{diag}(C)} < \alpha$$

This can be re-written as:

$$\frac{\text{diag}(C_{cross}^T \times (C_{obs} + E)^{-1} \times C_{cross})}{\text{diag}(C)} < \alpha$$

α is chosen to be 0.25 in the UKMO paper

Written by S. Chan, modified by J. Siddons.

This requires that the *kriging_weights* attribute is set.

Returns

constraint_mask – Constraint mask values, the left-hand-side of equation A14 from Morice et al. (2021). This is a vector of length *k_obs.size[0]*.

Return type

ndarray

References

[Morice_2021]: <https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2019JD032361>

get_kriging_weights()

Compute the Kriging weights from the flattened grid indices where there is an observation. Optionally add an error covariance to the covariance between observation grid points.

The Kriging weights are calculated as:

$$(C_{obs} + E)^{-1} \times C_{cross}$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, and C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points).

Sets the *kriging_weights* attribute.

Return type

None

get_uncertainty(full_covariance=False)

Compute the kriging uncertainty. This requires the attribute *kriging_weights* to be computed.

Parameters

full_covariance (bool) – Return the full Kriging covariance if True. Defaults to False.

Returns

uncert – The Kriging uncertainty or the full Kriging covariance.

Return type

ndarray

kriging_weights_from_inverse(inv)

Compute the Kriging weights from the flattened grid indices where there is an observation, using a pre-computed inverse of the covariance between grid-points with observations.

The Kriging weights are calculated as:

$$(C_{obs} + E)^{-1} \times C_{cross}$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, and C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points).

Sets the *kriging_weights* attribute.

Parameters

inv (ndarray) – The pre-computed inverse of the covariance between grid-points with observations. $(C_{obs} + E)^{-1}$

Return type

None

solve(*mean=0.0*)

Solves the simple Kriging problem. Computes the Kriging weights if the *kriging_weights* attribute is not already set. The solution to Kriging is:

$$(C_{obs} + E)^{-1} \times C_{cross} \times y$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points), and y are the observation values.

Parameters

mean (ndarray | float) – Constant, known, mean value of the system. Defaults to 0.0.

Returns

The solution to the simple Kriging problem (as a Vector, this may need to be re-shaped appropriately as a post-processing step).

Return type

ndarray

Examples

```
>>> SK = SimpleKriging(interp_covariance)
>>> SK.solve(obs, idx, error_covariance)
```

`glomar_gridding.kriging.kriging_simple(obs_obs_cov, obs_grid_cov, grid_obs, interp_cov, mean=0.0)`

Perform Simple Kriging assuming a constant known mean.

This function is deprecated in favour of SimpleKriging class. It will be removed in version 1.0.0.

Parameters

- **obs_obs_cov** (ndarray) – Covariance between all measured grid points plus the covariance due to measurements (i.e. measurement noise, bias noise, and sampling noise). Can include error covariance terms.
- **obs_grid_cov** (ndarray) – Covariance between the all (predicted) grid points and measured points. Does not contain error covariance.
- **grid_obs** (ndarray) – Gridded measurements (all measurement points averaged onto the output gridboxes).
- **interp_cov** (ndarray) – interpolation covariance of all output grid points (each point in time and all points against each other).
- **mean** (float | ndarray) – The constant mean of the output field.

Return type

tuple[ndarray, ndarray]

Returns

- **z_obs** (*np.ndarray[float]*) – Full set of values for the whole domain derived from the observation points using simple kriging.
- **dz** (*np.ndarray[float]*) – Uncertainty associated with the simple kriging.

9.3 Ordinary Kriging

class glomar_gridding.kriging.**OrdinaryKriging**(*covariance*, *idx*, *obs*, *error_cov=None*)

Class for OrdinaryKriging.

The equation for ordinary Kriging is:

$$(C_{obs} + E)^{-1} \times C_{cross} \times y$$

with a constant but unknown mean.

In this case, the C_{obs} , C_{cross} and y values are extended with a Lagrange multiplier term, ensuring that the Kriging weights are constrained to sum to 1.

The matrix C_{obs} is extended by one row and one column, each containing the value 1, except at the diagonal point, which is 0. The C_{cross} matrix is extended by an extra row containing values of 1. Finally, the grid observations y is extended by a single value of 0 at the end of the vector.

Parameters

- **covariance** (ndarray) – The spatial covariance matrix. This can be a pre-computed matrix loaded into the environment, or computed from a Variogram class or using Ellipse methods.
- **idx** (ndarray) – The 1d indices of observation grid points. These values should be between 0 and (N * M) - 1 where N, M are the number of longitudes and latitudes respectively. Note that these values should also be computed using “C” ordering in numpy reshaping. They can be computed from a grid using `glomar_gridding.grid.map_to_grid`. Each value should only appear once. Points that contain more than 1 observation should be averaged. Used to compute the Kriging weights.
- **obs** (ndarray) – The observation values. If there are multiple observations in any grid box then these values need to be averaged into one value per grid box.
- **error_cov** (ndarray | None) – Optionally add error covariance values to the covariance between observation grid points.

constraint_mask(*simple_kriging_weights=None*)

Compute the observational constraint mask (A14 in [Morice_2021]) to determine if a grid point should be masked/weights modified by how far it is to its near observed point

Note: typo in Section A4 in [Morice_2021] (confirmed by authors).

Equation to use is A14 is incorrect. Easily noticeable because dimensionally incorrect is wrong, but the correct answer is easy to figure out.

Correct Equation (extra matrix inverse for $C_{obs} + E$):

$$1 - \frac{\text{diag}(C - C_{cross}^T \times (C_{obs} + E)^{-1} \times C_{cross})}{\text{diag}(C)} < \alpha$$

This can be re-written as:

$$\frac{\text{diag}(C_{cross}^T \times (C_{obs} + E)^{-1} \times C_{cross})}{\text{diag}(C)} < \alpha$$

α is chosen to be 0.25 in the UKMO paper

Written by S. Chan, modified by J. Siddons.

This requires the Kriging weights from simple Kriging. If these are not provided as an input, then they are calculated.

Parameters

simple_kriging_weights (ndarray | None) – The Kriging weights for the equivalent simple Kriging system. error covariance.

Returns

constraint_mask – Constraint mask values, the left-hand-side of equation A14 from Morice et al. (2021). This is a vector of length $k_{obs.size}[0]$.

Return type

ndarray

References

[Morice_2021]: <https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2019JD032361>

extended_inverse(simple_inv)

Compute the inverse of a covariance matrix $S = C_{obs} + E$, and use that to compute the inverse of the extended version of the covariance matrix with Lagrange multipliers, used by Ordinary Kriging.

This is useful when one needs to perform BOTH simple and ordinary Kriging, or when one wishes to compute the constraint mask for ordinary Kriging which requires the Kriging weights for the equivalent simple Kriging problem.

The extended form of S is given by:

$$\begin{pmatrix} & & & 1 \\ & S & & \vdots \\ & & & 1 \\ 1 & \dots & 1 & 0 \end{pmatrix}$$

This approach follows Guttman 1946 10.1214/aoms/1177730946

Parameters

simple_inv (ndarray) – Inverse of the covariance between observation grid-points

Returns

Inverse of the extended covariance matrix between observation grid-points including the Lagrange multiplier factors.

Return type

ndarray

get_kriging_weights()

Compute the Kriging weights from the flattened grid indices where there is an observation. Optionally add an error covariance to the covariance between observation grid points.

The Kriging weights are calculated as:

$$(C_{obs} + E)^{-1} \times C_{cross}$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, and C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points).

In this case, the C_{obs} , C_{cross} and are extended with a Lagrange multiplier term, ensuring that the Kriging weights are constrained to sum to 1.

The matrix C_{obs} is extended by one row and one column, each containing the value 1, except at the diagonal point, which is 0. The C_{cross} matrix is extended by an extra row containing values of 1.

Sets the *kriging_weights* attribute.

Return type

None

get_uncertainty(*full_covariance=False*)

Compute the kriging uncertainty. This requires the attribute *kriging_weights* to be computed.

Parameters

full_covariance (bool) – Return the full Kriging covariance if True. Defaults to False.

Returns

uncert – The Kriging uncertainty or the full Kriging covariance.

Return type

ndarray

kriging_weights_from_inverse(*inv*)

Compute the Kriging weights from the flattened grid indices where there is an observation, using a pre-computed inverse of the covariance between grid-points with observations.

The Kriging weights are calculated as:

$$(C_{obs} + E)^{-1} \times C_{cross}$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, and C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points).

In this case, the inverse matrix must be computed from the covariance between observation grid-points with the Lagrange multiplier applied.

This method is appropriate if one wants to compute the constraint mask which requires simple Kriging weights, which can be computed from the unextended covariance inverse. The extended inverse can then be calculated from that inverse.

Sets the *kriging_weights* attribute.

Parameters

inv (ndarray) – The pre-computed inverse of the covariance between grid-points with observations. $(C_{obs} + E)^{-1}$

Return type

None

solve()

Solves the ordinary Kriging problem. Computes the Kriging weights if the *kriging_weights* attribute is not already set. The solution to Kriging is:

$$(C_{obs} + E)^{-1} \times C_{cross} \times y$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points), and y are the observation values.

In this case, the C_{obs} , C_{cross} and are extended with a Lagrange multiplier term, ensuring that the Kriging weights are constrained to sum to 1.

The matrix C_{obs} is extended by one row and one column, each containing the value 1, except at the diagonal point, which is 0. The C_{cross} matrix is extended by an extra row containing values of 1.

Returns

The solution to the ordinary Kriging problem (as a Vector, this may need to be re-shaped appropriately as a post-processing step).

Return type
ndarray

Examples

```
>>> OK = OrdinaryKriging(interp_covariance)
>>> OK.solve(obs, idx, error_covariance)
```

`glomar_gridding.kriging.kriging_ordinary(obs_obs_cov, obs_grid_cov, grid_obs, interp_cov)`

Perform Ordinary Kriging with unknown but constant mean.

This function is deprecated in favour of OrdinaryKriging class. It will be removed in version 1.0.0.

Parameters

- **obs_obs_cov** (ndarray) – Covariance between all measured grid points plus the covariance due to measurements (i.e. measurement noise, bias noise, and sampling noise). Can include error covariance terms, if these are being used.
- **obs_grid_cov** (ndarray) – Covariance between the all (predicted) grid points and measured points. Does not contain error covariance.
- **grid_obs** (ndarray) – Gridded measurements (all measurement points averaged onto the output gridboxes).
- **interp_cov** (ndarray) – Interpolation covariance of all output grid points (each point in time and all points against each other).

Return type
tuple[ndarray, ndarray]

Returns

- **z_obs** (*np.ndarray[float]*) – Full set of values for the whole domain derived from the observation points using ordinary kriging.
- **dz** (*np.ndarray[float]*) – Uncertainty associated with the ordinary kriging.

9.4 Perturbed Gridded Fields

An additional two-stage combined Kriging class is provided in the *stochastic* module. In this case, the *solve* method has an additional optional *simulated_state* argument, if this is set, then the value is used as the simulated system state used as the base of the perturbed field (from which observations are simulated), otherwise the value is computed. This allows for pre-computation of a sequence of simulated states as part of ensemble processing, for example.

class `glomar_gridding.stochastic.StochasticKriging(covariance, idx, obs, error_cov)`

Class for the combined two-stage Kriging approach following [Morice_2021] The first stage is to produce a gridded field from the observations using Ordinary Kriging. The second stage is to apply a perturbation.

The perturbation is constructed by first generating a simulated state from the covariance matrix. A set of simulated observations is drawn from the error covariance matrix. A simulated gridded field is then computed using Simple Kriging with the simulated observations as input. Finally, the perturbation is then the difference between the simulated gridded field and the simulated state. This perturbation is added to the gridded field from the first stage.

The equation for ordinary Kriging is:

$$(C_{obs} + E)^{-1} \times C_{cross} \times y$$

with a constant but unknown mean.

In this case, the C_{obs} , C_{cross} and y values are extended with a Lagrange multiplier term for the first stage, ensuring that the Kriging weights are constrained to sum to 1.

Additionally, the matrix C_{obs} is extended by one row and one column, each containing the value 1, except at the diagonal point, which is 0 for the first Ordinary Kriging stage. The C_{cross} matrix is extended by an extra row containing values of 1. Finally, the grid observations y is extended by a single value of 0 at the end of the vector.

Parameters

- **covariance** (ndarray) – The spatial covariance matrix. This can be a pre-computed matrix loaded into the environment, or computed from a Variogram class or using Ellipse methods.
- **idx** (ndarray) – The 1d indices of observation grid points. These values should be between 0 and $(N * M) - 1$ where N, M are the number of longitudes and latitudes respectively. Note that these values should also be computed using “C” ordering in numpy reshaping. They can be computed from a grid using `glomar_gridding.grid.map_to_grid`. Each value should only appear once. Points that contain more than 1 observation should be averaged. Used to compute the Kriging weights.
- **obs** (ndarray) – The observation values. If there are multiple observations in any grid box then these values need to be averaged into one value per grid box.
- **error_cov** (ndarray) – Error covariance values to the covariance between observation grid points.

constraint_mask()

Compute the observational constraint mask (A14 in [Morice_2021]) to determine if a grid point should be masked/weights modified by how far it is to its near observed point

Note: typo in Section A4 in [Morice_2021] (confired by authors).

Equation to use is A14 is incorrect. Easily noticeable because dimensionally incorrect is wrong, but the correct answer is easy to figure out.

Correct Equation (extra matrix inverse for $C_{obs} + E$):

$$1 - \frac{\text{diag}(C - C_{cross}^T \times (C_{obs} + E)^{-1} \times C_{cross})}{\text{diag}(C)} < \alpha$$

This can be re-written as:

$$\frac{\text{diag}(C_{cross}^T \times (C_{obs} + E)^{-1} \times C_{cross})}{\text{diag}(C)} < \alpha$$

α is chosen to be 0.25 in the UKMO paper

Written by S. Chan, modified by J. Siddons.

This requires the Kriging weights from simple Kriging, set as the *simple_kriging_weights* attribute.

Returns

constraint_mask – Constraint mask values, the left-hand-side of equation A14 from [Morice_2021]. This is a vector of length $k_{obs.size}[0]$.

Return type

ndarray

References

[Morice_2021]: <https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2019JD032361>

get_kriging_weights()

Compute the Kriging weights from the flattened grid indices where there is an observation. Optionally add an error covariance to the covariance between observation grid points.

The Kriging weights are calculated as:

$$(C_{obs} + E)^{-1} \times C_{cross}$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, and C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points).

In this case, the C_{obs} , C_{cross} and are extended with a Lagrange multiplier term, ensuring that the Kriging weights are constrained to sum to 1, for the first stage of the StochasticKriging process.

The matrix C_{obs} is extended by one row and one column, each containing the value 1, except at the diagonal point, which is 0. The C_{cross} matrix is extended by an extra row containing values of 1.

Sets the *kriging_weights* and *simple_kriging_weights* attributes.

Return type

None

get_uncertainty(full_covariance=False)

Compute the kriging uncertainty. This requires the attribute *kriging_weights* to be computed.

Parameters

full_covariance (bool) – Return the full Kriging covariance if True. Defaults to False.

Returns

uncert – The Kriging uncertainty or the full Kriging covariance.

Return type

ndarray

kriging_weights_from_inverse(inv)

Compute the Kriging weights from the flattened grid indices where there is an observation, using a pre-computed inverse of the covariance between grid-points with observations.

The Kriging weights are calculated as:

$$(C_{obs} + E)^{-1} \times C_{cross}$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, and C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points).

This method is appropriate if one wants to compute the constraint mask which requires simple Kriging weights, which can be computed from the unextended covariance inverse. The extended inverse can then be calculated from that inverse.

The inverse matrix is used to compute the inverse of the extended covariance matrix used for the first Ordinary Kriging stage.

Sets the *kriging_weights* and *simple_kriging_weights* attributes.

Parameters

inv (ndarray) – The pre-computed inverse of the covariance between grid-points with observations. $(C_{obs} + E)^{-1}$

Return type

None

set_simple_kriging_weights(*simple_kriging_weights*)

Set Simple Kriging Weights. For use in the second Simple Kriging stage for computing the simulated gridded field.

Sets the *simple_kriging_weights* attribute.

Parameters

simple_kriging_weights (ndarray) – The pre-computed *simple_kriging_weights* to use.

Return type

None

solve(*simulated_state=None*)

Solves the combined Stochastic Kriging problem. Computes the Kriging weights if the *kriging_weights* attribute is not already set.

Stochastic Kriging is a combined Ordinary and Simple Kriging approach. First a gridded field is generated for the observations using Ordinary Kriging. Secondly, a simulated gridded field is generated using Simple Kriging from a simulated state and simulated observations. The simulated state can be pre-computed or computed by this method. The simulated observations are drawn from the error covariance and located at the simulated state. The difference between the simulated gridded field and the simulated state is added to the gridded field.

The solution to Kriging is:

$$(C_{obs} + E)^{-1} \times C_{cross} \times y$$

Where C_{obs} is the spatial covariance between grid-points with observations, E is the error covariance between grid-points with observations, C_{cross} is the covariance between grid-points with observations and all grid-points (including observation grid-points), and y are the observation values.

In this case, the C_{obs} , C_{cross} and are extended with a Lagrange multiplier term, ensuring that the Kriging weights are constrained to sum to 1. For the Ordinary Kriging component.

The matrix C_{obs} is extended by one row and one column, each containing the value 1, except at the diagonal point, which is 0. The C_{cross} matrix is extended by an extra row containing values of 1. For the Ordinary Kriging component.

This additionally sets the following attributes:

- *gridded_field* - The unperturbed gridded field
- *simulated_grid* - The simulated gridded field
- *epsilon* - The perturbation

Parameters

simulated_state (ndarray | None) – Flattened simulated state, used as the location basis for the simulated observation draw. If this is not provided it will be calculated. Often it is better to pre-compute a series of states in advance, since this can be a time consuming step (drawing a single state takes approximately the same time as drawing 200).

Returns

The solution to the stochastic Kriging problem (as a Vector, this may need to be re-shaped appropriately as a post-processing step).

Return type

ndarray

Examples

```
>>> SK = StochasticKriging(interp_covariance)
>>> SK.solve(obs, idx, error_covariance)
```

```
glomar_gridding.stochastic.draw_from_cov(loc, cov, ndraws=1, sym_atol=1e-05, eigen_rtol=1e-06,
                                         eigen_thresh=1e-08)
```

Do a random multivariate normal draw using numpy, with a fallback to scipy.stats.multivariate_normal.rvs if that fails.

This function has similar API as GP_draw with less kwargs.

Warning/possible future scipy version may change this: It seems if one uses stats.Covariance, you have to have add [0] from rvs function. The above behavior applies to scipy v1.14.0

Parameters

- **loc** (ndarray) – the location for the normal dry
- **cov** (ndarray) – not a xarray/iris cube! Some of our covariances are saved in numpy format and not netCDF files
- **n_draws** (int) – number of simulations, this is usually set to 1 except during
- **testing** (unit)
- **sym_atol** (float) – absolute tolerance to check symmetry of cov
- **eigen_rtol** (float) – relative tolerance to negative eigenvalues
- **eigen_thresh** (float) – forced minimum value of eigenvalues if negative values are detected

Returns

draw – The draw(s) from the multivariate random normal distribution defined by the loc and cov parameters. If the cov parameter is not positive-definite then a new covariance will be determined by adjusting the eigen decomposition such that the modified covariance should be positive-definite.

Return type

ndarray

Examples

```
>>> A = np.random.rand(5, 5)
>>> cov = np.dot(A, A.T)
>>> draw_from_cov(np.zeros(5), cov, ndraws=5)
array([[ -0.35972806,  -0.51289612,   0.85307028,  -0.11580307,   0.6677707 ],
       [ -1.38214628,  -1.29331638,  -0.4879436 ,  -1.42310831,  -0.19369562],
       [ -1.04502143,  -1.97686163,  -2.058605  ,  -1.97202206,  -2.90116796],
       [ -1.97981119,  -2.72330373,   0.0088662 ,  -2.53521893,  -0.03670664],
       [  0.49948228,   0.54695988,   0.33864294,   0.53730282,   0.14743019]])
```

9.5 Outputs

The outputs to the solvers, `glomar_gridding.SimpleKriging.solve()` for example will be vectors, they should be re-shaped to the grid.

Outputs can also be re-mapped to the grid

`glomar_gridding.grid.assign_to_grid(values, grid_idx, grid, fill_value=nan)`

Assign a vector of values to a grid, using a list of grid index values.

Parameters

- **values** (ndarray) – The values to map onto the output grid.
- **grid_idx** (ndarray) – The 1d index of the grid (assuming “C” style ravelling) for each value.
- **grid** (DataArray) – The grid used to define the output grid.
- **fill_value** (Any) – The value to fill unassigned grid boxes. Must be a valid value of the input *values* data type.

Returns

out_grid – A new grid containing the values mapped onto the grid.

Return type

DataArray

SPLINE INTERPOLATION

The *spline* module brings an alternative approach to interpolation using thin-plate-spline methodology to perform interpolation. We provide a standard thin-plate-spline interpolator using the traditional thin-plate-spline kernel:

$$r^2 \times \log(r)$$

where r is the distance, as well as a spline method that applies to spherical geometry following the approach of [Wahba_Sphere].

The *Spline* classes all have the following methods:

- *fit* : To fit the spline to the training data (used to initialise the class instance). A smoothing parameter can be supplied as the *lam* argument, or it can be estimated using GCV minimisation.
- *predict* : To predict at new positions (or at the training positions) using the result from fitting the spline. The standard error can also be returned.
- *standard_error* : To compute the standard error of the prediction (following [Nychka]).

On using the *fit* method with *set_results = True* a *SplineResult* class is set to the *Spline* as the *result* attribute. This contains details about the fit, and optionally statistics about the fit.

10.1 Example Workflow

```
import numpy as np
import polars as pl

from glomar_gridding.spline import SphericalThinPlateSpline

# Load observations (latitude, longitude, tos)
obs = pl.read_csv("/path/to/obs.csv", ...)

# Get the training data from the observations
X_train = obs.select(["latitude", "longitude"]).to_numpy()
y_train = obs.get_column("tos").to_numpy()

# Define output grid (testing data) - in this example it
# is a 5 degree global (un-masked) grid
out_grid = np.mgrid[-87.5:90:5, -177.5:180:5]
X_test = out_grid.reshape(2, -1).T

# Initialise the SphericalThinPlateSpline class with the
# training data
```

(continues on next page)

(continued from previous page)

```

stps = SphericalThinPlateSpline(X_train, y_train)

# Fit with lambda estimation using GCV minimisation
stps.fit(set_results=True, compute_stats=True)

# Predict at new points and get the standard_error
y_pred, std_stps = stps.predict(X_test)

```

10.2 Spline

This is the abstract class form of the *Spline*, and should not be used directly and will raise an error if used. Instead one should use either the *ThinPlateSpline* or *SphericalThinPlateSpline* classes.

The documentation for this abstract class is included as it includes the documentation for several shared components between the *Spline* classes.

class glomar_gridding.spline.**Spline**(X, y, error_cov=None)

Generic class for Spline Interpolation and Smoothing approaches.

Do not use this default class.

Parameters

- **X** (ndarray) – The training data positions. Expected shape *[num points, num coords]*. Must be unique positions.
- **y** (ndarray) – The training data values. Expected to have the same *num points* values.
- **error_cov** (ndarray | None) – An optional error covariance. If unset then the default error covariance matrix is used - in most cases this will be the identity matrix.

A_matrix(M_inv=None, X=None, diag_only=False)

The A-matrix for computation of fit/prediction statistics.

Parameters

- **M_inv** (ndarray | None) – The inverse of the fitting matrix, used to compute some statistics. If not supplied, then the M_inv of the fitted result (if the ‘fitted’ attribute is True) is used.
- **X** (ndarray | None) – Optional positions for which to compute the A-matrix. If not supplied then the training positions are used.
- **diag_only** (bool) – Compute only the diagonal of the A-matrix, useful for computing simple statistics of the fit.

Returns

The A-matrix, or a vector containing the diagonal of the A-matrix.

Return type

ndarray

property default_error_cov: ndarray

Default error covariance

abstractmethod dist_func(positions=None)

Distances between each pair of points.

Parameters

positions (ndarray | None) – Optionally compute the distance between each position in this value to the training positions. Otherwise compute the distance between training positions.

Returns

dist – The distances.

Return type

ndarray

property error_cov: ndarray

Error Covariance Matrix

estimate_lambda_gcv(kwargs)**

Use Generalised Cross-Validation to estimate an optimal smoothing parameter.

Parameters

****kwargs** – Keyword arguments to pass to `minimize_scalar`.

Returns

The result of an optimisation for the smoothing parameter (lambda).

Return type

float

 See also

scipy.optimize.minimize_scalar

For optimisation options.

fit(lam=None, set_results=True, compute_stats=True, **kwargs)

Fit the Spline. With optional smoothing. Optionally save the results.

Parameters

- **lam** (float | None) – Optional smoothing parameter. If set to 0 the interpolation will be an **exact** interpolator. If unset, the value will be estimated using a GCV method.
- **set_results** (bool) – Assign the results to the *result* attribute and set the *fitted* attribute to True. This is generally advised.
- **compute_stats** (bool) – Add the statistics to the results attribute.
- ****kwargs** – Keyword parameters to pass to `minimize_scalar` if estimating the smoothing parameter using the GCV approach. Note, if 'lam' is not None then keyword arguments will not be used.

Returns

result – The result of fitting the Spline class using the training data used to initialise the Spline instance. This value is also set to the 'result' attribute if *set_results* is True.

Return type

SplineResult

 See also

scipy.optimize.minimize_scalar
For optimisation options.

get_K(*positions=None*)

Spatial structure matrix determined by the kernel of the Spline model.

Parameters

positions (ndarray | None) – Compute the kernel between the training positions and the values provided. If not set this will return the value from the kernel between the training positions with themselves.

Returns

result

Return type

ndarray

get_stats()

Get the statistics of the fitted result.

Return type

SplineResult

abstractmethod get_trend_basis(*positions*)

Trend Basis

Return type

ndarray

abstractmethod kernel(*distances*)

Spline Kernel

Return type

ndarray

predict(*X_test=None, compute_se=True, result=None*)

Predict for new positions.

Predicts using the weights and trend coefficients from the set *result* attribute, or from a provided *SplineResult* instance.

Parameters

- **X_test** (ndarray | None) – New positions use to predict using the *SplineResult*. Must be unique positions. If unset, the training data is used.
- **compute_se** (bool) – Optionally compute standard error of the prediction.
- **result** (*SplineResult* | None) – Optional *SplineResult* instance to use for prediction (instead of the set *result* attribute). If not set, the *result* attribute must be set and the *fitted* attribute must be True.

Return type

tuple[ndarray, ndarray | None]

Returns

- **prediction** (*numpy.ndarray*) – The predicted values at the test positions.
- **standard_error** (*numpy.ndarray* | None) – Standard error of the predictions at the input positions if *compute_se* is True, otherwise None.

property result: *SplineResult*

Interpolation result, generated from 'fit'

standard_error(*X=None, result=None*)

Standard Error of a prediction.

Calculation follows that of the R Fields package. [Nychka]

Parameters

- **X** (ndarray | None) – Optional positions, typically the same as those used for prediction. If unset then the standard error for the fitted training data is calculated.
- **result** (*SplineResult* | None) – Optional *SplineResult* instance to use for prediction (instead of the set *result* attribute). If not set, the *result* attribute must be set and the *fitted* attribute must be True.

Returns

Vector containing standard error of the prediction at the input positions.

Return type

ndarray

property trend_basis: ndarray

Trend Basis of Training Data

10.3 Thin-Plate-Splines

class glomar_gridding.spline.**ThinPlateSpline**(*X, y, error_cov=None, use_radbas_const=False*)

Thin Plate Spline Interpolation and Smoothing. This uses the standard kernel for thin-plate-spline interpolation:

$$r^2 \times \log(r).$$

Where *r* is the Euclidean distance between a pair of positions. Note that this is undefined for 0 distance. The kernel is instead defined to be 0 at 0 distance, which is the limit.

The methodology use the Euclidean distance metric, applied directly to the input position values.

The trend basis for this method is the positions and intercept. Used as additional predictors for the interpolation. Note, this method may not be appropriate on spherical geometry.

The methodology follows [Hutchinson].

Parameters

- **X** (ndarray) – The training data positions. Expected shape [*num points, num coords*]. Must be unique positions. It is generally recommended that the input positions should be scaled. There are numerous ways to approach this scaling, such as by the range, or the mean and standard deviation. Typically this depends on the input data and domain. It is assumed that the positions have been scaled appropriately.
- **y** (ndarray) – The training data values. Expected to have the same *num points* values.
- **error_cov** (ndarray | None) – An optional error covariance. If unset then the default error covariance matrix is used - in most cases this will be the identity matrix.
- **use_radbas_const** (bool) – Use the Radial Basis constant to scale the computed kernel. This value will scale the kernel to match that computed by the R fields package [Nychka]. This will effect the value of the smoothing parameter (*lam*).

References

[Hutchinson]

property default_error_cov: ndarray

Default Error Covariance, for thin plate spline this is the identity matrix.

dist_func(*positions=None*)

Pairwise Euclidean Distances between each pair of points.

Parameters

positions (ndarray | None) – Optionally compute the distance between each position in this value to the training positions. Otherwise compute the distance between training positions.

Returns

dist – The distances.

Return type

ndarray

fit(*lam=None, set_results=True, compute_stats=True, **kwargs*)

Fit the ThinPlateSpline.

With optional smoothing. Optionally save the results.

Parameters

- **lam** (float | None) – Optional smoothing parameter. If set to 0 the interpolation will be an **exact** interpolator. If unset, the value will be estimated using a GCV method.
- **set_results** (bool) – Assign the results to the *result* attribute and set the *fitted* attribute to True. This is generally advised.
- **compute_stats** (bool) – Add the statistics to the results attribute.
- ****kwargs** – Keyword parameters to pass to `minimize_scalar` if estimating the smoothing parameter using the GCV approach. Note, if ‘lam’ is not None then keyword arguments will not be used.

Returns

result – The result of fitting the Spline class using the training data used to initialise the Spline instance. This value is also set to the ‘result’ attribute if *set_results* is True.

Return type

SplineResult

➡ See also

scipy.optimize.minimize_scalar

For optimisation options.

get_trend_basis(*positions*)

Linear trend basis for universal kriging: [1, x, y, ...].

Parameters

positions (ndarray) – Compute the trend basis of these positions, otherwise the training positions will be used.

Returns

P – The trend-basis. An $(n \times m + 1)$ matrix of ones where n is the number of positions, and m is the number of coordinates.

Return type

ndarray

kernel(*distances*)

Thin Plate Spline RBF kernel.

This is

$$r^2 \times \log(r)$$

or 0 if $r == 0$. Where r is the pairwise distance.

Parameters

distances (ndarray) – The pairwise haversine distances.

Returns

The result of the ThinPlateSpline kernel.

Return type

ndarray

predict(*X_test=None, compute_se=True, result=None*)

Predict for new positions.

Predicts using the weights and trend coefficients from the set *result* attribute, or from a provided SplineResult instance.

Parameters

- **X_test** (ndarray | None) – New positions use to predict using the SplineResult. Must be unique positions. If unset, the training data is used.
- **compute_se** (bool) – Optionally compute standard error of the prediction.
- **result** (*SplineResult* | None) – Optional SplineResult instance to use for prediction (instead of the set *result* attribute). If not set, the *result* attribute must be set and the *fitted* attribute must be True.

Return type

tuple[ndarray, ndarray | None]

Returns

- **prediction** (numpy.ndarray) – The predicted values at the test positions.
- **standard_error** (numpy.ndarray | None) – Standard error of the predictions at the input positions if *compute_se* is True, otherwise None.

standard_error(*X=None, result=None*)

Standard Error of a prediction.

Calculation follows that of the R Fields package. [Nychka]

Parameters

- **X** (ndarray | None) – Optional positions, typically the same as those used for prediction. If unset then the standard error for the fitted training data is calculated.

- **result** (*SplineResult* | None) – Optional SplineResult instance to use for prediction (instead of the set *result* attribute). If not set, the *result* attribute must be set and the *fitted* attribute must be True.

Returns

Vector containing standard error of the prediction at the input positions.

Return type

ndarray

10.4 Spherical Thin-Plate-Splines

class glomar_gridding.spline.SphericalThinPlateSpline(*X*, *y*, *error_cov=None*)

Spherical form of Thin Plate Spline Interpolation and Smoothing.

Adapted from [[Wahba_Sphere](#)].

This approach is specifically designed to operate using spherical geometry, such as the earth. The results will be continuous over the full domain (including at boundaries such as 0, 2pi).

The Kernel used in this Spline interpolation is the solution to equation 3.3 and table 1 from [[Wahba_Sphere](#)] for $m = 2 (q_2)$.

$$q_m(z) = \int_0^1 (1-h)^m (1-2hz+h^2)^{-1/2} dh.$$

Where z is the cosine of the Haversine distance between two points.

To guarantee continuity of the result on spherical geometry the Haversine distance is used, and the trend basis is a vector of ones. This means that the positions are not used (directly) as predictors in this method (unlike the traditional thin-plate spline method). Consequently, this method is closer to ordinary kriging.

Parameters

- **X** (ndarray) – The training data positions, expected to be lat, lon in degrees. The input positions must be unique. Expected shape [*num points*, 2].
- **y** (ndarray) – The training data values. Expected to have the same *num points* values.
- **error_cov** (ndarray | None) – An optional error covariance. If unset then the default error covariance matrix is used - in this case $n \times$ identity matrix, where n is the number of positions in *X*.

References

[[Wahba_Sphere](#)]

property default_error_cov: ndarray

Default Error Covariance, n times identity matrix.

dist_func(*positions=None*)

Pairwise Haversine Distances between each pair of points.

Parameters

positions (ndarray | None) – Optionally compute the distance between each position in this value to the training positions. Otherwise compute the distance between training positions.

Returns

dist – The distances.

Return type
ndarray

fit(*lam=None, set_results=True, compute_stats=True, **kwargs*)

Fit the SphericalThinPlateSpline.

With optional smoothing. Optionally save the results.

Parameters

- **lam** (float | None) – Optional smoothing parameter. If set to 0 the interpolation will be an **exact** interpolator. If unset, the value will be estimated using a GCV method.
- **set_results** (bool) – Assign the results to the *result* attribute and set the *fitted* attribute to True. This is generally advised.
- **compute_stats** (bool) – Add the statistics to the results attribute.
- ****kwargs** – Keyword parameters to pass to `minimize_scalar` if estimating the smoothing parameter using the GCV approach. Note, if ‘lam’ is not None then keyword arguments will not be used.

Returns

result – The result of fitting the Spline class using the training data used to initialise the Spline instance. This value is also set to the ‘result’ attribute if *set_results* is True.

Return type

SplineResult

 See also

scipy.optimize.minimize_scalar

For optimisation options.

get_trend_basis(*positions*)

Trend basis for ordinary kriging: [1]. Note that this does not include the position components, since this could break continuity of the result over the sphere (e.g. at 0, 360 degrees boundary).

Parameters

positions (ndarray) – Compute the trend basis of these positions, otherwise the training positions will be used.

Returns

P – The trend-basis. An (n x 1) matrix of ones where n is the number of positions.

Return type

ndarray

kernel(*distances*)

Spherical Thin Plate Spline RBF kernel, this is q[2] from [[Wahba_Sphere](#)].

Parameters

distances (ndarray) – The pairwise haversine distances.

Returns

The result of the SphericalThinPlateSpline kernel.

Return type

ndarray

predict(*X_test=None, compute_se=True, result=None*)

Predict for new positions.

Predicts using the weights and trend coefficients from the set *result* attribute, or from a provided *SplineResult* instance.

Parameters

- **X_test** (ndarray | None) – New positions use to predict using the *SplineResult*. Must be unique positions. If unset, the training data is used.
- **compute_se** (bool) – Optionally compute standard error of the prediction.
- **result** (*SplineResult* | None) – Optional *SplineResult* instance to use for prediction (instead of the set *result* attribute). If not set, the *result* attribute must be set and the *fitted* attribute must be True.

Return type

tuple[ndarray, ndarray | None]

Returns

- **prediction** (numpy.ndarray) – The predicted values at the test positions.
- **standard_error** (numpy.ndarray | None) – Standard error of the predictions at the input positions if *compute_se* is True, otherwise None.

standard_error(*X=None, result=None*)

Standard Error of a prediction.

Calculation follows that of the R Fields package. [Nychka]

Parameters

- **X** (ndarray | None) – Optional positions, typically the same as those used for prediction. If unset then the standard error for the fitted training data is calculated.
- **result** (*SplineResult* | None) – Optional *SplineResult* instance to use for prediction (instead of the set *result* attribute). If not set, the *result* attribute must be set and the *fitted* attribute must be True.

Returns

Vector containing standard error of the prediction at the input positions.

Return type

ndarray

10.5 Spline Result

class glomar_gridding.spline.**SplineResult**(*lam, weights, trend_coefs, M_inv, y_fit, parent, compute_stats=True*)

Result class for a Spline, holds result attributes from the spline fitting.

Parameters

- **lam** (float) – The lambda value used in fitting, controls the smoothing.
- **weights** (ndarray) – The weights from fitting a Spline.
- **trend_coefs** (ndarray) – The trend coefficients from fitting a Spline.
- **M_inv** (ndarray) – The inverse of the fitting matrix, used to compute some statistics.

- **y_fit** (ndarray) – The fitted result of the Spline interpolation at the input coordinates.
- **parent** (_Interpolator) – The Spline instance used to generate the result. Requires ‘kernel’, ‘trend_basis’ and ‘y’ attributes to be set.
- **compute_stats** (bool) – Compute fitting statistics. Sets statistical attributes to the instance.

fit_stats()

Get statistics from a Spline fit. Requires the ‘parent’ attribute to be set and to be a Spline with ‘kernel’, ‘trend_basis’, and ‘y’ attributes to be set.

signal

Signal of the fit.

Type

float

error

Error.

Type

float

gcv

Generalised Cross-Validation.

Type

float

msr

Mean square error.

Type

float

var

Variance.

Type

float

sigma2

The sigma MLE for the fit.

Type

float

tau2

The squared Tau MLE for the fit.

Type

float

Return type

None

GRID CLASS

An alternative workflow is to use the *Grid* class added in version 1.1.0. The *Grid* class can be used to run the full Kriging process from defining an output grid to interpolating the observational data.

This class allows for easy masking of positions, so that masked positions are not infilled.

A new notebook was added (“Mask_example.ipynb”) to demonstrate the full process.

```
from glomar_gridding.grid import Grid

# Initialise, like with grid_from_resolution
grid = Grid.from_resolution(
    resolution=5,
    bounds=[(-90, 90), (-180, 180)],
    coord_names=["latitude", "longitude"],
    definition="left",
)

# Can now easily add a mask and only used unmasked positions on the grid
mask = ...
grid.add_mask(mask)

# Add distance matrix and covariance matrix
grid.distance_matrix()
grid.covariance_matrix(
    "matern",
    range=1300,
    psill=1.2,
    nu=1.5,
    nugget=0.0,
    method="sklearn",
)

# Add observations
obs = ...
grid.map_observations(obs, ...)

# Can ensure an external error covariance matrix is aligned to the grid
error_cov = ...
error_cov = grid.prep_covariance(error_cov)

# Perform Kriging
grid.kriging("ordinary", error_cov)
```

class glomar_gridding.grid.**Grid**(*coords*)

Grid Class.

Allows for construction of a grid to use in the Kriging process. A mask can be applied to the grid with the *add_mask* method. If a mask is applied then computed fields will account for the mask. This has the ability to improve performance of the Kriging process as masked positions will not be infilled.

The class is constructed from an instance of *xarray.Coordinates*, however it can also be constructed using the *from_resolution* method, which replicates the behaviour of *glomar_gridding.grid.grid_from_resolution()*.

Parameters

coords (Coordinates) – The coordinates of the output grid.

Examples

```
>>> grid = Grid(coordinates)
>>> grid.add_mask(mask)
>>> grid.distance_matrix() # Haversine by default
>>> grid.covariance("matern", psill=0.36, range=1300, nugget=0, nu=1.5)
>>> grid.map_observations(obs)
>>> grid.kriging("ordinary", error_covariance)
>>> masked_infilled = grid.krige.solve()
>>> infilled = grid.assign(masked_infilled)
```

add_mask(*mask*)

Add and apply a mask to the grid.

Parameters

mask (ndarray | MaskedArray) – The mask to apply, either an array of Booleans or a masked array.

Return type

None

assign_values(*values*, *grid_idx=None*, *fill_value=nan*, *apply_mask=True*)

Assign a vector of values to a grid, using a list of grid index values.

Parameters

- **values** (ndarray) – The values to map onto the output grid.
- **grid_idx** (ndarray | None) – The 1d index of the grid (assuming “C” style ravelling) for each value. If unset then it is assumed that the values are complete for the (possibly masked) grid.
- **fill_value** (Any) – The value to fill unassigned grid boxes. Must be a valid value of the input *values* data type.
- **apply_mask** (bool) – The input *grid_idx* represents the index for the masked grid, apply the mapping to the final grid index.

Returns

out_grid – A new grid containing the values mapped onto the grid.

Return type

DataArray

property coord_df: DataFrame

Convert the coordinates to a DataFrame

property coord_names: list[str]

Names of the coordinates

property coords: Coordinates

Grid coordinates

property covariance: ndarray | DataArray

Spatial Covariance matrix

covariance_matrix(variogram, **kwargs)

Compute a covariance matrix from the grid using the distance matrix and a variogram model, this requires the pre-computation of the distance matrix attribute (*dist*). If the grid is masked then the resulting *covariance* attribute (set by this method) will be filtered to align with the mask.

This sets the *covariance* and *variogram* attributes.

Parameters

- **variogram** (Literal['exponential', 'gaussian', 'matern', 'spherical']) – Lower-case name of the variogram model to use. One of “exponential”, “gaussian”, “matern”, or “spherical”. The *variogram* attribute will be set to an instance of the equivalent class from *glomar_gridding.variogram*.
- ****kwargs** – Keyword arguments for the variogram model.

Return type

None

distance_matrix(dist_func=<function haversine_distance_from_frame>, **dist_kwargs)

Calculate a distance matrix between all positions in a grid. Orientation of latitude and longitude will be maintained in the returned distance matrix.

One distances between unmasked coordinates will be calculated if the grid is masked.

This sets the *dist* attribute.

Parameters

- **grid** (xarray.DataArray) – A 2-d grid containing latitude and longitude indexes specified in decimal degrees.
- **dist_func** (Callable) – Distance function to use to compute pairwise distances. See *glomar_gridding.distances.calculate_distance_matrix* for more information. Defaults to Haversine distances.
- ****dist_kwargs** – Keyword arguments to pass to the distance function. This may include requirements for the name of specific coordinates, for example “lat_coord” and “lon_coord”.

Return type

None

Examples

```
>>> grid = Grid.from_resolution(
    resolution=5,
    bounds=[(-87.5, 90), (-177.5, 180)], # Lower bound is centre
    coord_names=["lat", "lon"]
)
>>> grid.to_distance_matrix(grid, lat_coord="lat", lon_coord="lon")
```

(continues on next page)

(continued from previous page)

```
>>> grid.dist
<xarray.DataArray 'dist' (index_1: 2592, index_2: 2592)> Size: 54MB
array([[ 0.,          24.24359308,   48.44112457, ...,
        19463.87158499, 19461.22915012, 19459.64166305],
       [ 24.24359308,   0.,          24.24359308, ...,
        19467.56390938, 19463.87158499, 19461.22915012],
       [ 48.44112457,   24.24359308,   0.,          ...,
        19472.29905588, 19467.56390938, 19463.87158499],
       ...,
       [19463.87158499, 19467.56390938, 19472.29905588, ...,
        0.,          24.24359308,   48.44112457],
       [19461.22915012, 19463.87158499, 19467.56390938, ...,
        24.24359308,   0.,          24.24359308],
       [19459.64166305, 19461.22915012, 19463.87158499, ...,
        48.44112457,   24.24359308,   0.          ]],
       shape=(2592, 2592))
Coordinates:
  * index_1  (index_1) int64 21kB 0 1 2 3 4 ... 2587 2588 2589 2590 2591
  * index_2  (index_2) int64 21kB 0 1 2 3 4 ... 2587 2588 2589 2590 2591
    lat_1    (index_1) float64 21kB -87.5 -87.5 -87.5 ... 87.5 87.5
    lon_1    (index_1) float64 21kB -177.5 -172.5 ... 172.5 177.5
    lat_2    (index_2) float64 21kB -87.5 -87.5 -87.5 ... 87.5 87.5 87.5
    lon_2    (index_2) float64 21kB -177.5 -172.5 ... 172.5 177.5
```

classmethod `from_resolution(resolution, bounds, coord_names, definition='center')`

Generate a grid from a resolution value, or a list of resolutions for given boundaries and coordinate names.

Note that all list inputs must have the same length, the ordering of values in the lists is assumed align.

The constructed grid will be regular, in the sense that the grid spacing is constant. However, the resolution in each direction can be different, allowing for finer resolution in some direction.

Parameters

- **resolution** (float | list[float]) – Resolution of the grid. Can be a single resolution value that will be applied to all coordinates, or a list of values mapping a resolution value to each of the coordinates.
- **bounds** (list[tuple[float, float]]) – A list of bounds of the form (*lower_bound*, *upper_bound*) indicating the bounding box of the returned grid. For a grid between -180 and 180 degrees in Longitude, one would set the bounds in this dimension to (-180, 180) (if using “left” for the ‘definition’ argument).
- **coord_names** (list[str]) – List of coordinate names in the same order as the bounds and resolution(s).
- **definition** (Literal['center', 'left']) – Each grid box is defined by the centre of the box in each coordinate. This argument defines whether the lower bound in each coordinate direction defines the “left”-edge or the “center” of the 1st grid-box.
 - “left”: The lower-bound can be used to define the left-side of each grid-box. The grid-box value will be the left-edge + 1/2 of resolution in each direction.
 - “center”: The lower-bound can be used to define the center of each grid-box. The grid-box value will be the center in each direction

Returns

grid – The grid defined by the resolution and bounding box.

Return type

Grid

Examples

```
>>> grid = Grid.from_resolution(
    resolution=5,
    bounds=[(-87.5, 90), (-177.5, 180)], # Lower bound is center
    coord_names=["lat", "lon"],
    definition="center",
)
>>> grid.grid
<xarray.DataArray (lat: 36, lon: 72)> Size: 21kB
array([[nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       ...,
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan]], shape=(36, 72))
Coordinates:
  * lat      (lat) float64 288B -87.5 -82.5 -77.5 ... 77.5 82.5 87.5
  * lon      (lon) float64 576B -177.5 -172.5 ... 172.5 177.5
```

property grid_idx: ndarray

All grid indices

property index_map: DataFrame

The mapping between mask and grid indices

kriging(*kriging_method*, *error_cov*=None)

Add a Kriging class object to the grid. Allowing for easy Kriging over a masked grid.

This sets the *krige* attribute with a instance of a `glomar_gridding.kriging.Kriging()` class object, which will have all of the attributes and methods of that class.

Parameters

- **kriging_method** (Literal['simple', 'ordinary', 'stochastic']) – Lower-case name of the Kriging class to use. One of “simple”, “ordinary”, or “stochastic”.
- **error_cov** (ndarray | DataArray | None) – Optional error covariance matrix. This will apply a smoothing effect to the observation points (as well as over the infilling). This can have one of the following sizes:
 - grid size (unmasked) times grid size (unmasked)
 - grid size (masked) times grid size (masked)
 - number of observations times number of observations

Return type

None

map_observations(*obs*, *obs_coords*=['lat', 'lon'], *sort*=True, *bounds*=None, *add_grid_pts*=True, *grid_prefix*='grid_', *apply_mask*=True, *obs_col*=None, *set_vals*=True)

Align an observation dataframe to the Grid.

Maps observations to the nearest grid-point, and sorts the data by the 1d index of the DataArray in a row-major format.

The grid defined by the latitude and longitude coordinates of the input DataArray is then used as the output grid of the Gridding process.

Parameters

- **obs** (DataFrame) – The observational DataFrame containing positional data with latitude, longitude values within the *obs_latname* and *obs_lonname* columns respectively. Observations are mapped to the nearest grid-point in the grid.
- **obs_coords** (list[str]) – Names of the column containing positional values in the input observational DataFrame.
- **sort** (bool) – Sort the observational DataFrame by the grid index
- **bounds** (list[tuple[float, float]] | None) – Optionally filter the grid and DataFrame to fall within spatial bounds. This list must have the same size and ordering as *obs_coords* and *grid_coords* arguments.
- **add_grid_pts** (bool) – Add the grid positional information to the observational DataFrame.
- **grid_prefix** (str) – Prefix to use for the new grid columns in the observational DataFrame.
- **apply_mask** (bool) – If the Grid is masked convert the grid index values to the masked grid values. This will drop any observations whose grid position is masked.

Returns

obs – Containing additional *grid_**, and *grid_idx* values indicating the positions and grid index of the observation respectively. The DataFrame is also sorted (ascendingly) by the *grid_idx* columns for consistency with the gridding functions.

Return type

DataFrame

Examples

```
>>> obs = pl.read_csv("/path/to/obs.csv")
>>> grid = Grid.grid_from_resolution(
    resolution=5,
    bounds=[(-87.5, 90), (-177.5, 180)], # Lower bound is centre
    coord_names=["lat", "lon"]
)
>>> obs = grid.map_observations(obs)
```

prep_covariance(covariance_matrix)

Apply the mask to a covariance matrix. This could be used to resize an error covariance matrix ahead of masked kriging, for example.

Parameters

covariance_matrix (ndarray | DataArray) – The covariance matrix.

Return type

ndarray | DataArray

remove_mask()

Remove the mask

Return type

None

select_bounds(*bounds*)

Filter the Grid by a set of bounds. Updates the *grid* attribute.

Parameters

- **bounds** (*list[tuple[float, float]]*) – A list of tuples containing the lower and upper bounds for each dimension.
- **variables** (*list[str]*) – Names of the dimensions (the order must match the bounds).

Return type

None

set_covariance(*covariance_matrix*)

Set a covariance matrix. This is automatically adjusted if the grid is masked.

Sets the *covariance* attribute.

Parameters

covariance_matrix (*ndarray | DataArray*) – The covariance matrix for the full (unmasked) grid.

Return type

None

set_distance_matrix(*distance_matrix*)

Set a distance matrix. This is automatically adjusted if the grid is masked.

Sets the *dist* attribute.

Parameters

distance_matrix (*ndarray | DataArray*) – The distance matrix for the full (unmasked) grid.

Return type

None

property shape: *tuple[int, ...]*

Shape of the grid

property size: *int*

Size of grid

AUTO REGRESSIVE FORECASTING AND KALMAN FILTER

12.1 Introduction and purpose

Most marine datasets are what we called a 2DVAR analysis - the current analysis is based on the present observations only (2D as in x, y ; in contrast of 3DVAR seen in historical weather forecasting literature: x, y, z).

This limits the amount of the observations one can use, resulting in higher uncertainty. In many cases, this is mitigated by temporally aggregating the observations (like doing monthly analysis instead of daily analysis) or blending multiple input sources (like high-level remote sensing data product). Sometimes one is stuck with such spotty observations; like one really wants a daily analysis based purely on daily observations. How we can reduce the uncertainties?

This leads to the introduction of time dimension into the analysis. This is what we called 4DVAR (x, y, z, t) weather forecasting (or 3DVAR if one is dealing with only surface variables - x, y, t).

A popular way to do so is via Kalman Filter – a recursive method that blends a forecast to the present state (aka prior) with present observations, resulting a posterior analysis that is more accurate with just the forecast or the observations. The method is widely used in autopilot and positioning. For Earth sciences, they are often associated with numerical weather prediction and reanalyses. Kalman filter has been used in producing more accurate weather forecasts and historical reanalysis by the ECWMF and the UK Met Office (https://www.met.rdg.ac.uk/~darc/training/ecmwf_collaborative_training/EnKF_AFowler.pdf).

12.2 Producing the prior using a simple autoregressive statistical model

This lead to the question how would one produce a forecast priors. The priors/first guesses are often based on outputs of dynamical models like in numerical weather prediction and reanalysis; even in satnav or autopilot, simple dynamics are included (dead-reckoning along roads). Such dynamics are often not part of the development of the gridded observations; GlomarGridding is a package for statistical modelling only.

A simple statistical approach is implemented here to produce such forecast, basing on the simple idea is that current condition (described by the variable a below) will generally tend toward climatology ($E(a)$). The simplest way to model that behavior is the local 1st order autoregressive model, in which the expected value for $a(t)$ is a linear function of the observed lag-1 correlation and its last observed value ($a(t-1)$), resulting in exponential relaxation toward $E(a)$, plus a normal-distributed uncertainty epsilon:

$$a_t = \Phi(a_{t-1} - E(a)) + \epsilon$$

This is often referred as the (non-vectorised) **AR1 model**. More sophisticated approach can be used like full vectorised form (requires full lagged autocovariance matrix), multi-lag, or autoregressive moving average (ARMA).

The AR1 model produces its own uncertainty. Input uncertainties can be incorporated following standard propagation of error rules. This is essential as uncertainty lie in the heart of Kalman Filter. Let say Φ is a diagonal matrix with local lag-1 correlation, $\Sigma_{\text{climatology}}$ is the diagonal matrix with climatological variance (no off-diagonal contribution; this

requires full vectorised AR1), Σ_t is a diagonal matrix of the diagonal of error covariance for the current observations, the full uncertainty can be estimated as:

$$\epsilon_{\text{AR1}} \sim \text{MVN}(0, \Phi \Sigma_t \Phi^T + \Sigma_{\text{climatology}} - \Phi \Sigma_{\text{climatology}} \Phi^T)$$

The above form is correct even for weights that has off-diagonal terms. In that case, Φ is no longer a diagonal autocorrelation matrix but the weights for fully vectorised multi-variate model $\mathbf{W}$; $\mathbf{W} = \Phi$ is a merely a special case.

$$a_t = \mathbf{W}(a_{t-1} - E(a)) + \epsilon$$

In the local and diagonal case, the above is equivalent to the below at each grid point, like the equations that one can find in Chapter 8 of [Wilks]:

$$\epsilon_{\text{AR1}} \sim \text{N}(0, \phi^2 \sigma_t^2 + (1 - \phi^2) \sigma_{\text{climatology}}^2)$$

[Wikipedia intro](#)

[Wikipedia intro](#)

For the diagonal Φ case, the lagged autocorrelation is easy to compute from observations; its application to statistical forecasting is highly stable and quick. The weights for the multi-variate correlated case are much more complicated. It is usually estimated by fitting each grid point with a regression model, taking advantage of the seemingly unrelated regression of vector autoregression. For high resolution datasets, fitting such models often results in overfitting and low predictive power. This can be mitigated using regularized regression. The one implemented here uses LASSO regression [Tibshirani_Lasso] via `scikit-learn`; LASSO generally leads to sparse $\mathbf{W}$, makes them easier to handle memory and computational cost via sparse matrix classes and function within `scipy`. What LASSO does is that it adds a penalty term to the least squares cost function that is proportional to the sum of the absolute values of all regression coefficients. The sum is scaled by a tunable hyperparameter; it is usually called λ in statistics literature (including the Wikipedia article) but is called α instead in `scikit-learn` documentation.

12.3 Uncertainty weighted average (Kalman Filter)

Kalman Filter, in its heart, is a weighted average – an average for two input streams: the prior/first-guess/forecast what the current state probably be, and what the most recent observations say. The weights are determined only by uncertainties as neither the first-guess and observations are perfect. In simple English, produce a better observations by taking multiple sources, weighting in favour the more accurate one.

Let say the uncertainty of the forecast and observations can be described by two different uni- or multi-variate zero-averaged normal distribution with some scalar variance / error covariance matrix.

$$\epsilon_{\text{forecast}} \sim \text{N}(0, \sigma_{\text{forecast}}^2) \text{ or } \text{MVN}(0, \Sigma_{\text{forecast}})$$

$$\epsilon_{\text{obs}} \sim \text{N}(0, \sigma_{\text{obs}}^2) \text{ or } \text{MVN}(0, \Sigma_{\text{obs}})$$

Assuming observations are on the same grid as the forecast (true for the exercise), the Kalman gain (weight for new observations) is:

Univariate:

$$K = \frac{\sigma_{\text{forecast}}^2}{\sigma_{\text{forecast}}^2 + \sigma_{\text{obs}}^2}$$

Multi-variate:

$$K = \Sigma_{\text{forecast}} (\Sigma_{\text{forecast}} + \sigma_{\text{obs}})^{-1} = \sigma_{\text{obs}}^{-1} (\Sigma_{\text{forecast}}^{-1} + \sigma_{\text{obs}}^{-1})^{-1}$$

in which K is solution to the following linear equation, this is how it is solved in the code using *numpy.linalg.solve*. No actual error covariance inversions is needed.

$$(\Sigma_{\text{forecast}} + \sigma_{\text{obs}})K^T = \Sigma_{\text{forecast}}$$

noting that Sigma is symmetrical but K is not (similar to Kriging weights).

Once the weights are known, the posterior analysis is a weighted average of forecast and observations

$$\tilde{a}(t) = K a_{\text{obs}} + (I - K) a_{\text{forecast}} = K(a_{\text{obs}} - a_{\text{forecast}}) + a_{\text{forecast}}$$

with posterior error covariance of:

$$\Sigma_{\text{analysis}} = (I - K)\Sigma_{\text{forecast}}$$

If observations are not on the same grid as the forecast, additional mapping multiplications are needed; that leads the usual form one will see in text books or Wikipedia.

The workflow here will be:

1. Produce an initial $t = 0$ analysis
2. Do a $t+1$ forecast
3. Krige the observations for $t + 1$
4. Blend the forecast with Kriging result using Kalman filter
5. Go back to step 2 using the posterior analysis and error covariance

12.4 Code and usage

The main processing involve two modules

- *autoregressive_kalman.forecast_linear_ar1*
- *autoregressive_kalman.inv_sigma_wgts*

There is a helping code module called *autoregressive_kalman.cov_diagonal* which can be used to separate the analysis between we expect to not to be spatially correlated with bits that are spatially correlated.

For the all the main classes, the computation uses a method that is called *compute_XXXXX*.

12.5 *autoregressive_kalman.forecast_linear_ar1*

For autoregressive forecasting, providing the weights and autocorrelation already computed.

```
class glomar_gridding.autoregressive_kalman.Autoregressive1ForecastUncorr(analysis,
                                                                    errcov_analysis,
                                                                    lag_1_autocor,
                                                                    clim_mean,
                                                                    clim_covar, er-
                                                                    rcov_analysis_is_sdev=False,
                                                                    clim_covar_is_sdev=False)
```

Class to compute AR1 forecast

Local/uncorrelated only: Lag-1 correlations are the weights and are in 1-D vector. All predictions are based on local autocorrelation only

Use Autoregressive1Forecast for full vector case.

Compute Lag-1 autoregressive forecast as a prior for Kalman filter.

Parameters

- **analysis** (ndarray) – 1D vector of analysis (such as Kriging and Kalman filter outputs) for $t=t$; dtype *float*, shape $(N,)$
- **errcov_analysis** (ndarray) – 1D error variance or 2D error covariance for analysis; if latter, it will only use the diagonal; dtype *float*, shape $(N,)$ | (N, N)
- **lag_1_autocor** (ndarray) – 1D vector of lag-1 autocorrelation; shape should be same as analysis; dtype *float*, shape $(N,)$
- **clim_mean** (ndarray) – 1D vector of climatological mean; shape should be same as analysis; dtype *float*, shape $(N,)$
- **clim_covar** (ndarray) – Climatological covariance; can be 1D or 2D; shape of this should either be $n \times n$ (like ellipse covariance) or n (vector of variance at each grid point); if a 2D matrix is provided, it will take the diagonal; dtype *float*, shape $(N,)$ | (N, N)
- **errcov_analysis_is_sdev** (bool) – Flag indicating if *errcov_analysis* is variance or standard deviation; this only applies to vectored form ($n \times 1$ shape) of *errcov_analysis*; default *False*, *errcov_analysis* is variance (like SST^{**2}) or full covariance matrix; *True* if standard deviation
- **clim_covar_is_sdev** (bool) – Flag indicating if *clim_covar* is variance or standard deviation; this only applies to vectored form ($n \times 1$ shape) of *clim_covar*; default *False*, *clim_covar* is variance (like SST^{**2}) or full covariance matrix; *True* if standard deviation

predict (*full_errcov_out=True*)

Compute AR1 forecast and estimate uncertainties

Speed: <https://stackoverflow.com/questions/44388358/python-numpy-matrix-multiplication-with-one-diagonal-matrix>This version uses scalar multiplication, less use of `np.diag`, and should be faster.**Parameters****full_errcov_out** (bool) – If *True*, new attribute *errcov* will be a diagonal matrix**forecast**

AR1 forecast

Type`numpy.ndarray`**errcov**

The error covariance for the forecast

Type`numpy.ndarray`

```
class glomar_gridding.autoregressive_kalman.Autoregressive1ForecastVector(analysis,  
                                                                           errcov_analysis,  
                                                                           weights,  
                                                                           clim_mean,  
                                                                           clim_covar, estimated_ar1_errcov=None)
```

Class to compute AR1 forecast in full vectorized form

Weights can be provided. It can be computed from non-lagged and lag-1 covariances, using classmethod *from_autocov*.Latter is possibly unstable. Typical spatial 0-lag covariances are computed using variogram techniques, and they may not work with lag-1 and there is no sure way to ensure $\langle x(t-1), x(t) \rangle @ \langle X(t), X(t) \rangle^{*-1}$ is stable. GloMarGridding has no functionality to compute lagged autocovariances.

AutoregressiveForecastUncorr is a special case of this class with the weights being the lag-1 autocorrelation.

Weights are only stable if its spectral radius is less than 1. Otherwise weights violate weak temporal stationarity assumption.

Instead usually weights are estimated separately using regression taking advantage of seemingly (un)related regression (SUR) method with additional use of regularization (like Lasso) to deal with weight stability, multi-collinearity, and insufficient sample sizes/too many covariates. Weights estimated this way often meets spectral radius < 1 requirement.

Regardless, this method is much more expensive. It is in theory more complete and deals with spatially correlated component of the analysis.

Compute lag-1 autoregressive forecast as a prior for Kalman filter.

Other notes:

(1) If *analysis* is a Kriging analysis, results here produces the same result as if applying spatiotemporal Kriging for $t+1$

The weight for spatiotemporal Kriging is: $W_Krige(t, t-1) = C(t, t-1) D.T @ INV(C(t, t) + E)$

That is the same as weights for AR1 multiplying on top of Kriging (t,t): $W_AR = C(t, t-1) @ INV(C(t, t))$
 $W_Krige(t,t) = C @ D.T @ INV(C(t, t) + E)$

Uncertainty: It can be shown that if *errcov_analysis* are Kriging covariance that:

$\sigma_{sq_eps} + \sigma_{sq_analysis} ==$ spatiotemporal Kriging for $t+1$

In practice, *analysis* will be computed recursively using Kalman filter.

(2) Uncertainty of multivariate autoregression: <https://math.stackexchange.com/questions/5004102/covariance-of-a-multivariate-autoregression>

Parameters

- **analysis** (ndarray) – 1D vector of analysis (such as Kriging and Kalman filter outputs) for $t=t$; dtype *float*, shape $(N,)$
- **errcov_analysis** (ndarray) – 2D error covariance for analysis; dtype *float*, shape (N, N)
- **weights** (ndarray | sparray) – 2D weights; dtype *float*, shape (N, N)
- **clim_mean** (ndarray) – 1D vector of climatological mean; shape should be same as analysis; dtype *float*, shape $(N,)$
- **clim_covar** (ndarray) – Climatological covariance; can be 1D or 2D; shape of this should either be $n \times n$ (like ellipse covariance) or n (vector of variance at each grid point); dtype *float*, shape (N, N)
- **estimated_ar1_errcov** (ndarray | None) – A precomputed estimate for $C - WCW^T$; defaults to None in which that is computed from *weights* and *clim_covar*

classmethod from_autocov(*analysis*, *errcov_analysis*, *lag_1_autocov*, *clim_mean*, *clim_covar*, *stability_perturbation=None*, *estimated_ar1_errcov=None*)

With provided *lag_1_autocov* (lag-1 autocovariance), estimate weights using $W = \text{lag1_autocov} @ \text{inv}(\text{clim_covar})$

$W = \langle x(t), x(t-1) \rangle @ \langle x(t), x(t) \rangle^{-1}$

$W @ \langle x(t), x(t) \rangle = \langle x(t), x(t-1) \rangle$

$\langle x(t), x(t) \rangle @ W.T = \langle x(t), x(t-1) \rangle$ <<< computed using `np.linalg.solve`

Both $\langle x(t), x(t) \rangle$ and $\langle x(t), x(t-1) \rangle$ are symmetric

This approach is often unstable as `lag1_autocov` and `clim_covar` are estimated separately, and there is no guarantee that the computed weights satisfy weak stationary requirement that the spectral radius of weights have to be less than 1.

Parameters

- **analysis** (ndarray) – 1D vector of analysis (such as Kriging and Kalman filter outputs) for $t=t$; dtype *float*, shape $(N,)$
- **errcov_analysis** (ndarray) – 2D error covariance for analysis; dtype *float*, shape (N, N)
- **lag_1_autocov** (ndarray) – 2D lag-1 autocovariance; dtype *float*, shape (N, N)
- **clim_mean** (ndarray) – 1D vector of climatological mean; shape should be same as analysis; dtype *float*, shape $(N,)$
- **clim_covar** (ndarray) – Climatological covariance; can be 1D or 2D; shape of this should either be $n \times n$ (like ellipse covariance) or n (vector of variance at each grid point); dtype *float*, shape (N, N)
- **stability_perturbation** (float | None) – A diagonal perturbation to be added `clim_covar` for the purpose of computing the weights. This may avoid problems that `clim_covar` being (close to) invertible.

predict(*check_wgt_stability=True, full_errcov_out=True, check_errcov_psd=True, **clip_kwargs*)

Compute VAR forecast using weights or covariances that are part of class.

Parameters

- **check_wgt_stability** (bool) – Check weights (in extension for the auto and climatological covariances) are stable; see `var_stability_check`
- **full_errcov_out** (bool) – Return full error covariance if True; defaults to *True*
- **check_errcov_psd** (bool) – Check if error covariance is positive semi-definite; fix if needed
- **clip_kwargs** (*dict* | None) – kwargs for `_check_and_fix_psd_errcov_by_clipping`

forecast

AR1 forecast

Type

numpy.ndarray

errcov

The error covariance for the forecast

Type

numpy.ndarray

var_stability_check(*warn_instead=True*)

A multivariate stationary AR model is only stable if and only if all eigvals of the weights have absolute values less than 1 <https://kevinkotze.github.io/ts-7-slide/#10> which is from <https://www.jstor.org/stable/j.ctv14jx6sm> Chapter 10 Proposition 10.1

If weights fail stability check, it implies expected value of anomalies can grow in time. Testing shows, one also gets wrong uncertainty estimates, like off diagonal term being (considerably) larger than diagonal terms; that's an invalid (error) covariance matrix, saying correlation > 1 which is impossible!

Parameters

- **warn_instead** (bool) – By default, this method only generates a warning attribute; if this is set to *True*, method will raise exception

12.6 *autoregressive_kalman.inv_sigma_wgts*

For Kalman Filtering, using inverse error covariance weighting (*sigma*) weights.

```
class glomar_gridding.autoregressive_kalman.KalmanOut(forecast_vector, obs_vector, errcov_forecast,
                                                    errcov_obs, cov_forecast_and_obs,
                                                    use_diag_only=True)
```

Class to compute blended forecast and observations variable names follows compute_inv_variance_wgt_mean_kalman

Note: If only one of the two error covariances matrices is a vector/diagonal matrix but the other one is not (square) and is much denser with comparable diagonals, KF and variance weighted means may strongly favour the vector/diagonal matrix, and give poor results (even if results are strictly correct). Check your underlying reason why one is doing that to make sure that is what one intended.

Vector autoregression should only match with full Kriging covariance (a fully spatially correlated Kalman filter); local time series autoregression would normally only match with local uncertainty of the Kriging covariance (i.e. the Kriging uncertainty) (i.e. effectively a 1D Kalman Filter).

Parameters

- **forecast_vector** (ndarray) – 1D vector of forecasts
- **obs_vector** (ndarray) – 1D vector of (gridded) observations
- **errcov_forecast** (ndarray) – 2D matrix of errcov for *forecast_vector*
- **errcov_obs** (ndarray) – 2D matrix of errcov for *obs_vector*
- **cov_forecast_and_obs** (ndarray | None) – Covariance between forecast/first-guess & observations In NWP-DA, this is usually assumed to be zero (aka observation and forecast errors are not correlated); set it to *None* if that is zero (and *None* is the default).
- **use_diag_only** (bool) – Ignore off-diagonals of *errcov_forecast*, *errcov_obs*, and *cov_forecast_and_obs* if set to *True*, default *True*.

predict()

Compute the inverse variance weighted average of *forecast_vector* & *obs_vector* using provided error covariances *errcov_forecast* & *errcov_obs*

This uses a form that requires only ONE matrix inverses and reciprocals (good!) and is more commonly seen in Kalman Filter guides (including the form uses in Wikipedia) https://en.wikipedia.org/wiki/Kalman_filter; Probably everyone hate matrix inverses... (for good reason).

1D vector form of Kalman gain:

$$\text{gain} = \text{forecast_err} / (\text{forecast_err} + \text{obs_err})$$

Matrix form of Kalman gain:

$$\text{gain} = \text{forecast_err} @ \text{inv}(\text{forecast_err} + \text{obs_err})$$

$$\text{gain} @ (\text{forecast_err} + \text{obs_err}) = \text{forecast_err}$$

$$(\text{gain} @ (\text{forecast_err} + \text{obs_err})).T = \text{forecast_err}.T$$

$$(\text{forecast_err} + \text{obs_err}).T @ \text{gain}.T = \text{forecast_err}.T$$

However, *forecast_err* and *obs_err* are symmetric; hence, solve set of linear equations that:

$$(\text{forecast_err} + \text{obs_err}) @ \text{gain}.T = \text{forecast_err}$$

wgt_mean

Posterior analysis

Type

numpy.ndarray

errcov

Posterior error covariance

Type

numpy.ndarray

kalman_gain_from_new_obs

Kalman gain (either a 2D matrix or 1D vector) aka the weights assigned to the new information

Type

numpy.ndarray

wgts_from_ar_forecast

Identity matrix or one vector minus Kalman gain aka the weights assigned to the prior/first-guess

Type

numpy.ndarray

sparse_approx_for_errcov(*sparse_threshold=1e-06, convert_to_sparse=True*)Setting small elements of *self.errcov_obs* and *errcov_forecast* to zero as defined by *sparse_threshold* value.

If there are small values in the diagonals.

This reduces memory footprint and may make block-splitting by *cov_diagonal* easier.It is not intended to be use for *use_diag_only mode*. This is only used to handle situations that the error covariances are big.**Parameters**

- **sparse_threshold** (float) – (Absolute) off-diagonal values that are to set to zero
- **convert_to_sparse** (bool) – Convert error covariances to *scipy.sparse.sparray* if *True*

```
class glomar_gridding.autoregressive_kalman.KalmanOutUncorrCorrSplit(forecast_vector,  
                                                                    obs_vector,  
                                                                    errcov_forecast,  
                                                                    errcov_obs,  
                                                                    cov_forecast_and_obs,  
                                                                    arr_to_decide_isolated_pts,  
                                                                    zero_threshold=1e-06)
```

Class to compute blended forecast and observations

This splits the error covariances into diagonal and non-diagonal bits

This is a wrapper for KalmanOut.

Parameters

- **forecast_vector** (ndarray) – 1D vector of forecasts
- **obs_vector** (ndarray) – 1D vector of (gridded) observations
- **errcov_forecast** (ndarray) – 2D matrix of errcov for *forecast_vector*
- **errcov_obs** (ndarray) – 2D matrix of errcov for *obs_vector*
- **cov_forecast_and_obs** (ndarray | None) – Covariance between forecast & observations
Set it to *None* if that is zero

- **arr_to_decide_isolated_pts** (ndarray) – The matrix (usually a covariance) to determine if the point is diagonally isolated. This can be the prior spatial covariance.
- **zero_threshold** (float) – The threshold applied to *arr_to_decide_isolated_pts* to decide if the row/column is diagonal.

predict()

Blend the results of uncorrelated and correlated streams

predict_corr()

Run *KalmanOutUncorrCorrSplit.corr_part.predict*; *corr_part* is an instance of *KalmanOut*

predict_uncorr()

Run *KalmanOutUncorrCorrSplit.uncorr_part.predict*; *uncorr_part* is an instance of *KalmanOut*

12.7 autoregressive_kalman.cov_diagonal

Diagonal and non-diagonal matrix tools

remove_diag_only_rows: - Auto-detect diagonal only rows/columns from an input matrix - Purge them to form a smaller matrix and saves the matrix

restore_diag_only_rows: - reverses the process, but one needs to give a filler value

diag_and_nondiag_rows_subsamplers: - get the subsampling matrix of diagonal and off-diagonal rows/columns - option to return the actual subsampled array

`glomar_gridding.autoregressive_kalman.cov_diagonal.diag_and_nondiag_rows_subsamplers(cov, zero_threshold=1e-06, return_subsampled_arr=True)`

Get the subsampling matrices for rows and columns with only diagonal-only elements and off-diagonal elements.

“Diagonal-only” elements are defined according to the *zero_threshold*. Any rows and columns that do not satisfy the “diagonal-only” definition is considered to have off-diagonal elements.

Returns a tuple with up to 4 matrices including the subsampling matrices.

Parameters

- **cov** (ndarray) – Covariance matrix with possible diagonal only elements
- **zero_threshold** (float) – The near-zero threshold
- **return_subsampled_arr** (bool) – Set to *True* if one wants the split subsampled covariances, otherwise function will just return the subsampling operator matrices

Return type

tuple[ndarray, None | ndarray, ndarray, None | ndarray]

Returns

- **d_off_diagonal** (*numpy.ndarray*) – Sampling matrix operator for the off-diagonal rows
- **the_denser_parts** (*numpy.ndarray* | *None*) – A (somewhat denser) subsampled by that matrix Set to *None* if *return_subsampled_arr* is *False*
- **d_diagonal_only** (*numpy.ndarray*) – Sampling matrix operator for the diagonal only rows
- **isolated_diag_vals** (*numpy.ndarray* | *None*) – Vector with diagonal values of those diagonal rows Set to *None* if *return_subsampled_arr* is *False*

```
glomar_gridding.autoregressive_kalman.cov_diagonal.remove_diag_only_rows(cov,  
                                                                           zero_threshold=1e-  
                                                                           06)
```

Subsampled the covariance matrix by the removal of diagonal-only elements. “Diagonal-only” elements are defined according to the `zero_threshold`

Parameters

- **cov** (ndarray) – covariance matrix with possible diagonal only elements
- **zero_threshold** (float) – The near-zero threshold

Return type

tuple[ndarray, ndarray]

Returns

- **new_cov_arr** (numpy.ndarray) – A new covariance without those diagonal-only elements
- **D** (numpy.ndarray) – The subsampling matrix that generates new_cov_arr

```
glomar_gridding.autoregressive_kalman.cov_diagonal.restore_diag_only_rows(trimmed_cov_arr,  
                                                                           D,  
                                                                           diag_fill_value=1.2)
```

Re-expanding subsampled covariance matrix prior to the removal of diagonal-only elements. Diagonal elements are filled with `diag_fill_value`.

Parameters

- **trimmed_cov_arr** (ndarray) – A trimmed numpy array that needs expanded
- **D** (ndarray | sparray) – The subsampling array that did the original purge (see `remove_diag_only_rows`); there should only be a single 1 or `True` per row (axis=0)
- **diag_fill_value** (float) – The diagonal fill value for restored rows and columns
- **atol** (float) – Instead of checking for exact 0s, this is the threshold to decide `diag_fill_value` replacement will occur

Returns

new_cov_arr – A larger (restored) covariance array

Return type

ndarray

12.8 autocorrelation

Compute lagged (lag-1) correlations

Estimation of lag-1 autocorrelation and other metrics for a single time series

Outputs can be organised to be used *AutoregressiveIForecastUncorr* in *forecast_linear_ar1*. It is not the most sophisticated and fancy computation (in a matter of fact very elementary).

Uses `numpy.vectorize` to extend the computation to multiple time series, but each time series are treated separately (i.e. no correlation between time series). Function applied along each row.

`get_anomalies`: `n -> n`: `vectorize` for an array with `m` rows `n` columns gives `m` rows and `n` columns

`get_obs_variance` etc: `n -> scalar`: `vectorize` for an array with `m` rows give `m` scalars.

The same behavior can be gotten using `np.apply_along_axis`

`glomar_gridding.autocorrelation.compute_autocorrelation.compute_lag_ts_metrics_from_da(da, lag=1)`

Compute the following (local) time series metrics - nth-lagged autocorrelation - sample variance

For lag > 1, the autoregressive regression coefficient depends on lagged correlation for shorter lags. See Yule-Walker equations for details.

This computation ignores spatial correlation. Time series metrics are local.

Parameters

- **da** (DataArray) – A 2+-dimension data array that presumes time is on axis-0; by climate data conventions, usually only ensemble number/realization should have axis number lower than time. e.g. Ens, time, z, y, x ... etc. Reorder your dimensions first!
- **lag** (int) – The lagged autocorrelation needed; if lag > 1, then your autoregressive model coefficients are different, see Yule-Walker equations for details.

Return type

tuple[DataArray, DataArray]

Returns

- **ar** (*xarray.DataArray*) – Data array with lagged autoregression Unit is always 1 For *da* with dimensions (t, x_1, ..., x_n) *ar* will have shape (x_1, ..., x_n)
- **variance** (*xarray.DataArray*) – Data array with sample variance Unit is (*unit_of_da*)**2 (variance, not standard deviation) Same shape as *ar*

`glomar_gridding.autocorrelation.compute_autocorrelation.get_anomalies(vec)`

Compute anomalies, 0-mean centering data samples.

Samples are NOT standardized by their standard deviation.

Parameters

vec (ndarray) – Array that needs be 0-averaged

Returns

anomalies – 0-mean-ed values

Return type

ndarray

`glomar_gridding.autocorrelation.compute_autocorrelation.get_anomalies_vectorize(arr)`

Vectorised *get_anomalies*

Similar precautions with *get_obs_variance_vectorize*: this assumes the convention used in climate data, *scikit-learn*, multivariate random num generator etc. (transpose *arr* automatically so it works properly with *np.vectorize*). See *get_obs_variance_vectorize* docstring

Parameters

arr (ndarray) – Array that needs be 0-averaged for each ROW; shape (*N*, *M*) (*N* samples, *M* features/independent variable)

Returns

anomalies – 0-averaged values along each ROW; shape (*N*, *M*) (*N* samples for *M* features)

Return type

numpy.ndarray

`glomar_gridding.autocorrelation.compute_autocorrelation.get_auto_corr(vec, lag=1)`

Compute lag-n autocorrelation

This is the “weight” for spatially uncorrelated autoregressive model.

Degrees of freedom: - This should be n minus lag minus 1 - For lag-1, there are n - 1 samples with an additional -1 from Bessel’s correction

Parameters

- **vec** (ndarray) – 1D vector that needs autoregression correlation
- **lag** (int) – order (lag) of AR

Return type

tuple[float, float, float, float]

Returns

- **ar_lag** (float) – The *lag*-th lagged correlation; this is autocovariance_n divided by (non-lagged) variance
- **autocovariance_lag** (float) – Same as in *ar_lag* but not normalised
- **variance** (float) – Standard deviation squared
- **sigma2_eps** (float) – Standard error of the autoregressive model associated with the same autocorrelation

`glomar_gridding.autocorrelation.compute_autocorrelation.get_auto_corr_vectorize(arr, lag=1)`

Vectorised `get_auto_corr`

Similar precautions with `get_obs_variance_vectorize`: this assumes the convention used in climate data, scikit-learn, multivariate random num generator etc. (transpose *arr* automatically so it works properly with `np.vectorize`); see `get_obs_variance_vectorize` docstring.

Parameters

- **arr** (ndarray) – A multi-dimension array that needs autoregression correlation; shape (*N*, *M*) (*N* samples, *M* features/independent variable)
- **lag** (int) – order (lag) of AR

Returns

- **ar_lag** (float) – The *n*-th lagged correlation; this is autocovariance_n divided by (non-lagged) variance; shape (*M*,) (*M* features)
- **autocovariance_lag** (float) – Same as in *ar_lag* but not standardised; shape (*M*,) (*M* features)
- **variance** (float) – Standard deviation squared; shape (*M*,) (*M* features)
- **sigma2_eps** (float) – Standard error of the autoregressive model associated with the same autocorrelation; shape (*M*,) (*M* features)

`glomar_gridding.autocorrelation.compute_autocorrelation.get_obs_variance(vec, ddof=1, compute_anomalies=True)`

Compute observed variance (standard deviation squared). Bessel correction applied by default.

Parameters

- **vec** (ndarray) – 1D vector that needs observed variance
- **ddof** (int) – Degrees of freedom correction, default 1 (Bessel’s correction); this is the similar to kwargs used in `numpy.cov` (*bias* and *ddof*)

- **compute_anomalies** (*bool*) – If *True*, the mean will be removed from *vec*

Returns

variance – standard deviation squared

Return type

float

```
glomar_gridding.autocorrelation.compute_autocorrelation.get_obs_variance_vectorize(arr,
                                                                                    ddof=1,
                                                                                    com-
                                                                                    pute_anomalies=True)
```

Running `get_obs_variance` for a multi-dimension array

Warning: naive `np.vectorize` may result operations on WRONG axis.

Nearly all climate data, functions like `scipy.stats.multivariate_norm` (aka scikit-learn convention) will need to be TRANSPOSED for the function to work as normally intended, which is automatically performed here. (aka (N, M) – N samples, M features/independent variable); this function (and its companion ones) ASSUMES this will be the shape of *arr*.

Try this function with 1D vector:

`vec.reshape(-1, 1)` (scikit-learn convention for single feature) gives (sample_size, 1) and WORKS

`vec.reshape(1, -1)` (scikit-learn convention for single sample) gives (1, sample_size) and DOES NOT WORK

Parameters

- **arr** (ndarray) – A multidimensional array that needs a vector of variance for each COLUMN; shape (N, M) (N samples, M features/independent variable)
- **ddof** (int) – Degrees of freedom correction, default 1 (Bessel's correction); this is the similar to kwargs used in `numpy.cov` (*bias* and *ddof*); this is the same keyword for `get_obs_variance`
- **compute_anomalies** (*bool*) – If *True*, the mean will be removed from *vec*; this is the same keyword for `get_obs_variance`

Returns

variance – standard deviation squared; shape (M,) (M features)

Return type

ndarray

```
class glomar_gridding.autocorrelation.LassoEstimate_AR1(data_sample, standardise=True, lam=0.01,
                                                         lasso_selection='random',
                                                         out_of_sample_residues=True,
                                                         hold_out_ratio=0.5, dtype=<class
                                                         'numpy.float32'>, lasso_kws=None)
```

A method to estimate of weights for Autoregressive1ForecastVector using regularized line-by-line / grid point-by-grid point (aka seemingly (un)related regression SUR) Lasso regression.

This method is significantly slower than using `compute_autocorrelation`, but handle spatial autocorrelation more properly.

This method is not directly compatible with covariance estimated using the ellipse method. No correlation-must-decrease-with-distance constrains when estimating the Lasso weights.

See: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Lasso.html as well. This is just a wrapper for that but will loop through *data_sample* over many rows.

Parameters

- **data_sample** (ndarray) – A 2D data sample array; different columns for different time; different rows for different grid points/locations; dtype *float*; shape (T, N) following netCDF/GRIB convention
- **standardise** (bool) – Should data_sample be standardised automatically (i.e. $A^* = A - E(A) / \sigma(A)$)? Standardise if *True*
- **lam** (float) – The lambda hyperparameter for Lasso regression; it is called *alpha* in *scikit-learn*, but it is more common to be called it *lambda* in the literature. Tuning of it is usually done by orders of magnitude 0.0001, 0.001, 0.01, 0.1, 1... etc (default of sklearn is 1) If lam=0, it becomes OLS; scikit-learn will probably raise warnings if one does that.
- **lasso_selection** (Literal['cyclic', 'random']) – *selection* in *scikit-learn lasso* class, see *scikit-learn* documentation
- **out_of_sample_residues** (bool) – Should the residues of the fit be estimated out of sample When set to *True*, data_sample is split by hold_out_ratio and data not used for training will be used to estimate the residues. Otherwise, residues are estimated in-sample.
- **hold_out_ratio** (float) – The ratio of data be used that is withheld for cross-validation. It is only used if *out_of_sample_residues* is *True*
- **dtype** (type) – *numpy* or python native float types that are used to store the outputs
- **lasso_kws** (dict | None) – Dict with (other) keyword arguments for *sklearn.linear_model.Lasso*; Do not include *alpha* or *selection* in lasso_kws; use *lasso_selection* and *lam* instead. If *alpha* or *selection* are in lasso_kws, they will be removed.

estimate_errcov(dof_adj=0)

Estimate error covariance

$$E = R @ R^T / (T^* - dof_{adj})$$

Parameters

dof_adj (float | int) – Additional degree of freedom adjustment for each line of seeming-unrelated regression

errcov

Error covariance for the regression

Type

np.ndarray

expand_coefficients(D, dtype=<class 'numpy.float32'>, fill_value_w=0.6, fill_value_errcov=0.36)

Expands coefficients according to the subsampling matrix D. D should be array of integers of 0s and 1s, and will be converted to sparse if not sparse already.

self.coefficient matrix will be converted to sparse if it has not been converted.

Sets attribute *expanded* to *True*.

Parameters

- **D** (ndarray | sparray) – A 2D array indicating unmasked location, should normally be sparse; If not an instance *sp.sparse.sparray*, it will converted to one; dtype *uint8* or *bool*; shape (N, M)
- **dtype** (type) – *numpy* or python native *float* types that are used for expanded array. Given the nature of this type of analysis double precision floats (float64) are overkill. Defaults to single precision floats (float32).
- **fill_value_w** (float) – The diagonal fill value for the weights; default 0.6

D

Adds *D* as an attribute, store in sparse if not already sparse

Type

scipy.sparse.sparray

fit(*convert_to_sparse=False*)

Fit the Lasso regression

The complete fitting procedure following line-by-line estimate of the regression coefficients based on seemingly (un)related regression approach

Parameters

convert_to_sparse (bool) – If True, runs `make_coeff_sparse` automatically, which converts coefficients to a scipy sparse matrix. Default False.

coefficients

The weights of the fits; shape (*N*, *N*)

Type

numpy.ndarray

residues

The residues of the fit; length of residue time series depends on *out_of_sample_residues* and *hold_out_ratio*; shape (*N*, *T**) in which $T^* \leq T$

Type

numpy.ndarray

expanded

A bool flag that indicates if coefficients have been expanded by the `expand_coefficients` method Set to *False*; it will only be set to *True* if method *expand_coefficients* is called later.

Type

bool

make_coeff_sparse()

Convert the weights to scipy sparse format

split_holdout()

Split data into training part and holdout cross-validation bits if *self.out_of_sample_residues* is *False*, no split will be used, and the full obs dataset will be used for training and residue estimation.

to_xarray_da_coefficients(*lats, lons, name='weights', attrs={'description': 'weights', 'units': '1'}*)

Convert expanded coefficients into a xarray DataArray instance

Example attrs for netCDF:

Parameters

- **lats** (ndarray) – Array of latitude values, such as `np.linspace(-19.5, 19.5, 40)`
- **lons** (ndarray) – Array of longitude values, such as `np.linspace(-179.5, 179.5, 360)`
- **name** (str) – Name of the variable
- **attrs** (dict | None) – Attributes to be added the *xarray.DataArray*

Returns

da – `self.coefficients` as a DataArray that has metadata and written easily to netCDF file

Return type

DataArray

to_xarray_da_errcov(*lats, lons, name='error_covariance', attrs=None*)

Convert estimated error covariance into a xarray DataArray instance

Example attrs:

Parameters

- **lats** (ndarray) – Array of latitude values, such as *np.linspace(-19.5, 19.5, 40)*
- **lons** (ndarray) – Array of longitude values, such as *np.linspace(-179.5, 179.5, 360)*
- **name** (str) – Name of the variable
- **attrs** (dict | None) – Attributes to be added the *xarray.DataArray*

Returns

da – errcov as a DataArray that has metadata and written easily to netCDF

Return type

xarray.DataArray

unexpand_coefficients()

Reverses `expand_coefficients`, by the subsampling matrix `D`. Removes the `D` attribute.

MISCELLANEOUS MODULES

13.1 Climatologies

Functions for mapping climatologies and computing anomalies

```
glommar_gridding.climatology.join_climatology_by_doy(obs_df, climatology_365, lat_col='lat',  
                                                    lon_col='lon', date_col='date', var_col='sst',  
                                                    clim_lat='latitude', clim_lon='longitude',  
                                                    clim_doy='doy', clim_var='climatology',  
                                                    temp_from_kelvin=True)
```

Merge a climatology from an xarray.DataArray into a polars.DataFrame using the **day of year** value and position.

This function accounts for leap years by taking the average of the climatology values for 28th Feb and 1st March for observations that were made on the 29th of Feb.

The climatology is merged into the DataFrame and anomaly values are computed.

Parameters

- **obs_df** (DataFrame) – Observational DataFrame.
- **climatology_365** (Dataset) – DataArray containing daily climatology values (for 365 days).
- **lat_col** (str) – Name of the latitude column in the observational DataFrame.
- **lon_col** (str) – Name of the longitude column in the observational DataFrame.
- **date_col** (str) – Name of the datetime column in the observational DataFrame. Day of year values are computed from this value.
- **var_col** (str) – Name of the variable column in the observational DataFrame. The merged climatology names will have this name prefixed to “_climatology”, the anomaly values will have this name prefixed to “_anomaly”.
- **clim_lat** (str) – Name of the latitude coordinate in the climatology DataArray.
- **clim_lon** (str) – Name of the longitude coordinate in the climatology DataArray.
- **clim_doy** (str) – Name of the day of year coordinate in the climatology DataArray.
- **clim_var** (str) – Name of the climatology variable in the climatology DataArray.
- **temp_from_kelvin** (bool) – Optionally adjust the climatology from Kelvin to Celsius if the variable is a temperature.

Returns

obs_df – With the climatology merged and anomaly computed. The new columns are “_climatology” and “_anomaly” prefixed by the *var_col* value respectively.

Return type

DataFrame

`glomar_gridding.climatology.read_climatology(clim_path, min_lat=-90, max_lat=90, min_lon=-180, max_lon=180, lat_var='lat', lon_var='lon', **kwargs)`

Load a climatology dataset from a netCDF file.

Parameters

- **clim_path** (str) – Path to the climatology file. Can contain format blocks to be replaced by the values passed to *kwargs*.
- **min_lat** (float) – Minimum latitude to load.
- **max_lat** (float) – Maximum latitude to load.
- **min_lon** (float) – Minimum longitude to load.
- **max_lon** (float) – Maximum longitude to load.
- **lat_var** (str) – Name of the latitude variable.
- **lon_var** (str) – Name of the longitude variable.
- ****kwargs** – Replacement values for the climatology path.

Returns

clim_ds – Containing the climatology bounded by the min/max arguments provided.

Return type

Dataset

13.2 Masking

Functions for applying masks to grids and DataFrames

`glomar_gridding.mask.get_mask_idx(mask, mask_val=nan, masked=True)`

Get the 1d indices of masked values from a mask array.

Parameters

- **mask** (DataArray) – The mask array, containing values indicated a masked value.
- **mask_val** (Any) – The value that indicates the position should be masked.
- **masked** (bool) – Return indices where values in the mask array equal this value. If set to False it will return indices where values are not equal to the mask value. Can be used to get unmasked indices if this value is set to False.

Return type

ndarray

Examples

```
>>> data = np.random.rand(4, 4)
>>> data[data > 0.65] = np.nan
>>> mask = xr.DataArray(data)
>>> get_mask_idx(mask)
array([[1],
       [3],
       [4],
```

(continues on next page)

(continued from previous page)

```
[5],
[8]])
```

`glomar_gridding.mask.mask_array(grid, mask, masked_value=nan, mask_value=True)`

Apply a mask to a DataArray.

The grid and mask must already align for this function to work. An error will be raised if the coordinate systems cannot be aligned.

Parameters

- **grid** (DataArray) – Observational DataArray to be masked by positions in the mask DataArray.
- **mask** (DataArray) – Array containing values used to mask the observational DataFrame.
- **masked_value** (Any) – Value indicating masked values in the DataArray.
- **mask_value** (Any) – Value to set masked values to in the observational DataFrame.

Returns

grid – Input xarray.DataArray with the variable masked by the mask DataArray.

Return type

DataArray

`glomar_gridding.mask.mask_dataset(dataset, mask, varnames, masked_value=nan, mask_value=True)`

Apply a mask to a DataSet.

The grid and mask must already align for this function to work. An error will be raised if the coordinate systems cannot be aligned.

Parameters

- **dataset** (Dataset) – Observational Dataset to be masked by positions in the mask DataArray.
- **mask** (DataArray) – Array containing values used to mask the observational DataFrame.
- **varnames** (str | list[str]) – A list containing the names of variables in the observational Dataset to apply the mask to.
- **masked_value** (Any) – Value indicating masked values in the DataArray.
- **mask_value** (Any) – Value to set masked values to in the observational DataFrame.

Returns

grid – Input xarray.Dataset with the variables masked by the mask DataArray.

Return type

Dataset

`glomar_gridding.mask.mask_from_obs_array(obs, datetime_idx)`

Infer a mask from an input array. Mask values are those where all values are NaN along the time dimension.

An example use-case would be to infer land-points from a SST data array.

Parameters

- **obs** (ndarray | DataArray) – Array containing the observation values. Records that are `numpy.nan` will count towards the mask, if all values in the datetime dimension are `numpy.nan`.

- **datetime_idx** (int) – The index of the datetime, or grouping, dimension. If all records at a point along this dimension are NaN then this point will be masked.

Returns

mask – A boolean array with dimension reduced along the datetime dimension. A True value indicates that all values along the datetime dimension for this index are numpy.nan and are masked.

Return type

ndarray | DataArray

```
glomar_gridding.mask.mask_from_obs_frame(obs, coords, value_col, datetime_col=None, grid=None,
                                         grid_coords=None)
```

Compute a mask from observations and an optional output grid..

Positions defined by the “coords” values that do not have any observations, at any datetime value in the “datetime_col”, for the “value_col” field are masked.

An example use-case would be to identify land positions from sst records.

If a grid is supplied, the observations are mapped to the grid and any positions from the grid that do not contain observations.

If no grid is supplied, then it is assumed that the observation frame represents the full grid, and any positions without observations are included with null values in the value_col.

Parameters

- **obs** (DataFrame) – DataFrame containing observations over space and time. The values in the “value_col” field will be used to define the mask.
- **coords** (str | list[str]) – A list of columns containing the coordinates used to define the mask. For example [“lat”, “lon”].
- **value_col** (str) – Name of the column containing values from which the mask will be defined.
- **datetime_col** (str | None) – Name of the datetime column. Any positions that contain no records at any datetime value are masked.
- **grid** (DataArray | None) – Optional grid, used to map observations so that empty positions can be identified. If not supplied, it is assumed that the observations frame contains the full grid, and includes nulls in empty positions.
- **grid_coords** (str | list[str] | None) – Optional grid coordinate names. Must be set if grid is set.

Return type

DataFrame

Returns

- *polars.DataFrame* containing coordinate columns and a Boolean “mask” column
- *indicating positions that contain no observations and would be a mask value.*

```
glomar_gridding.mask.mask_observations(obs, mask, varnames, masked_value=nan, mask_value=True,
                                       obs_coords=['lat', 'lon'], mask_coords=['latitude', 'longitude'],
                                       align_to_mask=False, drop=False,
                                       mask_grid_prefix='_mask_grid_')
```

Mask observations in a DataFrame subject to a mask DataArray.

Parameters

- **obs** (DataFrame) – Observational DataFrame to be masked by positions in the mask DataArray.
- **mask** (DataArray) – Array containing values used to mask the observational DataFrame.
- **varnames** (str | list[str]) – Columns in the observational DataFrame to apply the mask to.
- **masked_value** (Any) – Value indicating masked values in the DataArray.
- **mask_value** (Any) – Value to set masked values to in the observational DataFrame.
- **obs_coords** (list[str]) – A list of coordinate names in the observational DataFrame. Used to map the mask DataArray to the observational DataFrame. The order must align with the coordinates of the mask DataArray.
- **mask_coords** (list[str]) – A list of coordinate names in the mask DataArray. These coordinates are mapped onto the observational DataFrame in order to apply the mask. The ordering of the coordinate names in this list must match those in the obs_coords list.
- **align_to_mask** (bool) – Optionally align the observational DataFrame to the mask DataArray. This essentially sets the mask's grid as the output grid for interpolation.
- **drop** (bool) – Drop masked values in the observational DataFrame.
- **mask_grid_prefix** (str) – Prefix to use for the mask gridbox index column in the observational DataFrame.

Returns

obs – Input polars.DataFrame containing additional column named by the mask_varname argument, indicating records that are masked. Masked values are dropped if the drop argument is set to True.

Return type

DataFrame

13.3 Distances and Distance Matrices

Functions for calculating distances or distance-based covariance components.

Some functions can be used for computing pairwise-distances, for example via `squareform`. Some functions can be used as a distance function for `glommar_gridding.error_covariance.dist_weights`, accounting for the distance component to an error covariance matrix.

Functions for computing covariance using Matern Tau by Steven Chan (@stchan).

```
glommar_gridding.distances.calculate_distance_matrix(df, dist_func=<function
    haversine_distance_from_frame>,
    **dist_kwargs)
```

Create a distance matrix from a DataFrame containing positional information, typically latitude and longitude, using a distance function.

Available functions are *haversine_distance*, *euclidean_distance*. A custom function can be used, requiring that the function takes the form: `(tuple[float, float], tuple[float, float]) -> float`

Parameters

- **df** (DataFrame) – DataFrame containing latitude and longitude columns indicating the positions between which distances are computed to form the distance matrix
- **dist_func** (Callable) – The function used to calculate the pairwise distances. Functions available for this function are *haversine_distance* and *euclidean_distance*. A custom function

can be based, that takes as input two tuples of positions (computing a single distance value between the pair of positions). (tuple[float, float], tuple[float, float]) -> float

- ****dist_kwargs** – Keyword arguments to pass to the distance function.

Returns

dist – A matrix of pairwise distances.

Return type

ndarray

`glomar_gridding.distances.displacements(lats, lons, lats2=None, lons2=None, delta_x_method=None)`

Calculate east-west and north-south displacement matrices for all pairs of input positions.

The results are not scaled by any radius, this should be performed outside of this function.

Parameters

- **lats** (ndarray) – The latitudes of the positions, should be provided in degrees.
- **lons** (ndarray) – The longitudes of the positions, should be provided in degrees.
- **lats2** (ndarray | None) – The latitudes of the optional second positions, should be provided in degrees.
- **lons2** (ndarray | None) – The longitudes of the optional second positions, should be provided in degrees.
- **delta_x_method** (Optional[Literal['Met_Office', 'Modified_Met_Office']]) – One of “Met_Office” or “Modified_Met_Office”. If set to None, the displacements will be returned in degrees, rather than actual distance values. Set to “Met_Office” to use a cylindrical approximation, set to “Modified_Met_Office” to use an approximation that uses the average of the latitudes to set the horizontal displacement scale.

Return type

tuple[ndarray, ndarray]

Returns

- **disp_y** (*numpy.ndarray*) – The north-south displacements.
- **disp_x** (*numpy.ndarray*) – The east-west displacements.

`glomar_gridding.distances.euclidean_distance(df, lat_coord='lat', lon_coord='lon', to_radians=True, radius=6371.0)`

Calculate the Euclidean distance in kilometers between pairs of lat, lon points on the earth (specified in decimal degrees).

$$d = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2 + (z_2 - z_1)^2}$$

where

$$(x_n, y_n, z_n) = (R \cos(lat) \cos(lon), R \cos(lat) \sin(lon), R \sin(lat))$$

Parameters

- **df** (DataFrame) – DataFrame containing latitude and longitude columns indicating the positions between which distances are computed to form the distance matrix
- **lat_coord** (str) – Name of the column in the input DataFrame containing latitude values.
- **lon_coord** (str) – Name of the column in the input DataFrame containing longitude values.

- **to_radians** (bool) – Convert inputs from degrees to radians. The distance methodology assumes that the input latitude and longitudes are in radians so must be converted if they are in degrees.
- **radius** (float) – The radius of the sphere used for the calculation. Defaults to the radius of the earth in km (6371.0 km).

Returns

dist – The direct pairwise distance between the positions in the input DataFrame through the sphere defined by the radius parameter.

Return type

ndarray

References

<https://math.stackexchange.com/questions/29157/how-do-i-convert-the-distance-between-two-lat-long-points-into-feet-meters>
https://cesar.esa.int/upload/201709/Earth_Coordinates_Booklet.pdf

```
glomar_gridding.distances.haversine_distance_from_frame(df, lat_coord='lat', lon_coord='lon',
                                                         to_radians=True, radius=6371)
```

Calculate the great circle distance in kilometers between pairs of lat, lon points on the earth (specified in decimal degrees).

Parameters

- **df** (DataFrame) – DataFrame containing latitude and longitude columns indicating the positions between which distances are computed to form the distance matrix
- **lat_coord** (str) – Name of the column in the input DataFrame containing latitude values.
- **lon_coord** (str) – Name of the column in the input DataFrame containing longitude values.
- **to_radians** (bool) – Convert inputs from degrees to radians. The distance methodology assumes that the input latitude and longitudes are in radians so must be converted if they are in degrees.
- **radius** (float) – The radius of the sphere used for the calculation. Defaults to the radius of the earth in km (6371.0 km).

Returns

dist – The pairwise haversine distances between the inputs in the DataFrame, on the sphere defined by the radius parameter.

Return type

ndarray

```
glomar_gridding.distances.haversine_gaussian(df, R=6371.0, r=40, s=0.6)
```

Gaussian Haversine Model

Parameters

- **df** (DataFrame) – Observations, required columns are “lat” and “lon” representing latitude and longitude respectively.
- **R** (float) – Radius of the sphere on which Haversine distance is computed. Defaults to radius of earth in km.
- **r** (float) – Gaussian model range parameter
- **s** (float) – Gaussian model scale parameter

Returns

C – Distance matrix for the input positions. Result has been modified using the Gaussian model.

Return type

ndarray

`glomar_gridding.distances.inv_2d(mat)`

Compute the inverse of a 2 x 2 matrix

Return type

ndarray

`glomar_gridding.distances.mahal_dist_func(delta_x, delta_y, Lx, Ly, theta=None)`

Calculate tau from displacements, Lx, Ly, and theta (if it is known). For an array of displacements, for a set of scalar ellipse parameters, Lx, Ly, and theta.

Parameters

- **delta_x** (ndarray) – displacement to remote point as in: $(\text{delta_x}) i + (\text{delta_y}) j$ in old school vector notation
- **delta_y** (ndarray) – displacement to remote point as in: $(\text{delta_x}) i + (\text{delta_y}) j$ in old school vector notation
- **Lx** (float) – Lx, Ly scale (km or degrees)
- **Ly** (float) – Lx, Ly scale (km or degrees)
- **theta** (float | None) – rotation angle in radians

Returns

tau – Mahalanobis distance

Return type

ndarray

`glomar_gridding.distances.radial_dist(lat1, lon1, lat2, lon2)`

Computes a distance matrix of the coordinates using a spherical metric.

Parameters

- **lat1** (float) – latitude of point A
- **lon1** (float) – longitude of point A
- **lat2** (float) – latitude of point B
- **lon2** (float) – longitude of point B

Return type

float

`glomar_gridding.distances.rot_mat(angle)`

Compute a 2d rotation matrix from an angle.

The input angle must be in radians

Return type

ndarray

`glomar_gridding.distances.sigma_rot_func(Lx, Ly, theta)`

Equation 15 in Karspeck et al 2011 and Equation 6 in Paciorek and Schervish 2006, assuming Sigma(Lx, Ly, theta) locally/moving-window invariant or we have already taken the mean (Sigma overbar, PP06 3.1.1)

Lx, Ly - anisotropic variogram length scales theta - angle relative to lines of constant latitude theta should be radians, and the fitting code outputs radians by default

Returns

sigma – 2 x 2 matrix

Return type

ndarray

`glomar_gridding.distances.tau_dist(dE, dN, sigma)`

Eq.15 in Karspeck paper but it is standard formulation to the Mahalanobis distance https://en.wikipedia.org/wiki/Mahalanobis_distance 10.1002/qj.900

Return type

ndarray

`glomar_gridding.distances.tau_dist_from_frame(df)`

Compute the tau/Mahalanobis matrix for all records within a gridbox

Can be used as an input function for `observations.dist_weight`.

Eq.15 in Karspeck paper but it is standard formulation to the Mahalanobis distance https://en.wikipedia.org/wiki/Mahalanobis_distance 10.1002/qj.900

By Steven Chan - @stchan

Parameters

df (DataFrame) – The observational DataFrame, containing positional information for each observation (“lat”, “lon”), gridbox specific positional information (“grid_lat”, “grid_lon”), and ellipse length-scale parameters used for computation of *sigma* (“grid_lx”, “grid_ly”, “grid_theta”).

Returns

tau – A matrix of dimension n x n where n is the number of rows in *df* and is the tau/Mahalanobis distance.

Return type

ndarray

13.4 Covariance Tools and Eigenvalue Clipping

When estimating covariance matrices, using ellipse-based methods for example, the results may not be positive-definite. This can be problematic, for instance the Kriging equations may not be solvable as the inverse matrix cannot be computed, or simulated states cannot be constructed using `glomar_gridding.stochastic.scipy_mv_normal_draw()`. To account for this, *glomar_gridding* includes a few tools for *fixing* these covariance matrices. In general, these methods attempt to coerce the input matrix to be positive-definite by updating the eigenvalues and re-computing the matrix from these updated eigenvalues and the original eigenvectors. The indicators of an invalid covariance matrix include

1. Un-invertible covariance matrices with 0 eigenvalues
2. Covariance matrices with eigenvalues less than zero

This can typically be a consequence of

1. Multicollinearity: but nearly all very large cov matrices will have rounding errors to have this occur
2. Number of spatial points >> length of time series (for ESA monthly pentads: this ratio is about 150)
3. Covariance is estimated using partial data

In most cases, the most likely causes are 2 and 3.

There are a number of methods included in this module. In general, the approach is to adjust the eigenvalues to ensure small or negative eigenvalues are increased to some minimum threshold. The covariance matrix is then re-calculated using these modified eigenvalues and the original eigenvectors.

In general, the recommended approach is Original Clipping, see [*glomar_gridding.covariance_tools.eigenvalue_clip\(\)*](#).

Fixes:

1. Simple clipping - [*glomar_gridding.covariance_tools.simple_clipping\(\)*](#):

Cut off the negative, zero, and small positive eigenvalues; this is method used in statsmodels.stats.correlation_tools but the version here has better thresholds based on the accuracy of the eigenvalues, plus a iterative version which is slower but more stable with big matrices. The iterative version is recommended for SST/MAT covariances.

This is used for SST covariance matrices which have less dominant modes than MAT; it also preserves more noise.

Trace (aka total variance) of the covariance matrix is not conserved, but it is less disruptive than EOF chop off (method 3).

It is more difficult to use for covariance matrices with one large dominant mode because that raises the bar of accuracy of the eigenvalues, which requires clipping off a lot more eigenvectors.

Note, this will adjust all negative eigenvalues to be positive.

2. Original clipping - [*glomar_gridding.covariance_tools.eigenvalue_clip\(\)*](#):

Determine a noise eigenvalue threshold and replace all eigenvalues below using the average of them, preserving the original trace (aka total variance) of the covariance matrix, but this will require a full computation of all eigenvectors, which may be slow and cause memory problems

Note, this will adjust all negative eigenvalues to be positive.

Other methods not implemented here -

a. **shrinkage methods**

<https://scikit-learn.org/stable/modules/covariance.html>

b. **reprojection (aka Higham's method) [Higham]**

https://github.com/mikecroucher/nearest_correlation

<https://nhigham.com/2013/02/13/the-nearest-correlation-matrix/>

Author S. Chan.

Modified by J. Siddons.

`glomar_gridding.covariance_tools.check_symmetric(a, rtol=1e-05, atol=1e-08)`

Helper function for perturb_sym_matrix_2_positive_definite

Return type

bool

`glomar_gridding.covariance_tools.clean_small(matrix, atol=1e-05)`

Set small values (abs(x) < atol) in an matrix to 0

Return type

ndarray

`glomar_gridding.covariance_tools.csum_up_to_val(vals, target, reverse=True, niter=0, csum=0.0)`

Find csum and sample index that target is surpassed. Displays a warning if the target is not exceeded or the input *vals* is empty.

Can provide an initial *niter* and/or *csum* value(s), if working with multiple arrays in an iterative process.

If *reverse* is set, the returned index will be negative and will correspond to the index required for the non-reversed array. Reverse is the default.

Parameters

- **vals** (ndarray) – Vector of values to sum cumulatively.
- **target** (float) – Value for which the cumulative sum must exceed.
- **reverse** (bool) – Reverse the array. The index will be negative.
- **niter** (int) – Initial number of iterations.
- **csum** (float) – Initial cumulative sum value.

Return type

tuple[float, int]

Returns

- **csum** (float) – The cumulative sum at the index when the target has been exceeded.
- **niter** (int) – The index of the value that results in the cumulative sum exceeding the target.

Note

It is actually faster to compute a full cumulative sum with *np.cumsum* and then look for the value that exceeds the target. This is not performed in this function.

Examples

```
>>> vals = np.random.rand(1000)
>>> target = 301.1
>>> csum_up_to_val(vals, target)
```

`glomar_gridding.covariance_tools.eigenvalue_clip(cov, method='explained_variance', **kwargs)`

Denoise symmetric damaged covariance/correlation matrix *cov* by clipping eigenvalues

Explained variance or aspect ratio based threshold Aspect ratios is based on dimensionless parameters (number of independent variable and observation size)

$$q = N/T = \frac{\text{num of independent variable}}{\text{num of observation per independent variable}}$$

Does not give the same results as in *eig_clip*

explained_variance here does not have the same meaning. The trace of a correlation, by definition, equals the number of diagonal elements, which isn't intuitively linked to actual explained variance in climate science sense

This is done by KEEPING the largest explained variance in which (number of basis vectors to be kept) >> (number of rows) In ESA data, keeping 95% variance means keeping top ~15% of the eigenvalues

Parameters

- **cov** (ndarray) – Input covariance matrix to be adjusted to positive definite.

- **method** (Literal['explained_variance', 'Laloux_2000']) – Method used to identify the index of the eigenvalues to clip. If set to “explained_variance” then the sorted eigenvalues below the target variance are *clipped*. If “Laloux_2000” is set, then the method of [Laloux] is used.
- ****kwargs**

Returns

Adjusted covariance matrix.

Return type

ndarray

 See also

`glomar_gridding.covariance_tools.explained_variance_clip()`

 See also

`glomar_gridding.covariance_tools.laloux_clip()`

`glomar_gridding.covariance_tools.explained_variance_clip(cov, target_variance_fraction=0.95)`

Clip all EOFs beyond a certain level of explained variance Starting from the EOF with the largest eigenvalue (explained variance) descending down the eigenvalues until the cumulative sum of the eigenvalues just goes beyond (Trace(cov) x target_variance_fraction)

All the larger eigenvalues up to the above thresholds are left unchanged.

Eigenvalues beyond (smaller positive and negative ones) are clipped instead. If n is the first eigenvalue to be clipped, and there are N eigenvalues, sorted in descending order:

$$(\lambda)_{clipped} = (\lambda_n, \lambda_{n+1}, \dots, \lambda_N)$$

They now get a revised eigenvalue of $E(\lambda_{clipped})$. Checks are in place within `_eigenvalue_clip` to ensure the threshold is suitable.

The threshold itself is subjective, but they are based on making intelligent guesses what may be good enough to perturb the matrix to make it positive (semi) definite. This approach would always give an acceptable and closest stable matrix ([Higham_blog_nearest]) as long as the thresholds are set sensibly.

The automatic method (`laloux_clip`) (sort-of) does it for you, but it only works after the covariance is standardised (or the diagonal of the covariance matrix is constant). It is also an aggressive method which is based on noisiness and probability distribution of eigenvalues. Noting that closeness to the original matrix is our primary goal, not the noisiness of eigenvalues.

- [Jolliffe] recommends 70-90% for truncation
- Graphical guidances (“spectrum of eigenvalues”) are helpful ([Wilks], [Laloux])

Noting that [Wilks] and [Jolliffe] mostly concern with TRUNCATION not clipping. However, clipping is essentially a modified version of truncation but with the explained variance fully conserved, nor we cannot use covariances based on truncated EOFs; it will have 0 eigenvalues, many of them will become negative after floating point errors.

95% is somewhat higher than [Jolliffe] 70-90% guidance. Operational experience with 5x5 monthly data indicates even at 95%, you are only retaining 200-400 eigenvalues out of $36 \times 72 = 2592$ possible eigenvalues (if data is global). There are only ~10-20 negative eigenvalues plus the 2000-ish eigenvalues that sits out of 95%

threshold. The magnitude of the negative values are 2-plus order of magnitude smaller than the largest positive eigenvalues.

Parameters

- **cov** (ndarray) – Input covariance matrix to be adjusted to positive definite.
- **target_variance_fraction** (float) – Starting from the largest eigenvalue and descending, all eigenvalues larger than (the trace of cov x target_variance_fraction) are left unmodified. Eigenvalues beyond are modified.

Returns

Adjusted covariance matrix.

Return type

ndarray

See also

`glomar_gridding.covariance_tools.laloux_clip()`

References

- [Laloux]
- [Jolliffe]
- [Wilks]

`glomar_gridding.covariance_tools.laloux_clip(cov, num_grid_pts=None, num_time_pts=40)`

Estimate the largest eigenvalue that one will get from covariance or correlation matrices that are generated by random uncorrelated matrices aka the noise level of the eigenvalues.

Eq 0.3 in [Laloux] says eigenvalues of covariances generated by uncorrelated random vectors with constant variance have a max of:

$$\lambda_{max} = \sigma^2(1 + Q + 2\sqrt{Q})$$

in which:

$$Q = \frac{\text{num of features}}{\text{length of each feature}}$$

One can make a simple Monte-Carlo simulation to check that (a few lines of Python code will do). With variable (but realistic observed) variance, equation above is not exactly correct. However, Monte-Carlo simulations of λ_{max} show differences are not large; it becomes a somewhat larger, determined by the larger variances of the variable variances, but the estimate is correct up to the first significant digit.

Hence λ_{max} can be thought of a floor for non-noise EOFs, as this is the largest eigenvalue that can be generated by random uncorrelated data, and is only a function of Q (aka aspect ratio of the number of features with the length of each feature vector).

Any eigenvalues smaller than λ_{max} (including negative ones) are essentially noise and can possibly be trimmed.

This approach is aggressive. For climate sciences, Q can be large; for example if number of features are on the order of 1000-2000 (a global 5x5 grid land-only, sea-only, or of any terrain (up to 2592)) with 30-40 data points per feature (one value per year), the threshold can be reached within first 3-5 EOFs. Q would only worsen if number of features increases while number of values per feature decreases (aka higher resolution remote sensing data).

The explained variance approach (default), even if set to a generous value (say 70%; see [Wilks] chapter 11) would still keep a lot more EOFs. Noting that just because some (many) EOFs are noisy or degenerate, they are still part of the original covariance; they are actual bits of our best estimate to the data. It is desirable to minimise the perturbation applied to the covariance, keeping it as close as it is original; noise and degeneracy may be interesting by their own right (say in the physical interpret the EOFs – a contentious issue), but that is not our primary concern here. See [North] Section 5 and [Higham_blog_nearest] for further discussion.

Above requires q and σ be the same for all features/independent variables. A way to get to around that is standardise the covariance into correlation and to apply the trimming to the correlation. The trimmed correlation can then converted back to the covariance by putting the correct standard deviations back into the diagonal. This effectively forces/averages the variance to constant.

Parameters

- **cov** (ndarray) – Input covariance matrix to be adjusted to positive definite.
- **num_grid_pts** (int | None) – Number of spatial grid points for covariance, this should usually be the shape of the covariance matrix. If unset, this will default to the size of the input covariance matrix (`cov.shape[0]`).
- **num_time_pts** (int) – Length of the time series that is behind the covariance generation Default 40 (aka 40 Jans from 1981-2020), since this is what used originally when covariances are generated for modern satellite era data (early 1980-ish to 2020-ish). It is important this value to be set correctly.

Returns

Adjusted covariance matrix.

Return type

ndarray

➡ See also

`glomar_gridding.covariance_tools.explained_variance_clip()`

References

- [Laloux]
- [Bun]

`glomar_gridding.covariance_tools.perturb_cov_to_positive_definite(cov, threshold=1e-15)`

Force an estimated covariance matrix to be positive definite using the eigenvalue clipping with `statsmodels.stats.correlation_tools.cov_nearest` function.

Deprecated in favour of `glomar_gridding.covariance_tools.simple_clipping()`.

Parameters

- **cov** (ndarray) – The estimated covariance matrix that is not positive definite.
- **threshold** (Union[float, Literal['auto']]) – Eigenvalues below this value are set to 0. If the input is 'auto' then the value is determined using the floating-point precision and magnitude of the largest eigenvalues.

Returns

cov_adj – Adjusted covariance matrix

Return type

ndarray

 See also`glomar_gridding.covariance_tool.simple_clipping()`**Notes**

Other methods:

- <https://nhigham.com/2021/02/16/diagonally-perturbing-a-symmetric-matrix-to-make-it-positive-definite/>
- <https://nhigham.com/2013/02/13/the-nearest-correlation-matrix/>
- <https://academic.oup.com/imajna/article/22/3/329/708688>

`glomar_gridding.covariance_tools.simple_clipping(cov, threshold='auto', method='iterative')`

A modified version of: https://www.statsmodels.org/dev/generated/statsmodels.stats.correlation_tools.corr_nearest.html

Force an estimated covariance matrix to be positive definite using the eigenvalue clipping with `statsmodels.stats.correlation_tools.cov_nearest` function.

This is appropriate for covariance matrices which have less dominant modes; it also preserves more noise.

Trace (aka total variance) of the covariance matrix is not conserved, but it is less disruptive than EOF chop off.

Parameters

- **cov** (*numpy.ndarray*) – The estimated covariance matrix that is not positive definite.
- **threshold** (*Union[float, Literal['auto', 'statsmodels_default']*) – Eigenvalues below this value are set to 0. If the input is 'auto' then the value is determined using the floating-point precision and magnitude of the largest eigenvalues.

Return type

`tuple[ndarray, dict[str, Any]]`

Returns

- **cov_adj** (*numpy.ndarray*) – Adjusted covariance matrix
- **summary_dict** (*dict[str, Any]*) – A dictionary containing a summary of the input and results with the following keys:
 - "threshold"
 - "smallest_eigv"
 - "determinant"
 - "total_variance"

 See also`statsmodels.stats.correlation_tools.cov_nearest()`**Notes**

Other methods:

- <https://nhigham.com/2021/02/16/diagonally-perturbing-a-symmetric-matrix-to-make-it-positive-definite/>
- <https://nhigham.com/2013/02/13/the-nearest-correlation-matrix/>

- <https://academic.oup.com/imagina/article/22/3/329/708688>

`glomar_gridding.covariance_tools.validate_covariance(cov, sym_atol=1e-05)`

Validate that a covariance is square and symmetric.

If the input is symmetric return itself. If it is close to symmetric (within ‘sym_atol’) then return $(cov + cov.T) / 2$. Otherwise raise an error.

Parameters

- **cov** (ndarray) – Covariance Matrix
- **sym_atol** (float) – absolute tolerance to check symmetry of cov

Returns

cov – Symmetric square covariance matrix

Return type

ndarray

13.5 IO

`glomar_gridding` includes functionality for loading datasets or arrays from *netCDF* files using python format strings. This can be useful for loading pre-computed inputs for the Kriging process, for example covariance matrices or observations. The allowance for passing a string containing format components (e.g. python t-string) allows for dynamic configuration if processing a series of monthly inputs for example.

Also included is a function for recursively getting sub-keys from a python *dict* style object. This can be useful for working with *yaml* formatting configuration files for instance.

`glomar_gridding.io.get_recurse(config, *keys, default=None)`

Recursively get sub keys from a python dict object.

If a dictionary object contains keys whose values are themselves dictionaries, get a value from a sub dictionary by specifying the key-path to get to the desired value.

Equivalent to:

```
config[keys[0]][keys[1]]...[keys[n]]
```

Or:

```
config.get(keys[0]).get(keys[1])...get(keys[n])
```

Parameters

- **config** (dict) – The layered dictionary containing sub dictionaries.
- ***keys** – The sequence of keys to recurse through to get the final value. If any key in the sequence is not found, or is not a dictionary (and is not the final key), then the default value is returned.
- **default** (Any) – The default value, returned if the sequence of keys cannot be completed or the final key is not present.

Return type

Any

Returns

- *The value associated with the final key, or the default value if the final*

- *key cannot be reached.*

`glomar_gridding.io.load_array(path, var='covariance', **kwargs)`

Load an xarray.DataArray from a netCDF file. Can input a filename or a string to format with keyword arguments.

Parameters

- **path** (`str`) – Full filename (including path), or filename with replacements using `str.format` with named replacements. For example:
`/path/to/global_covariance_{month:02d}.nc`
- **var** (`str`) – Name of the variable to select from the input file
- ****kwargs** – Keywords arguments matching the replacements in the input path.

Returns

arr – An array containing the values of the variable specified by `var`

Return type

DataArray

`glomar_gridding.io.load_dataset(path, **kwargs)`

Load an xarray.Dataset from a netCDF file. Can input a filename or a string to format with keyword arguments.

Parameters

- **path** (`str`) – Full filename (including path), or filename with replacements using `str.format` with named replacements. For example:
`/path/to/global_covariance_{month:02d}.nc`
- ****kwargs** – Keywords arguments matching the replacements in the input path.

Returns

arr – The netcdf dataset as an xarray.Dataset.

Return type

Dataset

13.6 Utilities

Utility functions for *glomar_gridding*

exception `glomar_gridding.utils.ColumnNotFoundError`

Error class for Column Not Being Found

class `glomar_gridding.utils.MonthName(*values)`

Name of month from int

`glomar_gridding.utils.add_empty_layers(nc_variables, timestamps, shape)`

Add empty layers to a netcdf file. This adds a layer of zeros to the netCDF file.

Parameters

- **nc_variables** (`Iterable[Variable] | Variable`) – Name(s) of the variables to add empty layers to
- **timestamps** (`Iterable[int] | int`) – Indices to add empty layers
- **shape** (`tuple[int, int]`) – Shape of the layer to add

Return type

None

`glommar_gridding.utils.adjust_small_negative(mat, atol=1e-08)`

Adjusts small negative values below an absolute tolerance value in a matrix to 0.

Raises a warning if any small negative values are detected.

Parameters

- **mat** (ndarray) – Squared uncertainty associated with chosen kriging method Derived from the diagonal of the matrix
- **atol** (float) – Absolute tolerance value.

Returns

With negative values below an absolute tolerance replaced with 0.

Return type

ndarray

Examples

```
>>> arr = np.array([[1, -1e-10], [-1e-10, 1]])
>>> adjust_small_negative(arr, atol=1e-8)
array([[1., 0.],
       [0., 1.]])
```

`glommar_gridding.utils.batched(iterable, n, *(Keyword-only parameters separator (PEP 3102)), strict=False)`

Implementation of itertools.batched for use if python version is < 3.12.

Examples

```
>>> list(batched("ABCDEFGH", 3))
[("A", "B", "C"), ("D", "E", "F"), ("G", )]
```

`glommar_gridding.utils.check_cols(df, cols)`

Check that all columns in a list of columns are in a DataFrame.

Parameters

- **df** (DataFrame)
- **cols** (list[str]) – List of columns to check for in *df*

Raises

ColumnNotFoundError – If any columns in *cols* are not present in *df*. The missing columns are displayed in the error message.

Return type

None

`glommar_gridding.utils.check_shape(X, dim=2)`

Check X to be 2D This should follow the convention that Each column represent different xy grid points Each row represents another time Proper netCDF files should have the correct order of the dimensions: (E), T, (Z), Y, X So reshaping (T, Y, X) >> T, YX will be sufficient

`glommar_gridding.utils.check_sq_matrix(arr)`

Check if arr is 2D and square

Return type

bool

glomar_gridding.utils.**cor_2_cov**(*cor*, *variances*, *rounding*=None)

Compute covariance matrix from correlation matrix and variances

Parameters

- **cor** (ndarray) – Correlation Matrix
- **variances** (ndarray) – Variances to scale the correlation matrix.
- **rounding** (int | None) – round the values of the output

Return type

ndarray

glomar_gridding.utils.**cov_2_cor**(*cov*, *rounding*=None)Normalises the covariance matrices within the class instance and return correlation matrices <https://gist.github.com/wiso/ce2a9919ded228838703c1c7c7dad13b>**Parameters**

- **cov** (ndarray) – Covariance matrix
- **rounding** (int | None) – round the values of the output

Return type

ndarray

glomar_gridding.utils.**days_since_by_month**(*year*, *day*)Get the number of days since *year*-01-*day* for each month. This is used to set the time values in a netCDF file where temporal resolution is monthly and the units are days since some date.**Parameters**

- **year** (int) – Get a value for each month in this year.
- **day** (int) – Day of month for each returned datetime value in the sequence.

ReturnsContaining 12 values, one for each month in the year containing the number of days since *year*-01-*day*.**Return type**

ndarray

Examples

```
>>> days_since_by_month(1988, 14)
array([ 0, 31, 60, 91, 121, 152, 182, 213, 244, 274, 305, 335])
```

glomar_gridding.utils.**deg_to_km**(*deg*)

Convert degree latitude change to km

Parameters**deg** (float) – The difference in latitude in degrees**Returns**

The latitude difference in kilometers

Return type

float

`glomar_gridding.utils.deg_to_nm(deg)`

Convert degree latitude change to nautical miles

Parameters

deg (float) – The difference in latitude in degrees

Returns

The latitude difference in nautical miles

Return type

float

`glomar_gridding.utils.filter_bounds(df, bounds, bound_cols, closed='left')`

Filter a polars DataFrame based on a set of lower and upper bounds.

Parameters

- **df** (DataFrame) – The data to be filtered by the bounds
- **bounds** (list[tuple[float, float]]) – A list of tuples containing lower and upper bounds for a column
- **bound_cols** (list[str]) – A list of column names to be filtered by the bounds, the length of the bounds list must equal the length of the bound_cols list.
- **closed** (Union[Literal['left', 'right', 'both', 'none'], list[Literal['left', 'right', 'both', 'none']]]) – One of “both”, “left”, “right”, “none” indicating the closedness of the bounds. If the input is a single instance then all bounds will have that closedness. If it is a list of closed values then its length must match the length of the bounds list.

Returns

DataFrame filtered by the positional bounds

Return type

DataFrame

`glomar_gridding.utils.find_nearest(array, values)`

Get the indices and values from an array that are closest to the input values.

A single index, value pair is returned for each look-up value in the values list.

Parameters

- **array** (ndarray) – The array to search for nearest values.
- **values** (ndarray) – The values to look-up in the array.

Return type

tuple[list[int], ndarray]

Returns

- **idx_list** (list[int]) – The indices of nearest values
- **array_values_list** (list) – The list of values in array that are closest to the input values.

Examples

```
>>> array = [1.0, 2.5, 2.7, 2.1, 4.5]
>>> tests = [1.1, 4.4, 2.2]
>>> find_nearest(array, tests)
([np.int64(0), np.int64(4), np.int64(3)], array([1. , 4.5, 2.1]))
```


`glomar_gridding.utils.get_date_index(year, month, start_year)`

Get the index of a given year-month in a monthly sequence of dates starting from month 1 in a specific start year

Parameters

- **year** (int) – The year for the date to find the index of.
- **month** (int) – The month for the date to find the index of.
- **start_year** (int) – The start year of the date series, the result assumes that the date time series starts in the first month of this year.

Returns

index – The index of the input date in the monthly datetime series starting from the first month of year *start_year*.

Return type

int

Examples

```
>>> get_date_index(2009, 14, start_year=1988)
265
```

`glomar_gridding.utils.get_month_midpoint(dates)`

Get the month midpoint for a series of datetimes.

The midpoint of a month is the exact half-way point between the start and end of the month.

For example, the midpoint of January 1990 is 1990-01-16 12:00.

Return type

Series

`glomar_gridding.utils.get_pentad_range(centre_date)`

Get the start and date of a pentad centred at a centre date. If the pentad includes the leap date of 29th Feb then the pentad will include 6 days. This follows the * pentad convention.

The start and end date are first calculated from a non-leap year.

If the centre date value is 29th Feb then the pentad will be a pentad starting on 27th Feb and ending on 2nd March.

Parameters

centre_date (date) – The centre date of the pentad. The start date will be 2 days before this date, and the end date will be 2 days after.

Return type

tuple[date, date]

Returns

- **start_date** (*datetime.date*) – Two days before *centre_date*
- **end_date** (*datetime.date*) – Two days after *centre_date*

Examples

```
>>> get_pentad_range(date(2008, 2, 29))
(datetime.date(2008, 2, 27), datetime.date(2008, 3, 2))
```

`glomar_gridding.utils.get_spatial_mean(grid_obs, covx)`

Compute the spatial mean accounting for auto-correlation. See [Cornell]

Parameters

- **grid_obs** (ndarray) – Vector containing observations
- **covx** (ndarray) – Observation covariance matrix

Returns

spatial_mean – The spatial mean defined as $(1^T x C^{-1} x 1)^{-1} * (1^T x C^{-1} x z)$

Return type

float

References

[Cornell] https://www.css.cornell.edu/faculty/dgr2/_static/files/distance_ed_geostats/ov5.pdf

`glomar_gridding.utils.init_logging(file=None, level='DEBUG')`

Initialise the logger

Parameters

- **file** (str | None) – File to send log messages to. If set to None (default) then print log messages to STDOUT
- **level** (str) – Level of logging, one of: “debug”, “info”, “warn”, “error”, “critical”.

Return type

None

`glomar_gridding.utils.intersect_mtlb(a, b)`

Returns data common between two arrays, a and b, in a sorted order and index vectors for a and b arrays Reproduces behaviour of Matlab’s intersect function.

Parameters

- **a** (ndarray)
- **b** (ndarray)

Returns

- Intersection
- List of indices, where the common values are located, for array a
- List of indices, where the common values are located, for array b

Return type

tuple[ndarray, ndarray, ndarray]

Examples

```
>>> a = np.array([1, 2, 3])
>>> b = np.array([1, 1, 2, 5, 6])
>>> intersect_mtlb(a, b)
(array([1, 2]), array([0, 1]), array([0, 2]))
```

`glomar_gridding.utils.is_iter(val)`

Determine if a value is an iterable

Return type

bool

`glomar_gridding.utils.km_to_deg(km)`

Convert meridional km change to degree latitude

Parameters

km (float) – The meridional difference in kilometers

Returns

The meridional difference in degrees

Return type

float

`glomar_gridding.utils.make_diag(arr)`

Check arr if it is 1D or 2D square. If 1D, returns itself. If 2D-square, returns diagonal (i.e. vector). Anything else (2D non-square or 3+D), raise exception.

Parameters

arr (ndarray) – An array to be checked if it is 1D, 2D or some other dimensions Some classes like `numpy.ma.MaskedArray`, `scipy.sparse.sparray` and even `pandas.DataFrame` and `polars.DataFrame` also have the `shape` attribute, and would work. Shape should be (N,) or (N, N); otherwise, exception will be raised.

Returns

A 1D vector of shape (N,).

Return type

ndarray

`glomar_gridding.utils.make_square(arr)`

Check arr if it is 1D or 2D square. If 1D, returns `np.diag` of it (i.e. square matrix). If 2D square, returns itself. Anything else (2D non-square or 3+D), raise exception.

Parameters

arr (ndarray) – An array to be checked if it is 1D, 2D or some other dimension. Some classes like `numpy.ma.MaskedArray`, `scipy.sparse.sparray` and even `pandas.DataFrame` and `polars.DataFrame` also have the `shape` attribute, and would work. Shape should be (N,) or (N, N); otherwise, exception will be raised.

Returns

A 2D array of shape (N, N).

Return type

ndarray

`glomar_gridding.utils.mask_array(arr)`

Forces numpy array to be an instance of `np.ma.MaskedArray`

Parameters

arr (ndarray) – Can be masked or not masked

Returns

arr – array is now an instance of `np.ma.MaskedArray`

Return type

MaskedArray

`glomar_gridding.utils.select_bounds(x, bounds=[(-90, 90), (-180, 180)], variables=['lat', 'lon'])`

Filter an xarray.DataArray or xarray.Dataset by a set of bounds.

Parameters

- **x** (TypeVar(_XR_Data, DataArray, Dataset)) – The data to filter
- **bounds** (list[tuple[float, float]]) – A list of tuples containing the lower and upper bounds for each dimension.
- **variables** (list[str]) – Names of the dimensions (the order must match the bounds).

Returns

x – The input data filtered by the bounds.

Return type

TypeVar(_XR_Data, DataArray, Dataset)

`glomar_gridding.utils.sizeof_fmt(num, suffix='B')`

Convert numbers to kilo/mega... bytes, for interactive printing of code progress.

Parameters

- **num** (float) – The number (typically of bytes) to format
- **suffix** (str)

Returns

The formatted number using power of 1024 base.

Return type

str

Examples

```
>>> sizeof_fmt(123456789)
'117.7MiB'
```

`glomar_gridding.utils.uncompress_masked(compressed_array, mask, fill_value=0.0, apply_mask=False, dtype=None)`

Un-compress a compressed array using a mask.

Parameters

- **compressed_array** (ndarray) – The compressed array, originally compressed by the mask
- **mask** (ndarray) – The mask - a boolean numpy array
- **fill_value** (Any) – The value to fill masked points. If *apply_mask* is set, then this will be removed in the output.
- **apply_mask** (bool) – Apply the mask to the result, returning a MaskedArray rather than a ndarray.
- **dtype** (type | None) – Optionally set a dtype for the returned array, if not set then the dtype of the *compressed_array* is used.

Returns

uncompressed – The uncompressed array, masked points are filled with the *fill_value* if *apply_mask* is False. If *apply_mask* is True, then the result is an instance of `numpy.ma.MaskedArray` with the mask applied to the uncompressed result.

Return type
ndarray | MaskedArray

Examples

```
>>> arr = np.random.rand(16)
>>> mask = arr > 0.65
>>> arr = np.ma.masked_where(mask, arr).compressed()
>>> uncompress_masked(arr, mask, fill_value=-999.0)
array([ 2.79245414e-01, -9.99000000e+02,  3.93541024e-01, -9.99000000e+02,
        8.07814120e-03,  3.34164220e-01, -9.99000000e+02,  2.08200564e-01,
        3.32044850e-01,  1.83166093e-01, -9.99000000e+02,  2.57339943e-02,
        1.76017461e-01,  3.56673893e-01,  1.59393168e-01,  2.17047382e-01])
```

13.7 Variable Transformation

Transformation of variables including to approximately normal

`glomar_gridding.transformation.inv_weibull_to_normality(x_hat, c, use_kp11=False)`

The inverse of weibull_to_normality

Parameters

- **x_hat** (ndarray) – Variable that have been transformed
- **c** (float) – Weibull shape parameter
- **use_kp11** (bool) – Use the approximation described in Kulkarni and Powar 2011

Returns

x – Inverse-transformed x_hat

Return type

numpy.ndarray

`glomar_gridding.transformation.standardise_data(X, normalise=True, axis=0)`

Standardise the data samples by 1) subtracting the mean 2) divide by standard deviation

Parameters

- **X** (*np.ndarray*) – As in *data_sample* in *LassoEstimate_AR1* class; shape (N, M)
- **normalise** (*bool*) – If set to *True*, divide by sample standard deviation
- **axis** (*int*) – The axis that the standardisation is applied. *axis=0* follows the convention that each column is a separate time series, which is the standard for netCDF and GRIB files (i.e. time is usually the 1st dimension unless there is an ensemble dimension; spatial dimensions follow time dimension, aka the (T, YX...) convention).

Return type

tuple[ndarray, ndarray, ndarray]

Returns

- **X_standardised** (*numpy.ndarray*) – Standardised or zero-mean sample; shape (N, M)
- **sample_mean** (*numpy.ndarray*) – Sample mean of X computed along axis; shape (N,)
- **X_std** (*numpy.ndarray*) – The normalisation for *X_standardised*; if standardise is *True*, this is a vector of standard deviations; otherwise it a vector of 1s; shape (N,)

`glomar_gridding.transformation.weibull_to_normality(x, c, use_kp11=False)`

Box Cox transform to make Weibull-distributed variable to near normal. This includes wind speeds, which is often modelled with Weibull. The transformation lambda parameter depends on the shape parameter.

Reference: <https://doi.org/10.2307/2287172>

$$\hat{x} = \frac{x^\lambda - 1}{\lambda}$$

$$\lambda \approx 0.2654 \times \text{shape, if } x \sim \text{Weibull}(\text{shape}, \text{scale})$$

A variation to the above is proposed by Kulkarni and Powar 2011: <https://doi.org/10.1155/2011/863274>

$$\hat{x} = x^\lambda$$

$$\lambda \approx 0.2776 \times \text{shape}$$

Parameters

- **x** (ndarray) – Variable that needs transformation
- **c** (float) – Weibull shape parameter
- **use_kp11** (bool) – Use the approximation described in Kulkarni and Powar 2011

Returns

x_hat – Approximate normal transformed x

Return type

numpy.ndarray

BIBLIOGRAPHY

- [Bun] Bun, J., Bouchaud, J.-P., and Potters, M.: Cleaning large correlation matrices: Tools from Random Matrix Theory, *Physics Reports*, Volume 666, 2017, Pages 1-109, ISSN 0370-1573, <https://doi.org/10.1016/j.physrep.2016.10.005>.
- [Cornell] https://www.css.cornell.edu/faculty/dgr2/_static/files/distance_ed_geostats/ov5.pdf
- [Cressie] Cressie, N. A. C.: *Statistics for Spatial Data*, Wiley Series in Probability and Statistics, Wiley, 1 edn., ISBN 978-0-471-00255-0 978-1-119-11515-1, <https://doi.org/10.1002/9781119115151>, 1993.
- [Nychka] Nychka, D., Furrer, R., Paige, J. and Sain, S.: Package ‘fields’, <http://cran.r-project.org/web/packages/fields/fields.pdf>, 2015.
- [Higham] Higham, N. J., Strabić, N., and Šego, V.: Restoring Definiteness via Shrinking, with an Application to Correlation Matrices with a Fixed 255 Block, *SIAM Review*, 58, 245–263, <https://www.jstor.org/stable/24778894>, 2016.
- [Higham_blog_nearest] <https://nhigham.com/2021/01/26/what-is-the-nearest-positive-semidefinite-matrix/>
- [Huang] Huang, B., Yin, X., Menne, M. J., Vose, R., and Zhang, H.-M.: Improvements to the Land Surface Air Temperature Reconstruction in NOAA GlobalTemp: An Artificial Neural Network Approach, *Artificial Intelligence for the Earth Systems*, 1, e220 032, <https://doi.org/10.1175/AIES-D-22-0032.1>, 2022.
- [Hutchinson] Hutchinson, M. F., and Gessler, P. E.: Splines - more than just a smooth interpolator, *Geoderma*, 62(1-3), 45-67, [https://doi.org/10.1016/0016-7061\(94\)90027-2](https://doi.org/10.1016/0016-7061(94)90027-2), 1994.
- [IPCC] Intergovernmental Panel On Climate Change: *Climate Change 2021 – The Physical Science Basis: Working Group I Contribution to the Sixth Assessment Report of the Intergovernmental Panel on Climate Change*, Cambridge University Press, 1 edn., ISBN 978-1-00-915789-6, <https://doi.org/10.1017/9781009157896>, 2021.
- [Jolliffe] Jolliffe, I. T., *Principal Component Analysis*, pp 487, Springer, 2002
- [Kadow] Kadow, C., Hall, D. M., and Ulbrich, U.: Artificial intelligence reconstructs missing climate information, *Nature Geoscience*, 13, 408–413, <https://doi.org/10.1038/s41561-020-0582-5>, 2020
- [Karspeck] Karspeck, A. R., Kaplan, A., and Sain, S. R.: Bayesian modelling and ensemble reconstruction of mid-scale spatial variability in North Atlantic sea-surface temperatures for 1850-2008, *Quarterly Journal of the Royal Meteorological Society*, 138, 234–248, <https://doi.org/10.1002/qj.900>, 2012.
- [Laloux] Laloux, L., Cizeau, P., Potters, M. and Bouchaud, J.P.: Random matrix theory and financial correlations. *International Journal of Theoretical and Applied Finance*, 3(03), 391-397. <https://www.worldscientific.com/doi/abs/10.1142/S0219024900000255>, 2000
- [Lenssen] Lenssen, N. J. L., Schmidt, G. A., Hansen, J. E., Menne, M. J., Persin, A., Ruedy, R., and Zyss, D.: Improvements in the GISTEMP Uncertainty Model, *Journal of Geophysical Research: Atmospheres*, 124, 6307–6326, <https://doi.org/10.1029/2018JD029522>, 2019.

- [Morice_2012] Morice, C. P., Kennedy, J. J., Rayner, N. A., and Jones, P. D.: Quantifying uncertainties in global and regional temperature change using an ensemble of observational estimates: The HadCRUT4 data set: THE HADCRUT4 DATASET, *Journal of Geophysical Research: Atmospheres*, 117, n/a–n/a, <https://doi.org/10.1029/2011JD017187>, 2012
- [Morice_2021] Morice, C. P., Kennedy, J. J., Rayner, N. A., Winn, J. P., Hogan, E., Killick, R. E., Dunn, R. J. H., Osborn, T. J., Jones, P. D., and Simpson, 285 I. R.: An Updated Assessment of Near-Surface Temperature Change From 1850: The HadCRUT5 Data Set, *Journal of Geophysical Research: Atmospheres*, 126, e2019JD032 361, <https://doi.org/10.1029/2019JD032361>, 2021.
- [North] North, G. R., Bell, T. L., Cahalan, R. F., and Moeng, F. J.: Sampling Errors in the Estimation of Empirical Orthogonal Functions. *Monthly Weather Review*, 110(7), 699–706. [https://doi.org/10.1175/1520-0493\(1982\)110<0699:SEITEO>2.0.CO;2](https://doi.org/10.1175/1520-0493(1982)110<0699:SEITEO>2.0.CO;2), 1982.
- [PaciorekSchervish] Paciorek, C. J. and Schervish, M. J.: Spatial modelling using a new class of nonstationary covariance functions, *Environmetrics*, 17, 483–506, <https://doi.org/10.1002/env.785>, 2006
- [Rasmussen] Rasmussen, C. E. and Williams, C. K. I.: *Gaussian Processes for Machine Learning*, The MIT Press, ISBN 978-0-262-25683-4, <https://doi.org/10.7551/mitpress/3206.001.0001>, 2005.
- [RohdeBerkeley] Rohde, R. A. and Hausfather, Z.: The Berkeley Earth Land/Ocean Temperature Record, *Earth System Science Data*, 12, 3469–3479, <https://doi.org/10.5194/essd-12-3469-2020>, 2020
- [Thorne] Thorne, P. W., D. E. Parker, J. R. Christy, and C. A. Mears: Uncertainties in climate trends: Lessons from upper-air temperature records, *Bull. Am. Meteorol. Soc.*, 86, 1437–1442, <https://doi.org/10.1175/BAMS-86-10-1437>, 2005
- [Wahba_Sphere] Wahba, G: Spline Interpolation and Smoothing on the Sphere, *SIAM Journal on Scientific and Statistical Computing*, 2(1), 5-16, <https://doi.org/10.1137/0902002>, 1981
- [Wilks] Wilks, D.: *Statistical Methods in Atmospheric Sciences*, 2nd Edition, Elsevier, 2006
- [Tibshirani_Lasso] Tibshirani, Robert. 1996. Regression Shrinkage and Selection Via the Lasso. *Journal of the Royal Statistical Society Series B: Statistical Methodology* 58(1):267–88. doi:10.1111/j.2517-6161.1996.tb02080.x.

PYTHON MODULE INDEX

g

- `glomar_gridding.autocorrelation.compute_autocorrelation,`
82
- `glomar_gridding.autoregressive_kalman.cov_diagonal,`
81
- `glomar_gridding.climatology,` 89
- `glomar_gridding.covariance_tools,` 97
- `glomar_gridding.distances,` 93
- `glomar_gridding.error_covariance,` 33
- `glomar_gridding.interpolation_covariance,` 19
- `glomar_gridding.io,` 104
- `glomar_gridding.mask,` 90
- `glomar_gridding.transformation,` 113
- `glomar_gridding.utils,` 105

A

`A_matrix()` (*glomar_gridding.spline.Spline method*), 54
`add_empty_layers()` (in module *glomar_gridding.utils*), 105
`add_mask()` (*glomar_gridding.grid.Grid method*), 66
`adjust_small_negative()` (in module *glomar_gridding.utils*), 106
`assign_to_grid()` (in module *glomar_gridding.grid*), 51
`assign_values()` (*glomar_gridding.grid.Grid method*), 66
`Autoregressive1ForecastUncorr` (class in *glomar_gridding.autoregressive_kalman*), 75
`Autoregressive1ForecastVector` (class in *glomar_gridding.autoregressive_kalman*), 76

B

`batched()` (in module *glomar_gridding.utils*), 106

C

`c_ij_anisotropic_array()` (*glomar_gridding.ellipse.EllipseCovarianceBuilder method*), 30
`c_ij_isotropic_array()` (*glomar_gridding.ellipse.EllipseCovarianceBuilder method*), 30
`calc_cov()` (*glomar_gridding.ellipse.EllipseBuilder method*), 23
`calculate_cor()` (*glomar_gridding.ellipse.EllipseCovarianceBuilder method*), 31
`calculate_covariance_array()` (*glomar_gridding.ellipse.EllipseCovarianceBuilder method*), 31
`calculate_covariance_batched()` (*glomar_gridding.ellipse.EllipseCovarianceBuilder method*), 31
`calculate_covariance_loop()` (*glomar_gridding.ellipse.EllipseCovarianceBuilder method*), 31
`calculate_distance_matrix()` (in module *glomar_gridding.distances*), 93

`check_cols()` (in module *glomar_gridding.utils*), 106
`check_shape()` (in module *glomar_gridding.utils*), 106
`check_sq_matrix()` (in module *glomar_gridding.utils*), 106
`check_symmetric()` (in module *glomar_gridding.covariance_tools*), 98
`clean_small()` (in module *glomar_gridding.covariance_tools*), 98
`coefficients` (*glomar_gridding.autocorrelation.LassoEstimate_AR1 attribute*), 87
`ColumnNotFoundError`, 105
`compute_lag_ts_metrics_from_da()` (in module *glomar_gridding.autocorrelation.compute_autocorrelation*), 82
`compute_params()` (*glomar_gridding.ellipse.EllipseBuilder method*), 23
`constraint_mask()` (*glomar_gridding.kriging.OrdinaryKriging method*), 44
`constraint_mask()` (*glomar_gridding.kriging.SimpleKriging method*), 41
`constraint_mask()` (*glomar_gridding.stochastic.StochasticKriging method*), 48
`coord_df` (*glomar_gridding.grid.Grid property*), 66
`coord_names` (*glomar_gridding.grid.Grid property*), 66
`coords` (*glomar_gridding.grid.Grid property*), 67
`cor_2_cov()` (in module *glomar_gridding.utils*), 107
`correlated_components()` (in module *glomar_gridding.error_covariance*), 33
`cov_2_cor()` (in module *glomar_gridding.utils*), 107
`covariance` (*glomar_gridding.grid.Grid property*), 67
`covariance_matrix()` (*glomar_gridding.grid.Grid method*), 67
`csum_up_to_val()` (in module *glomar_gridding.covariance_tools*), 98

D

`D` (*glomar_gridding.autocorrelation.LassoEstimate_AR1 attribute*), 87

`days_since_by_month()` (in module `glomar_gridding.utils`), 107
`default_error_cov` (`glomar_gridding.spline.SphericalThinPlateSpline` property), 60
`default_error_cov` (`glomar_gridding.spline.Spline` property), 54
`default_error_cov` (`glomar_gridding.spline.ThinPlateSpline` property), 58
`deg_to_km()` (in module `glomar_gridding.utils`), 107
`deg_to_nm()` (in module `glomar_gridding.utils`), 107
`diag_and_nondiag_rows_subsamplers()` (in module `glomar_gridding.autoregressive_kalman.cov_diagonal`), 81
`displacements()` (in module `glomar_gridding.distances`), 94
`dist_func()` (`glomar_gridding.spline.SphericalThinPlateSpline` method), 60
`dist_func()` (`glomar_gridding.spline.Spline` method), 54
`dist_func()` (`glomar_gridding.spline.ThinPlateSpline` method), 58
`dist_weight()` (in module `glomar_gridding.error_covariance`), 34
`distance_matrix()` (`glomar_gridding.grid.Grid` method), 67
`draw_from_cov()` (in module `glomar_gridding.stochastic`), 51
E
`eigenvalue_clip()` (in module `glomar_gridding.covariance_tools`), 99
`EllipseBuilder` (class in `glomar_gridding.ellipse`), 23
`EllipseCovarianceBuilder` (class in `glomar_gridding.ellipse`), 29
`EllipseModel` (class in `glomar_gridding.ellipse`), 21
`errcov` (`glomar_gridding.autocorrelation.LassoEstimate_AR1` attribute), 86
`errcov` (`glomar_gridding.autoregressive_kalman.AutoregressiveIForecastGriding` attribute), 76
`errcov` (`glomar_gridding.autoregressive_kalman.AutoregressiveIForecastGriding` attribute), 78
`errcov` (`glomar_gridding.autoregressive_kalman.KalmanOnForecastGriding` attribute), 80
`error` (`glomar_gridding.spline.SplineResult` attribute), 63
`error_cov` (`glomar_gridding.spline.Spline` property), 55
`estimate_errcov()` (`glomar_gridding.autocorrelation.LassoEstimate_AR1` method), 86
`estimate_lambda_gcv()` (`glomar_gridding.spline.Spline` method), 55
`euclidean_distance()` (in module `glomar_gridding.distances`), 94
`expand_coefficients()` (`glomar_gridding.autocorrelation.LassoEstimate_AR1` method), 86
`expanded` (`glomar_gridding.autocorrelation.LassoEstimate_AR1` attribute), 87
`explained_variance_clip()` (in module `glomar_gridding.covariance_tools`), 100
`ExponentialVariogram` (class in `glomar_gridding.variogram`), 16
`extended_inverse()` (`glomar_gridding.kriging.OrdinaryKriging` method), 45
F
`filter_bounds()` (in module `glomar_gridding.utils`), 108
`find_nearest()` (in module `glomar_gridding.utils`), 108
`find_nearest_xy_index_in_cov_matrix()` (`glomar_gridding.ellipse.EllipseBuilder` method), 25
`fit()` (`glomar_gridding.autocorrelation.LassoEstimate_AR1` method), 87
`fit()` (`glomar_gridding.ellipse.EllipseModel` method), 21
`fit()` (`glomar_gridding.spline.SphericalThinPlateSpline` method), 61
`fit()` (`glomar_gridding.spline.Spline` method), 55
`fit()` (`glomar_gridding.spline.ThinPlateSpline` method), 58
`fit()` (`glomar_gridding.variogram.ExponentialVariogram` method), 16
`fit()` (`glomar_gridding.variogram.GaussianVariogram` method), 17
`fit()` (`glomar_gridding.variogram.MaternVariogram` method), 18
`fit()` (`glomar_gridding.variogram.SphericalVariogram` method), 16
`fit_ellipse_model()` (`glomar_gridding.ellipse.EllipseBuilder` method), 25
`five_knots()` (`glomar_gridding.spline.SplineResult` method), 63
`forecast` (`glomar_gridding.autoregressive_kalman.AutoregressiveIForecastGriding` attribute), 76
`forecast` (`glomar_gridding.autoregressive_kalman.AutoregressiveIForecastGriding` attribute), 78
`from_autocov()` (`glomar_gridding.autoregressive_kalman.AutoregressiveIForecastVector` class method), 77
`from_resolution()` (`glomar_gridding.grid.Grid` class method), 68

G

- GaussianVariogram (class in glo-mar_gridding.variogram), 17
- gcv (glomar_gridding.spline.SplineResult attribute), 63
- get_anomalies() (in module glo-mar_gridding.autocorrelation.compute_autocorrelation), 83
- get_anomalies_vectorize() (in module glo-mar_gridding.autocorrelation.compute_autocorrelation), 83
- get_auto_corr() (in module glo-mar_gridding.autocorrelation.compute_autocorrelation), 83
- get_auto_corr_vectorize() (in module glo-mar_gridding.autocorrelation.compute_autocorrelation), 84
- get_date_index() (in module glomar_gridding.utils), 108
- get_ellipse_params() (in module glo-mar_gridding.ellipse.estimate), 27
- get_K() (glomar_gridding.spline.Spline method), 56
- get_kriging_weights() (glomar_gridding.kriging.OrdinaryKriging method), 45
- get_kriging_weights() (glomar_gridding.kriging.SimpleKriging method), 42
- get_kriging_weights() (glomar_gridding.stochastic.StochasticKriging method), 48
- get_mask_idx() (in module glomar_gridding.mask), 90
- get_month_midpoint() (in module glo-mar_gridding.utils), 109
- get_obs_variance() (in module glo-mar_gridding.autocorrelation.compute_autocorrelation), 84
- get_obs_variance_vectorize() (in module glo-mar_gridding.autocorrelation.compute_autocorrelation), 85
- get_pentad_range() (in module glo-mar_gridding.utils), 109
- get_recurse() (in module glomar_gridding.io), 104
- get_spatial_mean() (in module glo-mar_gridding.utils), 109
- get_stats() (glomar_gridding.spline.Spline method), 56
- get_trend_basis() (glomar_gridding.spline.SphericalThinPlateSpline method), 61
- get_trend_basis() (glomar_gridding.spline.Spline method), 56
- get_trend_basis() (glomar_gridding.spline.ThinPlateSpline method), 58
- get_uncertainty() (glomar_gridding.kriging.OrdinaryKriging method), 46
- get_uncertainty() (glomar_gridding.kriging.SimpleKriging method), 42
- get_uncertainty() (glomar_gridding.stochastic.StochasticKriging method), 49
- get_weights() (in module glo-mar_gridding.error_covariance), 34
- glomar_gridding.autocorrelation.compute_autocorrelation module, 82
- glomar_gridding.autoregressive_kalman.cov_diagonal module, 81
- glomar_gridding.climatology module, 89
- glomar_gridding.covariance_tools module, 97
- glomar_gridding.distances module, 93
- glomar_gridding.error_covariance module, 33
- glomar_gridding.interpolation_covariance module, 19
- glomar_gridding.io module, 104
- glomar_gridding.mask module, 90
- glomar_gridding.transformation module, 113
- glomar_gridding.utils module, 105
- Grid (class in glomar_gridding.grid), 65
- grid_box_weighted_sum() (in module glo-mar_gridding.error_covariance), 35
- grid_from_resolution() (in module glo-mar_gridding.grid), 14
- grid_idx (glomar_gridding.grid.Grid property), 69
- grid_to_distance_matrix() (in module glo-mar_gridding.grid), 14

H

- haversine_distance_from_frame() (in module glo-mar_gridding.distances), 95
- haversine_gaussian() (in module glo-mar_gridding.distances), 95

I

- index_map (glomar_gridding.grid.Grid property), 69
- init_logging() (in module glomar_gridding.utils), 110
- intersect_mtlb() (in module glomar_gridding.utils), 110
- inv_2d() (in module glomar_gridding.distances), 96

`inv_weibull_to_normality()` (in module `glomar_gridding.transformation`), 113
`is_iter()` (in module `glomar_gridding.utils`), 110

J

`join_climatology_by_doy()` (in module `glomar_gridding.climatology`), 89

K

`kalman_gain_from_new_obs` (`glomar_gridding.autoregressive_kalman.KalmanOut` attribute), 80
`KalmanOut` (class in `glomar_gridding.autoregressive_kalman`), 79
`KalmanOutUncorrCorrSplit` (class in `glomar_gridding.autoregressive_kalman`), 80
`kernel()` (`glomar_gridding.spline.SphericalThinPlateSpline` method), 61
`kernel()` (`glomar_gridding.spline.Spline` method), 56
`kernel()` (`glomar_gridding.spline.ThinPlateSpline` method), 59
`km_to_deg()` (in module `glomar_gridding.utils`), 111
`kriging()` (`glomar_gridding.grid.Grid` method), 69
`kriging_ordinary()` (in module `glomar_gridding.kriging`), 47
`kriging_simple()` (in module `glomar_gridding.kriging`), 43
`kriging_weights_from_inverse()` (`glomar_gridding.kriging.OrdinaryKriging` method), 46
`kriging_weights_from_inverse()` (`glomar_gridding.kriging.SimpleKriging` method), 42
`kriging_weights_from_inverse()` (`glomar_gridding.stochastic.StochasticKriging` method), 49

L

`laloux_clip()` (in module `glomar_gridding.covariance_tools`), 101
`LassoEstimate_AR1` (class in `glomar_gridding.autocorrelation`), 85
`load_array()` (in module `glomar_gridding.io`), 105
`load_covariance()` (in module `glomar_gridding.interpolation_covariance`), 19
`load_dataset()` (in module `glomar_gridding.io`), 105

M

`mahal_dist_func()` (in module `glomar_gridding.distances`), 96
`make_coeff_sparse()` (`glomar_gridding.autocorrelation.LassoEstimate_AR1` method), 87
`make_diag()` (in module `glomar_gridding.utils`), 111
`make_square()` (in module `glomar_gridding.utils`), 111
`map_observations()` (`glomar_gridding.grid.Grid` method), 69
`map_to_grid()` (in module `glomar_gridding.grid`), 39
`mask_array()` (in module `glomar_gridding.mask`), 91
`mask_array()` (in module `glomar_gridding.utils`), 111
`mask_dataset()` (in module `glomar_gridding.mask`), 91
`mask_from_obs_array()` (in module `glomar_gridding.mask`), 91
`mask_from_obs_frame()` (in module `glomar_gridding.mask`), 92
`mask_observations()` (in module `glomar_gridding.mask`), 92
`MaternVariogram` (class in `glomar_gridding.variogram`), 17

module

`glomar_gridding.autocorrelation.compute_autocorrelation`, 82
`glomar_gridding.autoregressive_kalman.cov_diagonal`, 81
`glomar_gridding.climatology`, 89
`glomar_gridding.covariance_tools`, 97
`glomar_gridding.distances`, 93
`glomar_gridding.error_covariance`, 33
`glomar_gridding.interpolation_covariance`, 19
`glomar_gridding.io`, 104
`glomar_gridding.mask`, 90
`glomar_gridding.transformation`, 113
`glomar_gridding.utils`, 105
`MonthName` (class in `glomar_gridding.utils`), 105
`msr` (`glomar_gridding.spline.SplineResult` attribute), 63

N

`negative_log_likelihood()` (`glomar_gridding.ellipse.EllipseModel` method), 22
`negative_log_likelihood_function()` (`glomar_gridding.ellipse.EllipseModel` method), 23

O

`OrdinaryKriging` (class in `glomar_gridding.kriging`), 44

P

`perturb_cov_to_positive_definite()` (in module `glomar_gridding.covariance_tools`), 102
`predict()` (`glomar_gridding.autoregressive_kalman.AutoregressiveIForec` method), 76
`predict()` (`glomar_gridding.autoregressive_kalman.AutoregressiveIForec` method), 78

predict() (glomar_gridding.autoregressive_kalman.KalmanOutUncorrCorrSpline method), 79
 predict() (glomar_gridding.autoregressive_kalman.KalmanOutUncorrCorrSpline method), 81
 predict() (glomar_gridding.spline.SphericalThinPlateSpline method), 61
 predict() (glomar_gridding.spline.Spline method), 56
 predict() (glomar_gridding.spline.ThinPlateSpline method), 59
 predict_corr() (glomar_gridding.autoregressive_kalman.KalmanOutUncorrCorrSpline method), 81
 predict_uncorr() (glomar_gridding.autoregressive_kalman.KalmanOutUncorrCorrSpline method), 81
 prep_covariance() (glomar_gridding.grid.Grid method), 70
 prep_obs_for_kriging() (in module glomar_gridding.kriging), 40
R
 radial_dist() (in module glomar_gridding.distances), 96
 read_climatology() (in module glomar_gridding.climatology), 90
 remove_diag_only_rows() (in module glomar_gridding.autoregressive_kalman.cov_diagonal), 81
 remove_mask() (glomar_gridding.grid.Grid method), 70
 residues (glomar_gridding.autocorrelation.LassoEstimate_AR1 attribute), 87
 restore_diag_only_rows() (in module glomar_gridding.autoregressive_kalman.cov_diagonal), 82
 result (glomar_gridding.spline.Spline property), 56
 rot_mat() (in module glomar_gridding.distances), 96
S
 select_bounds() (glomar_gridding.grid.Grid method), 71
 select_bounds() (in module glomar_gridding.utils), 111
 set_covariance() (glomar_gridding.grid.Grid method), 71
 set_distance_matrix() (glomar_gridding.grid.Grid method), 71
 set_simple_kriging_weights() (glomar_gridding.stochastic.StochasticKriging method), 49
 shape (glomar_gridding.grid.Grid property), 71
 sigma2 (glomar_gridding.spline.SplineResult attribute), 63
 sigma_rot_func() (in module glomar_gridding.distances), 96
 SimpleKriging (class in glomar_gridding.kriging), 41
 size (glomar_gridding.grid.Grid property), 71
 sizeof_fmt() (in module glomar_gridding.utils), 112
 solve() (glomar_gridding.kriging.OrdinaryKriging method), 46
 solve() (glomar_gridding.kriging.SimpleKriging method), 43
 solve() (glomar_gridding.stochastic.StochasticKriging method), 50
 sparse_approx_for_errcov() (glomar_gridding.autoregressive_kalman.KalmanOutUncorrCorrSpline method), 80
 SphericalThinPlateSpline (class in glomar_gridding.spline), 60
 SphericalVariogram (class in glomar_gridding.variogram), 16
 Spline (class in glomar_gridding.spline), 54
 SplineResult (class in glomar_gridding.spline), 62
 split_holdout() (glomar_gridding.autocorrelation.LassoEstimate_AR1 method), 87
 standard_error() (glomar_gridding.spline.SphericalThinPlateSpline method), 62
 standard_error() (glomar_gridding.spline.Spline method), 57
 standard_error() (glomar_gridding.spline.ThinPlateSpline method), 59
 standardise_data() (in module glomar_gridding.transformation), 113
 StochasticKriging (class in glomar_gridding.stochastic), 47
T
 tau2 (glomar_gridding.spline.SplineResult attribute), 63
 tau_dist() (in module glomar_gridding.distances), 97
 tau_dist_from_frame() (in module glomar_gridding.distances), 97
 ThinPlateSpline (class in glomar_gridding.spline), 57
 to_xarray_da_coefficients() (glomar_gridding.autocorrelation.LassoEstimate_AR1 method), 87
 to_xarray_da_errcov() (glomar_gridding.autocorrelation.LassoEstimate_AR1 method), 87
 trend_basis (glomar_gridding.spline.Spline property), 57

U

`uncompress_cov()` (*glomar_gridding.ellipse.EllipseCovarianceBuilder* method), 31

`uncompress_masked()` (in module *glomar_gridding.utils*), 112

`uncorrelated_components()` (in module *glomar_gridding.error_covariance*), 36

`unexpand_coefficients()` (*glomar_gridding.autocorrelation.LassoEstimate_AR1* method), 88

V

`validate_covariance()` (in module *glomar_gridding.covariance_tools*), 104

`var` (*glomar_gridding.spline.SplineResult* attribute), 63

`var_stability_check()` (*glomar_gridding.autoregressive_kalman.AutoregressiveIForecastVector* method), 78

`variogram_to_covariance()` (in module *glomar_gridding.variogram*), 19

W

`weibull_to_normality()` (in module *glomar_gridding.transformation*), 113

`weighted_sum()` (in module *glomar_gridding.error_covariance*), 36

`wgt_mean` (*glomar_gridding.autoregressive_kalman.KalmanOut* attribute), 79

`wgts_from_ar_forecast` (*glomar_gridding.autoregressive_kalman.KalmanOut* attribute), 80